

# Informal learning in everyday human–LLM interaction

Zixin Chen<sup>1</sup>, Haotian Li<sup>2\*</sup>, Ziang Xiao<sup>3</sup>, Huamin Qu<sup>1</sup>, Xing Xie<sup>2</sup>

<sup>1</sup>The Hong Kong University of Science and Technology, Hong Kong, China.

<sup>2</sup>Microsoft Research Asia, Beijing, China.

<sup>3</sup>Johns Hopkins University, Baltimore, 21218, Maryland, United States.

\*Corresponding author(s). E-mail(s): [haotian.li@microsoft.com](mailto:haotian.li@microsoft.com);  
Contributing authors: [zchendf@connect.ust.hk](mailto:zchendf@connect.ust.hk); [ziang.xiao@jhu.edu](mailto:ziang.xiao@jhu.edu);  
[huamin@cse.ust.hk](mailto:huamin@cse.ust.hk); [xing.xie@microsoft.com](mailto:xing.xie@microsoft.com);

## Abstract

As LLMs become increasingly capable of completing tasks for users, a central concern is that working with them may erode the opportunities through which people develop their own capabilities. We analyse large scale human–LLM conversations to surface how users engage with LLMs in learning-oriented ways, revealing the prevalence and interactional mechanisms through which informal learning may occur in everyday AI use. Across 128,569 naturalistic conversations, we translated learning-science constructs into turn-level behavioural signatures. Cognitive engagement—visible uptake of assistant output—appeared in 31.9% of 491,685 user turns, whereas constructive engagement, the deepest observable form of learning-oriented engagement, appeared in 4.9%, showing that deeper sense-making was recurrent but selective. Our study further identifies factors associated with these forms of engagement. Scaffolded assistant support, present in 31.7% of assistant turns, consistently marked richer constructive participation, with associations varying by user framing, task ecology, support form, timing and prior user state. These findings make behavioural signatures of informal learning observable at scale and shift AI evaluation: from answer-delivery efficiency toward the preservation of cognitive opportunities for users to reason, test ideas and construct understanding in the course of everyday problem-solving.

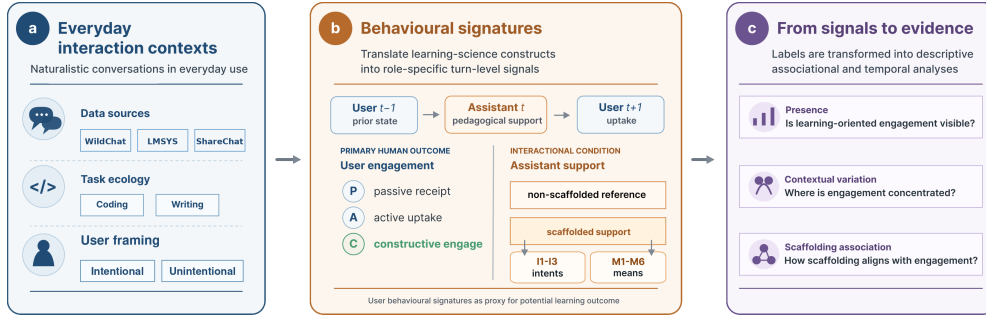
**Keywords:** human–AI interaction, informal learning, large language models, scaffolding, learner engagement

# 1 Introduction

Large language models (LLMs) are becoming embedded in the cognitive work of everyday life: people increasingly turn to them to write, code, search and reason through unfamiliar problems [1–3]. Much of human competence, however, is acquired precisely by doing such work: reasoning through a problem, attempting a solution, reworking it until it holds [4, 5]. A growing body of evidence warns that delegating that effort to machines can erode the very capabilities the work would otherwise build [6, 7]. Whether working with an assistant diminishes *users’ own cognitive work* or instead becomes an occasion for thought depends not on the technology but on how people engage with it: they may take an answer and move on, or use it to test their own ideas, seek out explanations and pursue a question to its root [8, 9]. Everyday human–LLM interaction has thus become a vast and mostly undesigned setting for informal learning—and, as AI does more of this work, an increasingly consequential one for whether people keep learning and thinking for themselves [10–12].

Yet this informal learning has gone largely unexamined. Existing research on LLMs and learning has focused instead on settings where learning support is deliberately designed, such as automated feedback, question generation, retrieval practice and tutoring systems [13–17]. These studies show that LLMs can enact instructional strategies when learning is the explicit design target [18, 19]; they tell us less about general-purpose LLM use, where users usually seek task progress rather than instruction. Everyday use is also harder to study: it follows no curriculum and yields no test score, so the learning at stake is easy to suspect but hard to see. Learning science offers a way to detect it: deeper forms of overt engagement are associated with stronger learning outcomes, so engagement can serve as an observable process indicator in informal, user-generated interactions whose goals and endpoints are not pre-specified by curricula or study protocols [20–22]. We therefore ask whether learning-oriented engagement is visible in everyday human–LLM interaction, where its deepest constructive form appears, and how assistant scaffolding is coupled with further constructive participation.

We analyse 128,569 naturalistic conversations from three public large-scale conversational corpora: WildChat, LMSYS Chat and ShareChat [23–25]. Across these corpora, we focus on coding and writing as two common task ecologies in which LLMs are used to produce, revise and troubleshoot knowledge work [1, 8, 26], yielding 491,685 user turns and 489,785 assistant turns. Our analyses move from engagement prevalence, to the contextual organization of constructive engagement, to how assistant scaffolding is associated with user participation at conversation-level, support-form and adjacent-turn scales. To make these questions measurable in naturalistic logs, we translate theories of engagement and scaffolding into role-specific, turn-level behavioural signature. User turns are coded by the depth of visible uptake of assistant output: *passive* receipt, *active* use or follow-up, and *constructive* engagement, where users elaborate, test, revise or extend ideas in the exchange [20, 21]. Assistant turns are coded by whether they primarily deliver the requested content or provide scaffolded support: contingent assistance that diagnoses difficulties, decomposes tasks, explains rationales, gives feedback or hints, or prompts reflection while leaving the user to continue the work [27–30]. Together, these signatures allow us to



**Fig. 1 Analytic framework for studying informal learning in everyday human-LLM interaction.** **a**, Public conversations are situated by data source, task ecology and user framing. **b**, Engagement and scaffolding constructs are translated into role-specific turn-level labels for user engagement and assistant support, with assistant support further characterized by support intent (I1-I3: metacognitive, cognitive and affective) and support form (M1-M6: feedback, hinting, instructing, explaining, modelling and questioning); Supplementary Table C1 provides plain-language examples of each label. **c**, These labels support analyses of engagement presence, contextual organization and support-engagement coupling at conversation-level, support-form and adjacent-turn scales. Labels indicate observable learning-oriented participation, not durable learning outcomes or causal instructional effects.

examine where users remain actively involved in reasoning with generated output, and how constructive participation is organized by task context, assistant support and local turn ordering.

Across the three corpora, everyday human-LLM interaction emerged as a genuine but selective and conditional site of learning. First, learning-oriented engagement was recurrent in everyday human-LLM interaction, but its deepest constructive form was selective: cognitive engagement appeared in nearly one third of user turns, whereas constructive engagement appeared in about five percent. Second, constructive engagement was ecologically organized: it was more common under explicit learning-oriented user framing, varied across task ecologies and became most visible in sustained exchanges where users returned to, interrogated and reshaped the assistant’s response. Third, assistant scaffolding marked richer constructive participation, but not as a uniform exposure. Its association depended on how support was positioned within the exchange: feedback-like and explanatory forms aligned most clearly with constructive engagement at the conversation level, while adjacent-turn analyses showed that support-engagement coupling in local sequences was shaped by prior user state and support form. Together, these findings locate everyday human-LLM interaction in a middle space between ordinary tool use and designed instruction. These results offer an empirical counterpoint to fears that AI use simply erodes human thinking: people have not ceded their thinking to the model but continue to engage constructively, if selectively. They suggest that the educational value of future AI assistants lies not in making every interaction tutor-like, but in calibrating when to answer, explain, give feedback or step back, preserving the cognitive agency through which people question, test, revise and build judgment while solving real problems.

**Table 1 Dataset and behavioural overview.** Summary statistics for the WildChat, LMSYS Chat and ShareChat task settings analysed in the study. Constructive engagement is reported as the deepest observable level within cognitive engagement and serves as the focal user-side measure in subsequent analyses.

| Setting           | Sample size   |            |                 | User framing          | Assistant support    | Cognitive engagement |                        | Emotional engagement |
|-------------------|---------------|------------|-----------------|-----------------------|----------------------|----------------------|------------------------|----------------------|
|                   | Conversations | User turns | Assistant turns | Intentional conv. (%) | Scaffolded turns (%) | Overall (%)          | Constructive level (%) | Emotional (%)        |
| WildChat coding   | 31,878        | 124,073    | 124,623         | 37.13                 | 51.03                | 50.21                | 9.23                   | 0.95                 |
| WildChat writing  | 39,534        | 163,705    | 164,824         | 25.80                 | 23.53                | 16.34                | 2.22                   | 0.57                 |
| LMSYS coding      | 32,114        | 107,065    | 105,174         | 41.49                 | 29.07                | 45.26                | 5.44                   | 0.91                 |
| LMSYS writing     | 21,023        | 77,906     | 76,279          | 20.64                 | 19.43                | 15.30                | 1.53                   | 0.21                 |
| ShareChat coding  | 2,481         | 11,541     | 11,522          | 45.63                 | 50.67                | 45.00                | 15.17                  | 3.38                 |
| ShareChat writing | 1,539         | 7,395      | 7,363           | 29.69                 | 22.42                | 33.05                | 5.18                   | 0.34                 |
| Total / pooled    | 128,569       | 491,685    | 489,785         | 32.12                 | 31.70                | 31.94                | 4.93                   | 0.74                 |

*Note.* User framing is reported as the percentage of conversations with explicit learning-oriented framing. Cognitive, constructive and emotional engagement are percentages of user turns; scaffolded support is computed over assistant turns. The constructive column uses all user turns as the denominator, whereas Fig. 2a decomposes cognitively engaged user turns into passive, active and constructive levels.

## 2 Results

The results follow the analytic sequence in Fig. 1: engagement prevalence, contextual organization and support–engagement coupling. We first quantify cognitive and constructive engagement in everyday human–LLM conversations, then locate constructive engagement across user framing, task ecology and interaction depth. We then turn to assistant support, asking whether scaffolded conversations contain richer constructive participation, which support forms carry this association and how support is ordered with prior and subsequent user engagement at the adjacent-turn scale.

### 2.1 Everyday human–LLM conversations contain learning-oriented engagement

Learning-oriented engagement was visible across all six task settings. Across 491,685 user turns, 31.9% showed cognitive engagement and 4.9% reached constructive engagement, the deepest observable user-side signal in the annotation scheme (Table 1). **Explicit affective expression was much less visible, (0.74% of user turns), consistent with the coding and writing focus of the sample.** Thus, everyday human–LLM conversations often contained observable user uptake of assistant responses, but deeper constructive sense-making remained selective rather than routine.

Figure 2a shows how cognitively engaged turns were composed. Active uptake was the dominant form in every setting, accounting for 62.0–86.4% of cognitively engaged turns, whereas constructive engagement accounted for 10.0–33.7% and passive receipt for 1.2–14.6%. This composition separates broad engagement from deeper sense-making. Cognitive engagement shows that everyday LLM use often involves users working with assistant output; active uptake captures use, application or follow-up; and constructive engagement identifies the cases in which users elaborate, test, revise or extend ideas. The remaining analyses therefore focus on where constructive engagement appears and how it is associated with user framing, task ecology, interaction depth and assistant support.

## 2.2 Constructive engagement is patterned by user framing, task ecology and sustained exchange

Having established that everyday human–LLM conversations contain learning-oriented engagement, we next examined where its constructive form became visible. Explicit learning-oriented framing provided the first contrast (Fig. 2b). In every setting, conversations framed around learning, understanding or exploration had higher cognitive and constructive user-turn rates than conversations without such framing. The cognitive-engagement lift ranged from +17.7 to +47.3 percentage points across settings, and the constructive-engagement lift ranged from +2.5 to +12.8 points. Yet constructive participation was also present in conversations without explicit learning-oriented framing. User framing therefore amplified constructive engagement, but did not define it: learning-relevant participation also emerged in task-oriented exchanges where learning was incidental to solving the problem at hand.

Task ecology provided a second source of variation. Across all three corpora, coding-oriented conversations showed higher cognitive and constructive engagement than writing-oriented conversations. Cognitive engagement ranged from 45.00–50.21% in coding conversations, compared with 15.30–33.05% in writing conversations; constructive engagement ranged from 5.44–15.17% in coding and 1.53–5.18% in writing. A conversation-level context model showed the same direction after including user framing, conversation-length bucket and dataset fixed effects, with coding task ecology positively associated with the presence of at least one constructive user turn (OR=1.89, 95% CI 1.71–2.08,  $p < .001$ ; Supplementary Table C12).

The coding–writing contrast points to task affordances that organize baseline opportunities for visible constructive participation. Coding tasks often externalize uncertainty through errors, constraints and executable tests, creating occasions for users to diagnose failures, refine assumptions, compare alternatives and test candidate solutions. Writing tasks also involve substantial revision and evaluation, but the interaction can often proceed through accepting, adapting or requesting a draft revision without requiring the user to articulate the rationale for the change [31–33]. The observed task difference is therefore consistent with the ecological interpretation: constructive engagement is more likely to appear when the task makes uncertainty, evaluation and revision explicit in the exchange.

The distinction between broad cognitive engagement and constructive engagement reinforces this task-ecology account. LMSYS Chat coding illustrates this distinction: it had high cognitive engagement, yet its cognitively engaged turns were dominated by active uptake rather than constructive elaboration (Fig. 2a). This pattern illustrates that users can work actively with assistant output—applying code, requesting repairs or asking for the next step—without making their reasoning, tests or revisions explicit enough to surface as constructive engagement. Taken together with the coding–writing contrast, this shows that task ecology patterned not only how often users worked with assistant output, but also whether that work became visible as articulated questioning, testing, refinement or extension.

Sustained exchange provided the third context in which constructive participation became visible (Fig. 2c). We operationalized sustained exchange with conversation-length buckets, asking whether a conversation contained at least one constructive

user turn. Across all six settings, conversations with more user turns were more likely to contain at least one constructive user turn. The gradient was steeper in coding-oriented settings, but the same direction appeared in writing-oriented settings. This pattern indicates that constructive engagement is most visible when everyday LLM use develops beyond a one-off request for an answer. In such exchanges, users have room to return to the assistant’s response, introduce new constraints, test implications, request revisions and reshape their understanding. Because this panel uses an any-event outcome, we treat the length gradient as a visibility pattern; a grouped-binomial constructive-rate sensitivity attenuated but preserved the positive length association (4–6 user turns: OR=1.13; 7+ user turns: OR=1.06; Supplementary Table C13).


Together, these results show that constructive engagement was ecologically organized rather than evenly distributed across everyday use. It was amplified by explicit learning-oriented framing, patterned by task affordances that create occasions for testing and revision, and most visible in sustained exchanges. A conversation-level context model with user framing, task ecology, conversation-length bucket and dataset fixed effects supported the same organization, and the rate sensitivity showed that the length association was attenuated but preserved after accounting for user-turn denominators (Supplementary Tables C12 and C13). Having located these user-side and contextual patterns, we next ask whether assistant scaffolding marks conversations with richer constructive participation.

### 2.3 Scaffolded support marks richer constructive participation

Having located the user-side and contextual conditions under which constructive engagement became visible, we next turned to assistant-side scaffolding: support that structures the user’s process rather than simply returning a deliverable, leaving room for diagnosis, repair or evaluation work [27–30].

Scaffolded support appeared in 31.7% of assistant turns (Table 1), making it a common interactional condition rather than a rare pedagogical exception. At the conversation level, conversations containing at least one scaffolded assistant turn had higher turn-weighted constructive user-turn ratios than reference conversations without scaffolded support in all six task settings (Fig. 3a). The constructive-ratio differences ranged from +0.99 to +7.34 percentage points across WildChat, LMSYS Chat and ShareChat, and all six contrasts were distinguished from zero in conversation-level bootstrap tests ( $p < .001$ ; Supplementary Table C10). Scaffolded conversations also showed more turns after the first assistant response in every setting (+1.42 to +3.38 turns; Fig. 3d), indicating that scaffolded support marked exchanges with more sustained participation. Together, these descriptive contrasts show that scaffolded support co-occurred with richer constructive participation across corpora and task settings.

The association remained positive after accounting for observed user framing and interactional context in covariate-adjusted models. In setting-specific models, scaffolded-support presence was associated with higher constructive-turn counts, with Poisson count ratios ranging from 1.569 to 2.491, and with higher odds that a conversation contained at least one constructive user turn, with logistic odds ratios ranging from 1.437 to 2.140 (Fig. 3b). The same positive direction appeared within both




**Fig. 2 Learning-oriented engagement varies by engagement form, user framing and conversation length.** **a**, Composition of cognitively engaged user turns by passive, active and constructive levels across the six task settings. **b**, Cognitive and constructive user-turn rates in conversations with versus without explicit learning-oriented framing; annotations show the framed-minus-unframed difference in percentage points. **c**, Share of conversations containing at least one constructive user turn by conversation-length bucket. Panel a uses cognitively engaged user turns as the denominator; panel b reports percentages of user turns; panel c reports conversation-level shares. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.

explicitly learning-oriented and non-explicitly framed conversations (Fig. 3c), indicating that the support–engagement association was not limited to interactions users announced as learning episodes. A rate sensitivity treating user turns as the exposure preserved positive scaffolded-support associations in all six settings (rate ratios 1.356–1.703; all  $p < .001$ ; Supplementary Table C11). Across these checks, scaffolded support remained a consistent conversation-level factor of richer constructive participation. This conversation-level signal motivates a finer analysis of the support forms through which scaffolded support was expressed.

## 2.4 Support form differentiates the scaffolding–engagement association

The conversation-level analyses established scaffolded support as a consistent factor of richer constructive participation. We next asked which features of scaffolded support carried this signal. For each scaffolded assistant turn, the annotation scheme recorded two parallel descriptors: a support-intent label indicating whether the turn primarily targeted metacognitive, cognitive or affective support (I1–I3), and one or more support-form labels indicating how the support was delivered—feedback, hinting, instructing, explaining, modelling or questioning (M1–M6). This two-axis scheme



**Fig. 3 Scaffolded support marks richer constructive participation across task settings.** **a**, Turn-weighted constructive user-turn ratios in conversations containing scaffolded support versus reference conversations without scaffolded support. **b**, Covariate-adjusted within-setting associations between scaffolded-support presence and constructive participation, reported as Poisson count ratios and logistic odds ratios. **c**, Poisson count ratios for scaffolded-support presence stratified by explicit versus non-explicit learning-oriented framing. **d**, Difference in the number of turns after the first assistant response between scaffolded and reference conversations. Numbers indicate scaffolded-minus-reference differences, except panels b and c, which report model estimates. Error bars in panels b and c show model-based 95% confidence intervals; bootstrap confidence intervals for the descriptive contrasts in panels a and d are reported in Supplementary Table C10. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.

follows scaffolding accounts that distinguish the purpose of assistance from the means through which assistance is provided [27–30]. Support intent offered limited contrast because scaffolded turns were predominantly cognitive in purpose (Supplementary Table C18). We therefore focused on the form axis, comparing scaffolded conversations containing each support form with scaffolded conversations without that form and reporting turn-weighted constructive user-turn ratio differences (Fig. 4a,b).


Support form sharply differentiated the scaffolded-support association. Feedback-like support (M1) showed the largest constructive-engagement contrast in both explicitly framed and non-explicitly framed conversations, with the association strongest under explicit learning-oriented framing (+14.7 versus +6.7 percentage points; framing difference, +8.0 points; Fig. 4a). Explaining (M4) showed the same direction (+6.7 versus +3.0 points; framing difference, +3.6 points), and hinting (M2) was positive but smaller. Instructing (M3) and questioning (M6), by contrast, were lower or negative, especially in explicitly framed conversations. Task-stratified estimates gave a complementary view (Fig. 4b): feedback-like support (M1) and explaining (M4) were



**Table 2 Summary of key support–engagement associations across the three corpora.** Constructive-ratio contrasts compare conversations with versus without scaffolded support using turn-weighted constructive user-turn ratios. Adjusted models report the association of scaffolded-support presence with constructive-turn count (Poisson count ratio) and with having at least one constructive turn (logit odds ratio). Adjacent-turn lift is the difference in the probability that the next user turn is constructive after scaffolded versus non-scaffolded reference assistant turns.

| Setting           | Constructive ratio<br>Has Scaf. / No Scaf. | Difference<br>(pp) | Depth diff.<br>(turns) | Poisson<br>count ratio | Logit<br>OR | Adjacent lift<br>(pp) |
|-------------------|--------------------------------------------|--------------------|------------------------|------------------------|-------------|-----------------------|
| WildChat coding   | 0.101 / 0.056                              | +4.6               | +2.23                  | 1.852                  | 1.761       | +2.68                 |
| LMSYS coding      | 0.065 / 0.040                              | +2.5               | +1.42                  | 1.622                  | 1.544       | +2.01                 |
| ShareChat coding  | 0.163 / 0.090                              | +7.3               | +3.36                  | 1.893                  | 2.140       | +6.21                 |
| WildChat writing  | 0.027 / 0.017                              | +1.0               | +3.13                  | 1.569                  | 1.437       | +1.13                 |
| LMSYS writing     | 0.022 / 0.011                              | +1.1               | +2.10                  | 1.985                  | 1.764       | +0.27                 |
| ShareChat writing | 0.066 / 0.034                              | +3.2               | +3.38                  | 2.491                  | 1.757       | +3.35                 |

positive in both coding- and writing-oriented task ecologies, whereas instructing (M3), modelling (M5) and questioning (M6) varied more strongly by task.



Support-form labels are non-exclusive; q values are Benjamini-Hochberg FDR-adjusted.

**Fig. 4 Support forms differentiate scaffolded interaction.** **a**, Constructive-engagement differences associated with each support form among scaffolded conversations, stratified by intentional versus unintentional user framing. **b**, The same support-form contrasts stratified by coding- versus writing-oriented task ecology. Bars are turn-weighted constructive user-turn ratio differences, comparing scaffolded conversations containing a support form with scaffolded conversations without that form; error bars show 95% confidence intervals, and bracket annotations show between-stratum differences and Benjamini–Hochberg FDR-adjusted q values within each panel family. Support-form labels correspond to M1–M6 and are not mutually exclusive.

The form-level pattern is distinctive because these support moves occurred inside general-purpose task assistance rather than within a designed lesson. In formal instruction, feedback, explanation and direct guidance are usually embedded in planned sequences with explicit learning goals. In everyday LLM use, the same forms are interleaved with task completion, so an assistant response can either close the problem by supplying the next step or leave part of the problem available for the user

to inspect, revise or test. Feedback-like support (M1) was strongly aligned with constructive engagement because it takes the user’s prior attempt as the object of the response, making a mismatch, partial success or error available for evaluation, repair or extension [34–36]. Explaining (M4) similarly opens the assistant’s answer by exposing mechanisms or rationales that users can test, adapt or connect to their own understanding [37, 38]. Hinting (M2) can leave solution work open, but its association was smaller. Instructing (M3), modelling (M5) and questioning (M6) can be useful for task progress, examples or clarification; in these logs, however, they were less consistently aligned with constructive elaboration. This is consistent with assistance-dilemma and scaffolding accounts in which highly executable steps, models to imitate or information-management moves can reduce the need for learners to generate, evaluate or revise the next idea [22, 39, 40]. Thus, in everyday human–LLM interaction, support form is informative not only as a pedagogical category but as an interactional position: it indicates whether assistance turns the answer into material for the user’s next cognitive move or completes the deliverable with little further reasoning visible.


The support-form supply profile adds task context to these associations (Supplementary Fig. D3). Explaining (M4) dominated scaffolded coding turns across the three corpora, whereas writing-oriented support drew more evenly on instructing (M3), explaining (M4) and questioning (M6). This support ecology clarifies the interactional environment in which the form-level associations arose: coding and writing differed both in which support forms aligned with constructive engagement and in which scaffolding repertoires users typically encountered.

Together, these analyses refine the conversation-level support result. Scaffolded support marked richer constructive participation, and this association was carried most clearly by particular support forms within task-specific scaffolding repertoires. The form-level result makes the temporal question sharper: if support matters partly through how it responds to the user’s ongoing work, then its adjacent-turn association with constructive engagement should depend on the user’s immediately preceding state. We therefore next examine support–engagement coupling at the adjacent-turn scale.

## 2.5 Support–engagement coupling is state- and form-dependent

The preceding analyses established two things: scaffolded conversations contained richer constructive participation, and this association was carried unevenly by support form. These results identify where assistant support is associated with constructive engagement, but not the interactional scale at which the association appears. We therefore turned to adjacent-turn ordering, analysing how the user’s current state, the assistant’s support form and the immediately following user turn were linked.

At this adjacent-turn scale, scaffolded assistant turns were more likely than reference assistant turns to be followed by constructive engagement in all six task settings (Fig. 5a). The next-turn constructive lift was positive in every setting, ranging from +0.27 to +6.21 percentage points. Five of the six adjacent-turn lifts were statistically distinguishable from zero in conversation-cluster bootstrap tests; LMSYS Chat writing was positive but weaker (+0.27 percentage points, 95% CI -0.01 to +0.56,  $p = .061$ ;



**Fig. 5 Support–engagement coupling is state- and form-dependent across task settings.** **a**, Next-turn constructive lift after scaffolded versus reference assistant turns. **b**, Probability of a constructive next user turn by prior user state and assistant support condition; annotations show scaffolded-minus-reference differences. **c**, Probability that the next assistant turn contains scaffolded support by prior user state. **d**, Pooled model/source fixed-effect odds ratios for focal support forms within prior user states. M1 feedback and M4 explaining are shown because Section 2.4 identified them as focal positive conversation-level forms; full M1–M6 interaction estimates and representative support-form cards are reported in Supplementary Table C17 and Supplementary Fig. D4. Error bars show 95% confidence intervals where displayed. Panel d reports adjacent-turn logistic models with conversation-cluster robust standard errors. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.

Supplementary Table C10). The temporal signal was therefore consistently positive across settings, but heterogeneous in magnitude.

This adjacent-turn coupling was organized by the user’s immediately preceding state. State-conditioned contrasts were generally positive but varied across prior constructive, active and passive turns (Fig. 5b). The user-to-assistant direction showed the same contingency: constructive and active user turns were more often followed by scaffolded support than passive turns, especially in coding-oriented conversations (Fig. 5c). These two directions suggest a coupled sequence rather than a one-way support pattern: engaged user states more often preceded scaffolded support, and scaffolded support was then more often followed by further constructive participation.

We then asked whether the form of support shaped this adjacent-turn coupling once prior user state and adjacent-turn context were modelled jointly. We fitted integrated logistic regressions predicting whether the next user turn was constructive,

combining prior user state, assistant-scaffolding descriptors, task and framing variables, turn position and model/source fixed effects. The scaffolding block improved fit beyond user state and context (full scaffolding block LR  $\chi^2_7 = 1339.6$ ,  $p < .001$ ; Supplementary Table C16), and prior-state  $\times$  support-form interactions improved fit further (pooled model/source fixed effects: LR  $\chi^2_{18} = 297.3$ ,  $p < .001$ ; Supplementary Table C17). These tests show that the immediate support–engagement association was not reducible to the user’s preceding state alone; it depended jointly on what the user had just made visible and how the assistant supported the next step.

Figure 5d highlights feedback-like support and explaining, the two support forms that carried the clearest positive signal in the conversation-level analysis; complete M1–M6 interaction estimates and paraphrased support-form cards are reported in the Appendix (Supplementary Table C17; Supplementary Fig. D4). Explaining provided the more stable adjacent-turn signal, predicting constructive follow-up after prior constructive, active and passive user turns (ORs 1.72, 2.12 and 2.88, respectively; all  $p < .001$ ). Feedback showed a different immediate role: its positive next-turn association was concentrated after prior passive turns, rather than after turns in which users were already active or constructive. This divergence clarifies the conversation-level result: feedback marked exchanges in which user attempts were available for evaluation and repair, whereas explanation more consistently provided material for the user’s immediate next move. Support form therefore mattered through its position in the unfolding user–assistant sequence, not as a context-free pedagogical category.

Taken together, the temporal analyses locate support–engagement coupling in adjacent-turn sequences. Constructive participation was most clearly organized where the user’s visible state, the assistant’s support condition and the form of support were considered together. Learning-oriented engagement in everyday human–LLM interaction is therefore assembled in situated next-turn sequences in which the user’s visible work and the assistant’s support form jointly shape whether constructive participation continues.

### 3 Discussion

#### Everyday human–LLM interaction as an informal learning ecology

By tracing engagement in large-scale public logs, this study shows that everyday human–LLM interaction can be analysed as an empirical ecology of informal learning. These conversations are organized around users’ own tasks, yet they contain behavioural traces of learning-oriented work. Section 2.1 establishes the depth structure of this participation: cognitive engagement appeared in nearly one third of user turns, whereas constructive engagement appeared in about five percent. This distinction matters because broad uptake and constructive sense-making mark different relationships to AI output. Uptake shows that users incorporate, apply or extend an assistant response in the task at hand; constructive engagement marks the more selective moments in which users make their reasoning visible by articulating constraints, testing implications, revising assumptions or extending explanations. In this

way, informal learning becomes observable not as a declared intention, but as a pattern of participation within ordinary interaction.

The contextual results in Section 2.2 show how this participation was organized. Explicit learning-oriented framing made constructive engagement more common, but task-oriented exchanges also contained constructive turns. Task ecology shaped which forms of sense-making became visible: coding conversations often exposed failure, constraints and tests, whereas writing conversations made revision, adaptation and evaluative choice visible when users brought those concerns into the exchange. Sustained conversations added temporal opportunity, allowing earlier responses to become points of return, comparison and revision. Together, these findings define everyday human–LLM interaction as an informal learning ecology: user orientation, task affordances and interaction depth shape when ordinary problem solving becomes an occasion to build, test and extend understanding.

### **Scaffolding as situated support**

The assistant-side findings show how support enters this ecology. As shown in Section 2.3, conversations containing scaffolded support had richer constructive participation across all six task settings, and this association remained positive after adjustment for observed conversation-level context. This pattern is most informative when scaffolded support is read as a relation between the assistant response and the user’s next possible action. In everyday task-oriented LLM use, an answer can do more than provide a usable deliverable: it can expose a rationale, error, mechanism, constraint or partial solution that remains available for inspection. This view extends scaffolding theory, which emphasizes contingent assistance and the gradual transfer of responsibility [27, 28, 40], into general-purpose AI interaction, where task completion and opportunities for sense-making are interleaved within the same exchange.

The support-form results specify what this situated support looked like in the logs. As shown in Section 2.4, feedback-like support and explanation were most consistently aligned with constructive engagement. Feedback makes a prior attempt evaluable: it turns a user’s partial solution, diagnosis or draft into something that can be repaired, refined or extended [34, 35]. Explanation makes reasons and mechanisms inspectable, giving users material for comparison, self-explanation and transfer [37, 38]. Other forms, including instruction, modelling and questioning, supported task progress, examples and clarification, with more task-contingent alignment to visible constructive elaboration. The form-level contrasts point to the warrant that support attaches to generated content: some forms organize action, whereas others make visible the grounds on which an answer, attempt or mechanism can be used. For general-purpose AI assistance, this shifts scaffolding from the surface form of a response to the epistemic status of the output it produces: fluent answers become learning-relevant when they are also judgeable answers, with grounds, limits or failure points that can be inspected and adapted within the task.

## Learning continuity across formal and everyday contexts

The temporal findings clarify what everyday AI assistance can add to a broader learning system. As shown in Section 2.5, support–engagement coupling was most interpretable in adjacent user–assistant–user sequences, where the user’s preceding state and the assistant’s support form shaped whether constructive participation continued. This scale gives everyday AI use a role that differs from formal instruction. Formal learning environments organize progression through curricula, sequencing, assessment, shared standards and opportunities to abstract beyond the immediate task. Everyday AI assistance organizes re-entry: it brings explanations, feedback and prior reasoning back into the moments where people apply knowledge under practical constraints.

This distinction matters because much learning depends on whether knowledge can be reactivated when it is needed. A concept explained in a course, tutorial or earlier exchange becomes educationally consequential when it can be recognized again in a new problem, adapted to a different constraint or used to diagnose a breakdown. Human–LLM exchanges provide occasions for that re-entry. A user encounters a concrete difficulty; the assistant supplies a rationale, comparison or correction; the user can then relate that material to the task at hand. The value of such episodes lies in making knowledge usable within action, where understanding is tested against the demands of a real problem.

This complementarity suggests a different role for personal AI assistants in education. Schools, courses and curricula provide progression, accountability and shared criteria. AI assistants can provide continuity across contexts: carrying explanations from one task to another, preserving traces of prior attempts, helping users notice recurring uncertainties and supporting consolidation after repeated encounters with a concept. In this role, lifelong AI assistance can function as a connective layer between structured learning and everyday practice, giving memory and personalization to the informal learning that emerges while people work.

## AI-mediated work and the social organization of expertise

The broader implication is that AI-mediated work may reshape the pathways through which expertise is formed. Much professional learning develops through participation in the practices of a domain: proposing first attempts, encountering error, receiving feedback, revising choices, explaining decisions and seeing consequences unfold [41, 42]. These activities are often treated as parts of production, yet they also carry an apprenticeship function. The results provide an empirical entry point into this process in AI-mediated settings: learning-oriented participation was measurable, selective and structured by the conditions of interaction. As LLMs become embedded in writing, coding, search and problem solving, the formation of expertise becomes tied to how these environments organize participation in the practices through which judgment is built.

AI also changes the repertoire of capacities through which people enter expert practice. When systems can generate fluent drafts, code, explanations and recommendations, expertise may depend increasingly on formulating good problems, recognizing

relevant principles, interrogating generated solutions, judging evidence, revising under constraints and remaining accountable for consequences. These capacities extend beyond generic AI literacy; they are forms of situated judgment that develop through repeated participation in domain work. Schools, workplaces and platforms will therefore shape future expertise through the routines they build around AI: what novices practise, where feedback enters, how responsibility is assigned and how mistakes become occasions for development.

A science of everyday AI learning should therefore study cognitive apprenticeship as a distributed property of sociotechnical systems. The outcomes of interest extend beyond immediate productivity or single-session engagement to the formation of future practitioners: whether repeated AI-mediated routines cultivate judgment, autonomy and domain competence over time. The social stakes are distributional. Intelligent assistance can broaden access to expert-like feedback, explanation and revision, but the developmental value of that access will depend on institutional design: who participates in diagnosis, who practises revision, who observes consequences and who is asked to justify decisions. The long-term educational significance of AI will depend on how societies organize these pathways into expertise, and on whether AI-mediated work sustains broad participation in the practices through which judgment, autonomy and domain competence are formed.

## Limitations

Several boundaries define the interpretation of these findings. The analyses use public conversational logs and characterize population-level interactional patterns rather than longitudinal learner histories. They do not observe users' prior knowledge, motivation, goals outside the conversation or later learning outcomes. The engagement labels capture behavioural manifestations of learning-oriented participation, not direct evidence of retention, transfer or skill development. Assistant support was also observed rather than assigned, so the associations reported here should be interpreted as structured patterns in naturalistic interaction rather than causal effects of scaffolding. Finally, the analysis focuses on English-language coding and writing conversations from three public corpora; other domains, languages, interfaces and institutional settings may organize different opportunities for engagement and support. These boundaries motivate the next stage of evidence: connecting the behavioural signatures identified here to independent measures of learning, autonomy and transfer in controlled and longitudinal settings.

## Conclusion

As LLMs become infrastructure for everyday work and study, the central educational question extends beyond how individuals learn from intelligent systems to how learning opportunities are organized in AI-mediated societies. This study shows that ordinary human–LLM interaction already contains measurable behavioural signatures of informal learning, and that deeper constructive engagement emerges selectively through task ecologies, support forms and short interactional sequences. The long-term significance of AI may therefore depend not only on the knowledge or productivity it

delivers, but also on how societies organize learning opportunities: who encounters them, how often they arise and whether institutions sustain them over time.

## 4 Methods

### 4.1 Study design and analytic scope

This study analyses naturalistic public human–LLM conversations as observational digital traces of everyday work and study, following the logic of large-scale digital-trace and computational social-science analyses [43, 44]. We use three public conversational corpora: WildChat-4.8M [23, 45] (public non-toxic release, reported below as WildChat), LMSYS Chat-1M [24] (reported below as LMSYS Chat) and ShareChat [25] (strict-English subset, reported below as ShareChat). Within these corpora, we focus on coding and writing, two prominent AI-mediated knowledge-work ecologies in which users can delegate task completion to the assistant or remain actively involved in diagnosis, evaluation and revision. Usage and workplace studies identify writing and software-development tasks as high-salience domains of everyday LLM use and AI-mediated knowledge work [1–3, 26]. The two domains also differ in how they make learning-oriented participation visible: coding often externalizes uncertainty through errors, constraints, execution failures and iterative testing [8, 31], whereas writing often involves drafting, revision and evaluative choice whose rationale may remain implicit unless users articulate it [32, 33].

The design is descriptive and associational. We treat conversations as naturally occurring records of user–assistant collaboration and analyse them at three scales: conversations, turns and adjacent user–assistant–user sequences. User turns provide engagement-depth measures, assistant turns provide support measures and adjacent sequences provide local ordering. The analyses estimate where learning-oriented engagement is visible, where constructive engagement is concentrated and how scaffolded support is associated with further constructive participation across conversation-level, support-form and adjacent-turn scales. Dataset and task ecology are treated as substantive dimensions of the interaction environment. The claims are therefore about observable behavioural signatures and interactional organization in public logs, with causal and durable-learning questions reserved for experimental or longitudinal designs [46].

### 4.2 Corpus construction and task settings

We constructed the analytic corpus from three public human–LLM conversation releases: WildChat-4.8M [23, 45], LMSYS Chat-1M [24] and ShareChat [25]. For each source, we converted the public records into role-ordered conversation records with user and assistant turns. We then applied corpus-specific language and minimum-turn screens to retain conversations within the English-oriented analysis scope and with at least four message turns. Coding- and writing-oriented conversations were identified with LLM-assisted semantic task filters. Retained conversations were then processed with the shared annotation pipeline for conversation-level user framing and turn-level labels for user engagement, scaffolded support, support intent and support form.



This construction process yielded six analytic task settings: WildChat coding, WildChat writing, LMSYS Chat coding, LMSYS Chat writing, ShareChat coding and ShareChat writing. The final corpus contains 128,569 conversations and 981,470 total turns, including 491,685 user turns and 489,785 assistant turns. WildChat and LMSYS Chat contribute the largest public interaction logs, whereas ShareChat contributes a smaller complementary sample.

Supplementary Table A1 reports the corpus-specific raw-to-final filtering ledger, including source records, eligibility screens, task-filtering steps and final analytic samples. Supplementary Table A2 records retained provenance artifacts and script-hash prefixes. Supplementary Table A3 summarizes validation and audit evidence for task, language, user-framing and turn-level labels.

### 4.3 Operational constructs

#### 4.3.1 Task ecology, user framing and interaction depth

Each conversation is represented as an ordered sequence of user and assistant turns together with available metadata. The main contextual variables are data source, task ecology, user framing, topic and language. Task ecology and user framing are analytic variables generated by the corpus-construction and annotation pipeline; topic, language and model/source metadata are retained where available. User framing is a conversation-level LLM-assisted label that captures whether the opening request or early exchange positions the interaction as an occasion for understanding, exploration or learning, rather than only for completing the immediate task.

Interaction depth summarizes how much exchange unfolded after the initial request. We use turn counts and turn ordering as descriptive measures of opportunity for follow-up, repair and elaboration, rather than as manipulated exposures. Post-answer depth is measured as the number of turns after the first assistant response. When analyses use whether a conversation contains at least one constructive turn, interaction depth is interpreted as a visibility condition: longer exchanges provide more occasions for constructive participation to appear. The corresponding rate-based sensitivity checks are described with the statistical analyses.

#### 4.3.2 User engagement

The unit of annotation is the conversational turn. User turns are labelled with a depth hierarchy for uptake of assistant output, drawing on engagement theory and the ICAP distinction between passive, active and constructive participation [20, 21]. Passive engagement (P) captures reception, acknowledgement or minimal continuation without substantial transformation. Active engagement (A) captures task-relevant use of, or follow-up on, supplied information, such as applying an answer, asking for a reformatted response or requesting the next step. Constructive engagement (C) captures articulated reasoning with generated output, such as testing an implication, revising a hypothesis or assumption, articulating a constraint, transferring a mechanism or extending an explanation. A user turn labelled P, A or C contributes to the overall cognitive-engagement measure. Explicit affective expression is labelled in parallel when a user turn expresses affect in text; this label characterizes the affective

profile of the exchange alongside the cognitive-depth hierarchy. Operational definitions, plain-language examples and validation metrics are provided in Supplementary Tables A4–C5.

Constructive engagement is the focal user-side measure in the main analyses because it identifies where generated output becomes material for the user’s own reasoning, testing and revision. We represent it as a turn-level C indicator, a conversation-level count, an indicator for at least one constructive user turn and a constructive-turn ratio. The mapping from these measures to the descriptive, conversation-level, support-form and adjacent-turn analyses is summarized in Supplementary Tables C6–C9.

### 4.3.3 Assistant support

Assistant turns are labelled by their role in the user’s ongoing work. The top-level support distinction separates reference responses from scaffolded support. Reference responses primarily provide the requested information, deliverable or stand-alone answer. Scaffolded support captures assistant behaviour that structures, regulates, diagnoses or redirects the user’s process while preserving a role for the user’s own reasoning. This definition follows the scaffolding tradition, especially contingency, transfer of responsibility and potential fading [27, 28, 30, 40].

Scaffolded support is the focal assistant-side measure in the main analyses. At the turn level, assistant turns are represented as scaffolded support or reference response. At the conversation level, we use an indicator for whether a conversation contains at least one scaffolded assistant turn. Scaffolded turns are further characterized by support-intent labels (I1 metacognitive, I2 cognitive and I3 affective support) and non-mutually exclusive support-form labels (M1 feedback, M2 hinting, M3 instructing, M4 explaining, M5 modelling and M6 questioning). The intent labels describe the broad purpose of support, whereas the form labels describe how support is delivered in the turn. Support-form analyses condition on scaffolded turns and compare scaffolded conversations containing each support-form label with scaffolded conversations that do not contain that label. Because support-form labels are multi-label descriptors, these contrasts trace how particular forms of scaffolding pattern with constructive participation within the broader scaffolded-support ecology.

Operational definitions and plain-language examples are provided in Supplementary Tables A4 and C1; the complete de-identified case transcripts are provided in Supplementary Tables C2–C3. Validation metrics for scaffolded support and support-form labels are reported in Supplementary Tables C4 and C5. The mapping from assistant-side measures to conversation-level, support-form and adjacent-turn analyses is summarized in Supplementary Tables C6–C9.

## 4.4 Annotation workflow and validation

The source corpora do not contain native labels for learning-oriented engagement, assistant scaffolding, support intent or support form. We therefore constructed these

labels through an iterative annotation workflow that combined theory-grounded codebook development, independent human review, prompt calibration, parser checks, targeted post-processing and production-scale LLM-assisted labelling.

***Codebook development and domain adaptation.***

The codebook was drafted from learning-science and scaffolding constructs: user-side engagement followed engagement theory and the ICAP distinction between passive, active and constructive participation [20, 21], whereas assistant-side support followed scaffolding accounts emphasizing contingency, structuring of activity and transfer of responsibility [27, 28, 30, 40]. The initial operational codebook was developed in the coding task ecology, where errors, execution failures, repair attempts and iterative testing provided dense boundary cases for distinguishing user uptake depth and assistant support. Two trained author-reviewers independently labelled overlapping coding validation cases across three rounds, reviewed disagreements and produced consensus or adjudicated human-confirmed labels. These coding calibration rounds covered 658 task turns, and refined recurring boundaries including passive versus active uptake, active versus constructive engagement, direct answer delivery versus scaffolded support, explicit affective expression and co-occurring support-form labels.

After the coding codebook stabilized, we adapted the same construct categories to writing. The adaptation preserved the passive-active-constructive user-engagement hierarchy, the reference-response versus scaffolded-support distinction, and the support-intent and support-form label sets, but added writing-specific decision rules and few-shot examples for drafting, rewriting, tone adjustment, evaluative choice and rationale articulation. We then conducted a writing-specific validation pass on 210 reviewed conversations to verify that the adapted codebook could be applied to writing exchanges. These review rounds anchored the final operational boundaries before production-scale annotation. The final operational definitions are reported in Supplementary Table A4; paraphrased label examples and two complete de-identified case transcripts are provided in Supplementary Tables C1- C3. Human-human agreement metrics for this codebook-operability layer are reported in Supplementary Table C4, and Supplementary Section A.4 describes the review-round sizes and adjudication workflow.

***Prompt calibration and production labelling.***

LLM-assisted annotation was used as the scalable production procedure after the domain-adapted codebook, prompt templates and structured output schemas had been calibrated against human-reviewed examples. Prompt outputs were first tested on a label-enriched pilot calibration set of 400 conversations, stratified by task ecology with 200 coding and 200 writing conversations. The pilot was assembled from preliminary LLM labels across all annotation prompt families, with quotas requiring at least 30 candidate turn-level instances for every target label, and was supplemented with randomly sampled background conversations without target positive labels. Because sampling was conducted at the conversation level but label quotas were defined at the turn level, the 400 conversations yielded a substantially larger set of reviewed turns for testing label boundaries.

Prompt templates and schemas were revised against observed disagreement patterns, schema failures and boundary errors. After these revisions, a balanced 200-turn human validation check, with 100 coding and 100 writing examples, was used to test the refreshed user-turn label boundaries and identify remaining recurring false-positive patterns. The resulting updates were incorporated into the frozen codebook–schema–parser–post-processing workflow used for final analysis labels. WildChat labels reflect the targeted update pass applied after validation; LMSYS Chat and ShareChat were then processed with the same refreshed workflow. Production labelling used the same construct definitions, prompt versions, schema checks and post-processing logic across corpora.

### ***Validation design.***

Validation evidence was mapped to the stage of the workflow it evaluated. The first layer assessed codebook operability through human–human agreement. This layer tested whether trained reviewers could apply the finalized construct definitions after coding codebook development and writing-domain adaptation. The coding codebook–operability check combined three validation rounds covering 658 task turns. The writing-domain validation used 210 reviewed conversations after the coding-derived codebook had been adapted with writing-specific decision rules and few-shot examples. Human–human agreement for this layer is reported in Supplementary Table C4; Supplementary Section A.4 describes the review rounds, adjudication process and domain-adaptation workflow.

The second layer assessed production-label quality after the frozen codebook–schema–parser–post-processing workflow had been applied. These production-label audits compared final LLM-assisted labels with human-confirmed labels and were stratified across the six corpus-by-task settings. The user-framing audit reviewed 450 first user turns, with 75 cases per setting: 50 production-positive and 25 production-negative cases. The user-engagement audit reviewed 600 user turns, with 100 cases per setting and quota enrichment by production label: 50 constructive, 20 active, 15 passive and 15 no-engagement cases. The assistant-side audit reviewed 180 assistant turns, with 30 cases per setting: 24 production-scaffolded and 6 reference-response cases. Support-form positives were enriched within the scaffolded cases, and support-form labels were treated as non-exclusive. Production-label audit metrics are reported in Supplementary Table C5.

The third layer consisted of targeted post-production checks for recurring boundary cases identified during validation and prompt calibration. A final user-turn check reviewed 100 coding and 100 writing examples after the prompt and boundary-rule updates. After the whole process, the WildChat analytic sample contained 31,878 coding conversations and 39,534 writing conversations. LMSYS Chat and ShareChat were processed with the same refreshed workflow. Supplementary Section A.4 reports the post-update checks and sampling details, and Supplementary Table A3 summarizes validation and audit evidence for task ecology, language filtering, user framing and turn-level labels.

### *Agreement metrics.*

Agreement is summarized with per-label F1, Matthews correlation coefficient and Gwet’s AC1 because the labels are sparse, prevalence-imbalanced and sometimes multi-label [47, 48]. We report agreement separately for the two main validation layers. Supplementary Table C4 reports human–human codebook agreement, which evaluates whether trained reviewers could apply the operational definitions consistently during codebook calibration and writing-domain adaptation. Supplementary Table C5 reports production-label agreement, which evaluates the final LLM-assisted labels after schema checks, parser checks and post-processing.

The consolidated production-label audits support the labels used for the main claims. Intentional user framing reached F1=0.857; constructive-versus-non-constructive user engagement reached F1=0.922; broad scaffolded support reached F1=0.912; and support-form labels ranged from F1=0.691 to 0.938. The audit samples intentionally enriched positive and boundary-relevant cases to stress-test classification boundaries; prevalence estimates are computed from the full production labels reported in the Results 2.

## 4.5 Statistical analyses

### *Descriptive and contextual estimates.*

Descriptive analyses estimated engagement prevalence, scaffolding prevalence, passive/active/constructive composition, explicit affective-expression rates and interaction-depth summaries within each task setting. Contextual analyses compared observed user framing, task ecology and interaction-depth groups. For analyses using whether a conversation contained at least one constructive user turn, we paired the event indicator with a rate-based sensitivity model to account for the number of user turns available for observing constructive participation. The sensitivity model was a grouped-binomial model in which the outcome was the constructive user-turn count out of total user turns, with user framing, task ecology, length bucket and dataset fixed effects as predictors.

### *Conversation-level support associations.*

Conversation-level support analyses compared conversations containing at least one scaffolded assistant turn with reference conversations without scaffolded support. Constructive-ratio contrasts used turn-weighted ratios within each comparison group, defined as total constructive user turns divided by total user turns, with uncertainty estimated from 1,000 conversation-level bootstrap resamples. Covariate-adjusted models reported in Section 2.3 used setting-specific Poisson generalized linear models for raw constructive-turn counts and companion logistic models for whether a conversation contained at least one constructive user turn [49, 50]. The primary Poisson model included conversation length as a covariate; offset-rate and quasi-Poisson sensitivities are reported in the Supplement. The resulting estimates are interpreted as covariate-adjusted associations in naturalistic conversation logs.

### ***Support-form analyses.***

Support-form analyses conditioned on scaffolded conversations. For each support-form label, we compared scaffolded conversations containing that label with scaffolded conversations that did not contain that label, and reported turn-weighted constructive user-turn ratio differences. Because support-form labels are multi-label descriptors, these contrasts trace how particular forms of scaffolding pattern with constructive participation within the broader scaffolded-support ecology. Bracket tests in Fig. 4 compare displayed strata for each support form; q values use Benjamini–Hochberg false-discovery-rate adjustment within each displayed family of six support-form comparisons [51].

### ***Adjacent-turn analyses.***

Temporal analyses focused on adjacent-turn ordering. We compared the probability that the next user turn was constructive after scaffolded versus reference assistant turns, estimated the same contrast within prior user states and estimated the probability that the next assistant turn was scaffolded after different prior user states. Adjacent-turn contrasts used 1,000 conversation-cluster bootstrap resamples because multiple turn pairs can come from the same conversation.

Integrated adjacent-turn regressions combined prior user-state features, assistant-scaffolding features, support-form descriptors and adjacent-turn context in one logistic model predicting whether the next user turn was constructive. These models included task and framing variables, turn position and dataset fixed effects; a sensitivity replaced dataset fixed effects with recoverable model/source-family fixed effects where available. Standard errors were clustered by conversation. We report broad scaffolded-support models separately from support-form decompositions because support-form labels are descriptors nested within scaffolded support. Nested likelihood-ratio block tests first add broad scaffolded-support presence and then test whether support-form descriptors add signal within scaffolded support. Supplementary Tables C6–C9 and C8 give the complete estimand map, covariate sets, model specifications and sensitivity checks.

## **4.6 LLM use and reporting boundaries**

Large language models were used in four roles. First, LLMs were the systems under investigation: the data are public human–LLM conversations from WildChat, LMSYS Chat and ShareChat. Available model/source identifiers were retained as observational metadata for descriptive and robustness analyses, not as randomized model comparisons.

Second, LLMs were used as measurement tools. GPT-4o-family Azure OpenAI deployments were used for coding-versus-writing task filtering and English-oriented filtering. Azure OpenAI GPT-5.1-family deployments were used for production annotation: `gpt-5.1` for the main user-turn and assistant-turn labels, and `gpt-5.1-chat` for final writing-oriented annotation. These calls assigned user framing, engagement, scaffolding, support intent and support-form labels under fixed prompts, schema/parser checks, retry logic, post-processing rules and human validation.

Third, LLMs were used during annotation development to test prompts and inspect boundary cases against human-reviewed examples; pilot outputs were not used as production labels or prevalence estimates. Fourth, LLM-based tools assisted with code drafting, debugging, workflow organization and language editing, but not as substitutes for human validation, audit-label adjudication, statistical decisions or result interpretation.

All analyses are aggregate turn- or conversation-level analyses. Model/source labels are observational descriptors, so model-related comparisons are interpreted as robustness analyses rather than causal effects.

## Data availability

This study analyses public human–LLM conversation corpora released by their original data providers: WildChat, LMSYS Chat-1M and ShareChat. Readers should obtain the raw public conversation data from the original releases and comply with the corresponding dataset licences, terms of use and redistribution restrictions. The manuscript does not redistribute verbatim raw conversation logs, user identifiers or linked user histories. The Supplementary includes two complete de-identified illustrative transcripts derived from retained conversations and lightly edited to preserve turn sequence, task structure and annotation decisions while reducing privacy and licence risk.

Derived data sufficient to verify the reported aggregate figures, tables and statistical summaries are archived on Zenodo at <https://doi.org/10.5281/zenodo.20995946>. The same release is also available in the GitHub repository at <https://github.com/CinderD/Informal-learning-in-everyday-human-LLM-interaction> (tag v0.1.0, commit 31e10cfea930). The release includes the numeric values underlying the main figures, de-identified conversation-level and adjacent-turn analytic labels without raw message text, table source files, confidence intervals, p values, false-discovery-rate q values, model-output summaries, provenance documentation and SHA-256 checksums. Any request for additional derived label files that are not included in the public release is assessed against the original dataset licences and privacy constraints.

## Code availability

Custom code used to export source data and de-identified label tables, compute statistical analyses and generate figures and tables is archived with the derived-data release on Zenodo at <https://doi.org/10.5281/zenodo.20995946> and is also available in the GitHub repository at <https://github.com/CinderD/Informal-learning-in-everyday-human-LLM-interaction> (tag v0.1.0, commit 31e10cfea930). The repository excludes API keys, raw message text and files that cannot be redistributed under the original corpus terms. The archived release also includes annotation prompt templates, output schemas, parser checks, post-processing scripts and production manifests, excluding API keys and raw message text. Reproducing the full annotation pipeline from raw conversations requires user-provided API credentials and access to the raw public corpora from their original sources.

## Ethics statement

The study is a secondary analysis of public, naturalistic human–LLM conversation datasets. The authors did not recruit participants, intervene in conversations or attempt to identify users. Analyses are reported at aggregate turn or conversation level. Operational examples are de-identified and paraphrased or lightly edited; the two complete illustrative transcripts preserve the original turn structure and annotation decisions but are not redistributed verbatim raw logs. Human validation was used to assess annotation quality on reviewed examples and did not involve interaction with the original dataset users.

## Competing interests

H.L. and X.X. are employees of Microsoft Research Asia. The remaining authors declare no competing interests.

## Author contributions

Author contributions are organized using the CRediT taxonomy. The contribution categories for this study are conceptualization, study design, data curation, annotation pipeline development, validation, statistical analysis, figure generation, writing of the original draft, manuscript review and editing, supervision and funding acquisition. Author initials and final role allocation: [to be added after the author list and order are confirmed]. All listed authors will approve the submitted version and will be accountable for their contributions.

## Acknowledgements and funding

Acknowledgements and funding information, including grant numbers where applicable: [to be added]. Funders had no role in study design, data collection, analysis, interpretation, manuscript preparation or publication decisions unless stated otherwise.

## LLM use statement

Large language models were objects of study and were also used as annotation and writing-assistance tools. Production labels were generated with LLM-assisted workflows and checked through parser rules, targeted update scripts and human validation. The authors also used LLM-based tools to assist with code drafting, debugging, workflow organization and language editing. No LLM or AI tool is listed as an author, and the authors take responsibility for the final manuscript, analyses and interpretation. Additional details are provided in Methods.



## Reporting summary and editorial policy checklist

A Nature Portfolio Reporting Summary and the journal-required editorial policy checklist have been prepared as separate submission files. The manuscript and supplementary information report the observational design, sampling and filtering rules, annotation workflow, human validation evidence, statistical model specifications, bootstrap procedures, multiple-comparison handling and data/code availability statements needed to assess reproducibility.

## Figure source data

Source data for the numeric main figures are provided in `source_data/figure2_source_data.csv`, `source_data/figure3_source_data.csv`, `source_data/figure4_source_data.csv` and `source_data/figure5_source_data.csv`. Figure 1 is a conceptual framework figure and has no numeric source data. Supplementary table and regression source files are provided in `tables/` and `outputs/integrated_regression/`.

## References

- [1] Chatterji, A. *et al.* How people use ChatGPT. Tech. Rep. Working Paper 34255, National Bureau of Economic Research (2025).
- [2] Noy, S. & Zhang, W. Experimental evidence on the productivity effects of generative artificial intelligence. *Science* **381**, 187–192 (2023).
- [3] Brynjolfsson, E., Li, D. & Raymond, L. Generative AI at work. *The Quarterly Journal of Economics* **140**, 889–942 (2025).
- [4] Anzai, Y. & Simon, H. A. The theory of learning by doing. *Psychological review* **86**, 124 (1979).
- [5] Kolb, D. A. *Experiential learning: Experience as the source of learning and development* (FT press, 2014).
- [6] Risko, E. F. & Gilbert, S. J. Cognitive offloading. *Trends in cognitive sciences* **20**, 676–688 (2016).
- [7] Gerlich, M. Ai tools in society: Impacts on cognitive offloading and the future of critical thinking. *Societies* **15**, 6 (2025).
- [8] Barke, S., James, M. B. & Polikarpova, N. Grounded copilot: How programmers interact with code-generating models. *Proceedings of the ACM on Programming Languages* **7**, 85–111 (2023).
- [9] Kobiella, C., Flores López, Y. S., Waltenberger, F., Draxler, F. & Schmidt, A. “If the machine is as good as me, then what use am I?”—how the use of ChatGPT

- changes young professionals' perception of productivity and accomplishment. *Proceedings of the CHI Conference on Human Factors in Computing Systems* 1–16 (2024).
- [10] Eraut, M. Informal learning in the workplace. *Studies in Continuing Education* **26**, 247–273 (2004).
  - [11] Marsick, V. J. & Watkins, K. *Informal and incidental learning in the workplace (Routledge Revivals)* (Routledge, 2015).
  - [12] Schugurensky, D. The forms of informal learning: Towards a conceptualization of the field. Tech. Rep. NALL Working Paper 19, Centre for the Study of Education and Work, OISE/UT, Toronto (2000). URL <https://hdl.handle.net/1807/2733>.
  - [13] Kasneci, E. *et al.* ChatGPT for good? On opportunities and challenges of large language models for education. *Learning and Individual Differences* **103**, 102274 (2023).
  - [14] Meyer, J. *et al.* Using LLMs to bring evidence-based feedback into the classroom: AI-generated feedback increases secondary students' text revision, motivation, and positive emotions. *Computers and Education: Artificial Intelligence* **6**, 100199 (2024).
  - [15] Arif, T., Asthana, S. & Collins-Thompson, K. Generation and assessment of multiple-choice questions from video transcripts using large language models. *Proceedings of the Eleventh ACM Conference on Learning @ Scale* 530–534 (2024).
  - [16] Jin, H., Lee, S., Shin, H. & Kim, J. Teach AI how to code: Using large language models as teachable agents for programming education. *Proceedings of the CHI Conference on Human Factors in Computing Systems* 1–28 (2024).
  - [17] Lim, H., Choi, D. & Hong, H. Identify design problems through questioning: Exploring role-playing interactions with large language models to foster design questioning skills. *Companion Publication of the 2024 Conference on Computer-Supported Cooperative Work and Social Computing* 598–602 (2024).
  - [18] Terzimehić, N., Bühler, B. & Kasneci, E. Conversational ai as a catalyst for informal learning: An empirical large-scale study on llm use in everyday learning. *Computers and Education: Artificial Intelligence* 100634 (2026).
  - [19] VanLehn, K. The relative effectiveness of human tutoring, intelligent tutoring systems, and other tutoring systems. *Educational Psychologist* **46**, 197–221 (2011).
  - [20] Fredricks, J. A., Blumenfeld, P. C. & Paris, A. H. School engagement: Potential of the concept, state of the evidence. *Review of Educational Research* **74**, 59–109

- (2004).
- [21] Chi, M. T. & Wylie, R. The ICAP framework: Linking cognitive engagement to active learning outcomes. *Educational Psychologist* **49**, 219–243 (2014).
  - [22] Koedinger, K. R. & Aleven, V. Exploring the assistance dilemma in experiments with cognitive tutors. *Educational Psychology Review* **19**, 239–264 (2007).
  - [23] Zhao, W. *et al.* WildChat: 1M ChatGPT interaction logs in the wild (2024). URL <https://openreview.net/forum?id=Bl8u7ZRlbM>.
  - [24] Zheng, L. *et al.* LMSYS-Chat-1M: A large-scale real-world LLM conversation dataset (2023). [arXiv:2309.11998](https://arxiv.org/abs/2309.11998).
  - [25] Yan, Y., Nguyen, T., Su, B., Lieffers, M. & Le, T. ShareChat: A dataset of chatbot conversations in the wild (2025). [arXiv:2512.17843](https://arxiv.org/abs/2512.17843).
  - [26] Handa, K. *et al.* Which economic tasks are performed with AI? Evidence from millions of Claude conversations (2025). [arXiv:2503.04761](https://arxiv.org/abs/2503.04761).
  - [27] Wood, D., Bruner, J. S. & Ross, G. The role of tutoring in problem solving. *Journal of Child Psychology and Psychiatry* **17**, 89–100 (1976).
  - [28] Van de Pol, J., Volman, M. & Beishuizen, J. Scaffolding in teacher–student interaction: A decade of research. *Educational Psychology Review* **22**, 271–296 (2010).
  - [29] Winkler, R., Hobert, S., Salovaara, A., Söllner, M. & Leimeister, J. M. Sara, the lecturer: Improving learning in online education with a scaffolding-based conversational agent. *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems* 1–14 (2020).
  - [30] Quintana, C. *et al.* A scaffolding design framework for software to support science inquiry. *The Journal of the Learning Sciences* **13**, 337–386 (2004).
  - [31] Keuning, H., Jeuring, J. & Heeren, B. A systematic literature review of automated feedback generation for programming exercises. *ACM Transactions on Computing Education* **19**, 1–43 (2018).
  - [32] Flower, L. & Hayes, J. R. A cognitive process theory of writing. *College Composition and Communication* **32**, 365–387 (1981).
  - [33] Sommers, N. Revision strategies of student writers and experienced adult writers. *College Composition and Communication* **31**, 378–388 (1980).
  - [34] Hattie, J. & Timperley, H. The power of feedback. *Review of Educational Research* **77**, 81–112 (2007).

- [35] Shute, V. J. Focus on formative feedback. *Review of Educational Research* **78**, 153–189 (2008).
- [36] Nicol, D. J. & Macfarlane-Dick, D. Formative assessment and self-regulated learning: A model and seven principles of good feedback practice. *Studies in Higher Education* **31**, 199–218 (2006).
- [37] Chi, M. T. H., Bassok, M., Lewis, M. W., Reimann, P. & Glaser, R. Self-explanations: How students study and use examples in learning to solve problems. *Cognitive Science* **13**, 145–182 (1989).
- [38] Renkl, A. Learning from worked-out examples: A study on individual differences. *Cognitive Science* **21**, 1–29 (1997).
- [39] Aleven, V., Stahl, E., Schworm, S., Fischer, F. & Wallace, R. Help seeking and help design in interactive learning environments. *Review of Educational Research* **73**, 277–320 (2003).
- [40] Reiser, B. J. Scaffolding complex learning: The mechanisms of structuring and problematizing student work. *The Journal of the Learning Sciences* **13**, 273–304 (2004).
- [41] Brown, J. S., Collins, A. & Duguid, P. Situated cognition and the culture of learning. *Educational researcher* **18**, 32–42 (1989).
- [42] Collins, A., Brown, J. S. & Newman, S. E. Cognitive apprenticeship: Teaching the crafts of reading, writing, and mathematics. *Knowing, learning, and instruction* 453–494 (2018).
- [43] Lazer, D. *et al.* Computational social science. *Science* **323**, 721–723 (2009).
- [44] Salganik, M. J. *Bit by Bit: Social Research in the Digital Age* (Princeton University Press, 2017).
- [45] Allen Institute for AI. WildChat-4.8M. <https://huggingface.co/datasets/allenai/WildChat-4.8M> (2025).
- [46] Hernán, M. A. & Robins, J. M. *Causal Inference: What If* (Chapman & Hall/CRC, Boca Raton, 2020).
- [47] Matthews, B. W. Comparison of the predicted and observed secondary structure of T4 phage lysozyme. *Biochimica et Biophysica Acta* **405**, 442–451 (1975).
- [48] Gwet, K. L. Computing inter-rater reliability and its variance in the presence of high agreement. *British Journal of Mathematical and Statistical Psychology* **61**, 29–48 (2008).

- [49] McCullagh, P. & Nelder, J. A. *Generalized Linear Models* 2nd edn, Vol. 37 of *Monographs on Statistics and Applied Probability* (Chapman & Hall, London, 1989).
- [50] Agresti, A. *Categorical Data Analysis* 3rd edn (John Wiley & Sons, Hoboken, NJ, 2013).
- [51] Benjamini, Y. & Hochberg, Y. Controlling the false discovery rate: a practical and powerful approach to multiple testing. *Journal of the Royal statistical society: series B (Methodological)* **57**, 289–300 (1995).

## Supplementary information

Supplementary information is organized to support reproducibility and interpretation. Appendix A records the corpus ledger, annotation ontology, production workflow, validation evidence, denominator rules and model specifications. Appendix B reports robustness checks and interpretive boundaries. Appendices C and D contain the supporting tables and figures referenced from the main text and Appendix A–B.

## Appendix A Reproducibility protocol and construct ledger

### A.1 Corpus provenance and filtering ledger

The main Methods summarize the three public corpora and six task settings. This section provides the auditable selection ledger behind those settings. The analyses use WildChat coding, WildChat writing, LMSYS Chat coding, LMSYS Chat writing, ShareChat coding and ShareChat writing. The analytic corpus contains 128,569 conversations and 981,470 total turns. WildChat contributes 31,878 coding conversations and 39,534 writing conversations; LMSYS Chat contributes 32,114 coding conversations and 21,023 writing conversations; ShareChat contributes 2,481 coding conversations and 1,539 writing conversations.

Corpus construction followed a common sequence: source records were converted to role-ordered user–assistant conversations; conversations were screened for English-oriented content and at least four message turns; coding and writing task ecologies were assigned with LLM-assisted semantic filters; and retained conversations were processed by the shared annotation workflow. ShareChat required an additional strict-English screen after task filtering. Public dataset-card turns, raw CSV message rows and role-normalized analysis turns are not identical counting units, so the ShareChat ledger reports platform-url conversations and role-labelled message rows before deriving user and assistant turns.

**Table A1 Corpus filtering and classification pipeline.** The table records the auditable steps used to construct the six main analytic settings. Task domain and user framing are analysis variables generated by the filtering and annotation pipeline, not native corpus metadata.

| Corpus     | Source records and first screen                                                                                                                                                                                                                                                                                                                                                                      | Language and minimum-turn filter                                                                                                                                                                                                                                                                         | Task classification source                                                                                                                                                                                                                                                                                                | Final analytic sample          | User framing and audit trail                                                                                                                                                                                                                                                                                                                   |
|------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| WildChat   | Public WildChat-4.8M release. Raw conversations are obtained from the original release rather than redistributed with the manuscript. The retained analysis archive records the refreshed task-filter outputs, user-turn relabel/update logs and downstream source-data exports used for this manuscript version.                                                                                    | Retained task-filter outputs have at least four message turns. The task filters were English-oriented semantic filters; source language metadata were retained for analysis but were not used as an additional strict post-hoc exclusion in the main WildChat sample.                                    | LLM-assisted semantic task filters labelled conversations as coding_english or writing_english with score threshold $\geq .70$ after keyword prefiltering. The retained filter outputs contain 31,879 coding conversations and 39,534 writing conversations; all retained records had task-filter scores above threshold. | 31,878 coding; 39,534 writing. | User framing came from the conversation-level learning_intent annotation (INTENTIONAL versus UNINTENTIONAL); not from native metadata. The final WildChat user-turn update pass scanned 31,879 coding and 39,534 writing conversations; one coding conversation with missing learning_intent was omitted from Level 2 and downstream analyses. |
| LMSYS Chat | LMSYS Chat-1M source release: 1,000,000 conversations.                                                                                                                                                                                                                                                                                                                                               | The conversion retained 777,453 English conversations and excluded 222,547 non-target-language conversations. Among retained English conversations, 236,372 had at least four turns.                                                                                                                     | The same LLM-assisted semantic task filters were run on the English pool with the minimum-turn screen applied inside the filter. The coding filter selected 32,114 coding conversations and the writing filter selected 21,023 writing conversations.                                                                     | 32,114 coding; 21,023 writing. | User framing came from the downstream learning_intent annotation. The conversion log also records redaction, moderation and model/source metadata where available; these fields were retained for robustness checks rather than used as task labels.                                                                                           |
| ShareChat  | ShareChat Hugging Face release. The dataset card reports 142,808 public conversations and 660,293 public turns across five platforms. The local conversion starts from the five final-language-filtered CSV files, which contain 1,199,899 message rows and 129,584 unique platform-url conversations; role-normalization retained 1,196,969 rows mapped to user or bot (user=600,222; bot=596,747). | The first screen retained 69,915 conversations with at least four message turns. After task filtering, a strict-English screen retained only conversations with final conversation-level language flag English; conversations flagged non-English or mixed/non-English by this final flag were excluded. | The same LLM-assisted semantic filters selected 3,494 coding and 1,663 writing conversations before the strict-English screen. The strict-English screen retained 2,481 of 3,494 coding conversations and 1,539 of 1,663 writing conversations.                                                                           | 2,481 coding; 1,539 writing.   | User framing came from the downstream learning_intent annotation. The strict-English summary recorded zero malformed JSON files; excluded counts were 1,013 coding and 124 writing conversations.                                                                                                                                              |

*Note.* A “conversation” is a role-ordered user-assistant record after corpus-specific conversion. The minimum-turn screen uses message count, not user-assistant round count. For ShareChat, the public dataset-card “turns” are not the same unit as the local message-row count used in conversion; our pipeline uses one row per role-labelled message and then counts user and assistant turns after role normalization. The term “strict-English” is reserved for the additional ShareChat final-language screen; WildChat and LMSYS Chat used English-oriented source or semantic filters as described above. Human validation was organized by task ecology rather than independently powered within each corpus; Supplementary Table C4 reports the domain-stratified agreement metrics.

Supplementary Table [A1](#) gives the raw-to- final filtering ledger for each corpus. Supplementary Table [A2](#) records the retained provenance artifacts and script-hash prefixes. Supplementary Table [A3](#) separates corpus-construction filters from analysis labels and reports validation or audit evidence for task ecology, user framing, language filtering and turn-level annotation.



**Table A2 Corpus construction provenance artifacts.** The table records the retained non-text provenance artifacts used to audit corpus construction. SHA-256 values are shown as 12-character prefixes for run scripts or manuscript-source scripts retained in the analysis archive. Raw conversation text remains governed by the original dataset releases and is not redistributed with the manuscript source.

| Corpus     | Source release and conversion unit                                                                                                                                                                                             | Selection ledger retained for this manuscript                                                                                                                                                                                                                                                                                          | Script/source-data provenance                                                                                                                                                                         |
|------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| WildChat   | WildChat-4.8M public release; role-ordered user/assistant messages after task filtering.                                                                                                                                       | Public release → English-oriented coding/writing semantic task filters with minimum four message turns → refreshed task-filter outputs (31,879 coding; 39,534 writing) → user-turn relabel/update pass → analytic sample (31,878 coding; 39,534 writing).                                                                              | WildChat relabel/update run script SHA-256 prefix <b>c8eecef55858</b> ; source-data exporter <b>206aaae44791</b> ; camera-ready figure script <b>d9ecf7118f10</b> .                                   |
| LMSYS Chat | LMSYS Chat-1M public release; one conversation record per public chat, converted to role-ordered messages.                                                                                                                     | 1,000,000 source conversations → 777,453 English conversations → 236,372 English conversations with at least four turns → semantic task filters → analytic sample (32,114 coding; 21,023 writing).                                                                                                                                     | Coding filter launcher <b>f54c079fda08</b> ; writing filter launcher <b>be970385e72e</b> ; after-filter annotation/analysis launcher <b>326e11709c41</b> ; source-data exporter <b>206aaae44791</b> . |
| ShareChat  | ShareChat Hugging Face release; public card reports 142,808 conversations and 660,293 public turns. Local conversion uses five final-language-filtered CSVs with one row per message and platform-url as the conversation key. | Five final-language-filtered CSVs (1,199,899 message rows; 129,584 platform-url conversations) → role-normalized user/bot rows (1,196,969 rows) → 69,915 conversations with at least four message turns → semantic task filters (3,494 coding; 1,663 writing) → strict-English screen → analytic sample (2,481 coding; 1,539 writing). | Coding filter launcher <b>f2ff9ca543b8</b> ; writing filter launcher <b>2788818044e7</b> ; after-filter annotation/analysis launcher <b>501e26e5a728</b> ; source-data exporter <b>206aaae44791</b> . |

**Table A3 Filtering and classifier validation audit.** The table separates corpus-construction filters from analysis labels and reports the strongest validation or audit evidence available in the current artifacts. Metrics from the writing-210 validation set are domain-level checks, whereas the user-framing audit deliberately sampled across corpus-by-task settings.

| Variable or filter                                        | How it was assigned                                                                                                                                                                                                                                        | Validation or audit evidence                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Interpretation boundary                                                                                                                                                                                                                                                                      |
|-----------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Task ecology: coding versus writing                       | LLM-assisted semantic task filters selected conversations matching coding-oriented or writing-oriented use after minimum-turn screening. The filters were English-oriented and used a score threshold of 0.70 after keyword prefiltering where applicable. | The provenance ledger records retained counts and selection steps for each corpus (Supplementary Table A2). In the writing-210 validation artifacts, the conversation-topic field showed high human-LLM agreement ( $F1=0.973$ , $MCC=0.872$ , $Gwet\ AC1=0.936$ ; $support=210$ ), but this is a domain-level check rather than a full cross-corpus task-filter audit.                                                                                                                                                                                                                                                                                                                                                                        | Task ecology is treated as an analytic classification of conversation content, not native dataset metadata or randomized assignment.                                                                                                                                                         |
| User framing: intentional versus unintentional            | Conversation-level LLM-assisted <b>learning_intent</b> annotation marked conversations explicitly oriented toward understanding, exploration or learning versus those without explicit learning-oriented framing.                                          | A human-verified positive-oversampled audit reviewed 450 first user turns across all six corpus-by-task settings (75 per setting; 50 production-positive and 25 production-negative). Against the production labels, overall intentional-framing $F1$ was 0.857, $MCC$ was 0.677 and $Gwet\ AC1$ was 0.671; task-stratified $F1$ was 0.924 for coding and 0.779 for writing.                                                                                                                                                                                                                                                                                                                                                                   | User framing is a measured textual signal of explicit learning orientation, not a stable learner trait, motivation measure or causal exposure. Because the audit oversampled production-positive cases, it validates classification boundaries rather than estimating population prevalence. |
| Language filtering                                        | WildChat and LMSYS Chat used English-oriented source or semantic filters. ShareChat received an additional strict-English screen after task filtering, retaining only records with final conversation-level language flag <b>Eng11sh</b> .                 | The ShareChat strict-English audit recorded zero malformed JSON files and explicit exclusions after task filtering: 1,013 coding and 124 writing conversations were removed by the final language screen. In the writing-210 validation artifacts, language showed near-perfect agreement ( $F1=0.998$ , $MCC=0.705$ , $Gwet\ AC1=0.995$ ; $support=210$ ), with $MCC$ lower because nearly all reviewed records were English.                                                                                                                                                                                                                                                                                                                 | The term strict-English is used only for ShareChat’s final post-task language screen. Mixed or non-English ShareChat records were excluded from the main analytic sample.                                                                                                                    |
| Turn-level labels: user engagement and scaffolded support | Production annotation used the final codebook, parser checks and post-processing rules. Constructive engagement and scaffolded support are the primary user-side outcome and assistant-side exposure.                                                      | Domain-stratified human-human codebook validation shows strong agreement for constructive engagement (coding $F1=0.85$ ; writing $F1=0.89$ ) and scaffolded support (coding $F1=0.87$ ; writing $F1=0.97$ ), with full per-label codebook metrics in Supplementary Table C4. Human-confirmed production audits across all six corpus-by-task settings showed acceptable or better reliability for user engagement and assistant labels: the 600-case user audit had constructive-vs-non-constructive $F1=0.922$ , $MCC=0.840$ and $Gwet\ AC1=0.837$ ; the 180-case assistant audit had scaffolded-support $F1=0.912$ , $MCC=0.590$ and $Gwet\ AC1=0.791$ , with support-form $F1$ values ranging from 0.691 to 0.938 (Supplementary Table C5). | These labels support the main association and temporal analyses. The production audits are positive-oversampled and validate label boundaries rather than prevalence.                                                                                                                        |

For each retained conversation, the pipeline stored the ordered turn sequence, role labels, available topic and language metadata, learning-intent indicators and model/source metadata where available. Model metadata are used only for robustness checks because model assignment in naturalistic public conversations is observational and may reflect time, user population, task mix and logging context.

## **A.2 Annotation ontology and operational examples**

The codebook translates learning-science constructs into role-specific behavioural signatures observable in conversation logs. User turns and assistant turns are labelled with different constructs because they play different roles in the interaction. Supplementary Table A4 gives the operational definitions and boundary rules. Supplementary Table C1 gives paraphrased examples that illustrate the difference between engagement levels and support forms without reproducing raw user text. Supplementary Tables C2 and C3 provide complete de-identified illustrative transcripts, one coding-oriented and one writing-oriented, derived from retained conversations and lightly edited while preserving turn sequence, task structure and annotation rationales.

**Table A4 Annotation codebook definitions.** The table defines the labels used to translate learning-science constructs into observable turn-level signals. These labels are behavioural annotations of conversation text; they are not measures of durable learning, learner identity or instructional efficacy.

| Family            | Label            | Definition                                                                                                                                            | Boundary rule                                                                                        |
|-------------------|------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------|
| User engagement   | Passive (P)      | Reception, acknowledgement or minimal continuation without visible transformation of the assistant response.                                          | Includes short confirmations or requests to continue; excludes task-relevant follow-up or reasoning. |
| User engagement   | Active (A)       | Task-relevant action, follow-up, application or manipulation of supplied information.                                                                 | The user does something with the response, but does not test, revise or extend an idea.              |
| User engagement   | Constructive (C) | Highest-depth observable user engagement: elaborating, testing, revising, transferring or extending an idea or explanation.                           | Requires visible user reasoning or transformation; not inferred from long text alone.                |
| User expression   | Emotional (E)    | Explicit affective expression, such as frustration, gratitude, anxiety or excitement.                                                                 | Coded separately from cognitive engagement and interpreted as visible expression in the log.         |
| Assistant support | Scaffolding      | Assistant support that structures, diagnoses, regulates or redirects the user's process while leaving cognitive or task responsibility with the user. | Distinguished from reference responses that simply provide an answer, draft or deliverable.          |
| Support intent    | I1 metacognitive | Support aimed at planning, monitoring, reflection or regulation of the user's work.                                                                   | Captures the purpose of support, not its surface form.                                               |
| Support intent    | I2 cognitive     | Support aimed at concepts, mechanisms, task knowledge or problem-solving content.                                                                     | Most scaffolded support in this corpus is cognitive in purpose.                                      |
| Support intent    | I3 affective     | Support aimed at encouragement, reassurance or affective regulation.                                                                                  | Used only when affective support is substantively present, not for polite tone alone.                |
| Support form      | M1 feedback      | Evaluative information about the user's prior attempt, diagnosis or understanding.                                                                    | Requires a target in the user's preceding contribution.                                              |
| Support form      | M2 hinting       | A cue, prompt or partial path that helps the user proceed without fully resolving the task.                                                           | Leaves the next inferential or task step with the user.                                              |
| Support form      | M3 instructing   | Explicit procedural guidance, ordered steps or strategy for doing the task.                                                                           | More directive than hinting; may still count as scaffolding when it structures process.              |
| Support form      | M4 explaining    | A rationale, mechanism or causal account that clarifies why something works or fails.                                                                 | Requires explanatory content, not merely a restatement of an answer.                                 |
| Support form      | M5 modelling     | A reusable pattern, worked example or demonstration intended to be adapted.                                                                           | Distinguished from simply giving the final requested product.                                        |
| Support form      | M6 questioning   | A question that advances diagnosis, reflection, planning or problem solving.                                                                          | Excludes generic clarification questions that do not support the user's process.                     |

User-turn labels describe observable engagement. Passive engagement (P) marks acknowledgement, reception or minimal continuation without visible transformation. Active engagement (A) marks task-relevant action or follow-up, such as applying supplied information or requesting a concrete modification. Constructive engagement (C) marks deeper work with ideas, such as testing an implication, revising a hypothesis, articulating a constraint, transferring a mechanism or extending an explanation. Emotional engagement is labelled separately and used descriptively because visible affective expression was rare.

Assistant-turn labels describe support. Scaffolded support denotes assistant behaviour that structures, regulates, diagnoses or redirects the user’s work while leaving some responsibility with the user. Non-scaffolded reference responses provide the requested information or deliverable without materially structuring the user’s process. Scaffolded turns can also receive support-intent labels and support-form labels. Support-intent labels distinguish metacognitive support (I1), cognitive support (I2) and affective support (I3). Support-form labels distinguish feedback (M1), hinting (M2), instructing (M3), explaining (M4), modelling (M5) and questioning (M6). Support-form labels are non-mutually exclusive because a single assistant turn can both explain a mechanism and provide an example or hint. Supplementary Tables C18 and C19 report support-intent and support-form profiles across the six task settings.

### A.3 Production annotation workflow

The production workflow had three reproducible stages. First, the authors developed a theory-grounded codebook and prompt templates on validation examples and disagreement cases. This stage identified recurring failure modes, including over-labelling brief acknowledgements as active engagement, treating any long assistant answer as scaffolding, collapsing directive instructions into scaffolding without checking whether responsibility remained with the user and over-labelling sparse support forms. Second, parser checks enforced the expected schema, normalized label variants and flagged missing or malformed outputs. Third, production-scale annotation applied the frozen schema to the six analytic settings, followed by narrowly specified rule-based corrections for known boundary errors.

Production annotation used Azure OpenAI chat-completions deployments from the GPT-5.1 family for turn-level user and assistant labels, with GPT-5.1-chat used for the final writing-oriented calls. LLM-assisted semantic filtering used GPT-4o-family deployments. The released analysis code and archived annotation metadata retain the full prompt templates, functional prompt role, schema, deployment family, API version, retry settings and parser checks. In the manuscript and supplement we refer to prompt function rather than internal prompt-path names, because those path names are run-management identifiers rather than scientific constructs.

The production calls used structured JSON outputs where supported. The final writing calls used JSON-object or JSON-schema response constraints, 60-second request timeouts, maximum completion budgets of 300 tokens for conversation-level labels and 350 tokens for turn-level labels and did not pass an explicit temperature parameter. The final coding-from-filtered-conversation runner likewise did not pass an explicit temperature parameter. English-language checks used temperature 0.0 and

semantic task checks used temperature 0.1. Retry policies were specified in code: coding runs used up to five attempts with exponential backoff for rate limits and fixed-delay retry for other transient failures; writing runs used up to six attempts, disabled SDK-level retries, used exponential backoff with jitter for rate limits, timeouts and server errors and fell back from JSON-schema response format to JSON-object format if unsupported. Content-filter blocks returned empty analyses and were logged rather than silently imputed.

Post-processing was part of the predefined production annotation workflow, not an after-the-fact statistical adjustment. We use the term for two transparent operations. Schema-level processing normalized output variants and recovered fields that were present in the model explanation but missing from the structured JSON when the mapping was unambiguous. Boundary-rule processing corrected recurring, narrow false-positive patterns identified during validation, especially cases where brief user acknowledgements were labelled as Active or Constructive, generic emotional tone was labelled as Emotional or direct answer delivery was labelled as scaffolded support. Coding outputs did not use a broad post-processing chain in the final filtered-conversation runner, but WildChat user turns underwent the targeted user-turn review described in Section A.4. Writing outputs used the retained production chain for assistant scaffolding recovery, support-detail filling and strict user Active, Constructive and Emotional boundary filtering. API keys and service endpoints are intentionally not reproduced in the manuscript.

#### A.4 Human validation and production-label reliability

Human review was used during codebook development, prompt iteration and final validation. Two trained human reviewers with human–AI interaction and LLM-interaction research backgrounds reviewed validation cases. The purpose of human review changed over the development cycle: early review identified codebook boundary problems and prompt failure modes, whereas the final validation tables report agreement after the codebook, parser checks and production workflow had been stabilized. The final agreement metrics should therefore be read as evidence for the frozen workflow, not as a claim that high agreement was achieved at the start of development.

Validation evidence is reported in three layers. First, human–human agreement assesses whether trained reviewers could apply the codebook consistently; Supplementary Table C4 reports this codebook-operability layer. Second, production-label checks compare final or post-processed LLM-assisted production labels with human-confirmed labels; Supplementary Table C5 consolidates these human–LLM agreement checks. Third, targeted post-production audits test whether narrow rule-based reviews reduced known residual false positives after prompt and parser development. Coding validation combined three review rounds, with 658 task turns and 643 turns after deduplication. Writing validation used 210 reviewed conversations with final human-reviewed labels. Because many labels are sparse or prevalence-imbalanced, agreement is reported with per-label F1, Matthews correlation coefficient and Gwet’s AC1 rather than with a single global agreement statistic.

The final user-turn check contained 100 coding examples and 100 writing examples reviewed after the targeted user-turn review. These examples were used to identify residual false positives and calibrate narrow rule-based corrections, especially for Active and Passive user-turn labels. The final WildChat user-turn review scanned 31,878 coding conversations and 39,534 writing conversations, applying 1,815 coding corrections and 21 writing corrections. LMSYS Chat and ShareChat were then processed with the same refreshed annotation workflow, construct definitions and post-processing logic. The main manuscript reports the labels produced by these final production runs.

Because user framing carries an important contextual role, the consolidated production-label audit includes a human-verified check for the conversation-level learning-intent label. This audit sampled 450 first user turns, with 75 cases from each corpus-by-task setting and with production-positive cases deliberately oversampled. Overall intentional-framing agreement with the production label was  $F1=0.857$ ,  $MCC=0.677$  and  $Gwet\ AC1=0.671$ ; task-stratified  $F1$  was 0.924 for coding and 0.779 for writing. This audit validates the label boundary for user framing but is not used to estimate population prevalence because of the positive-oversampled design.

For user engagement production labels, the consolidated audit includes a 600-case check across the six corpus-by-task settings, with 100 user turns per setting and production constructive cases deliberately oversampled. The final labels were human-confirmed after item-level review. Against these human-confirmed labels, the production constructive-versus-non-constructive distinction had  $F1=0.922$ ,  $MCC=0.840$  and  $Gwet\ AC1=0.837$ . One-vs-rest metrics for the four user engagement levels were also high:  $F1=0.922$  for constructive, 0.883 for active, 0.941 for passive and 0.930 for none (Supplementary Table C5). Because the sample is positive-oversampled, it validates label boundaries and production-label reliability rather than estimating population prevalence.

For assistant production labels, the consolidated audit includes a 180-case check across the six corpus-by-task settings, with 30 assistant turns per setting and scaffolded-support and support-form positives deliberately oversampled. The final labels were human-confirmed after item-level review. Broad scaffolded-support agreement was  $F1=0.912$ ,  $MCC=0.590$  and  $Gwet\ AC1=0.791$ . Support-form one-vs-rest  $F1$  ranged from 0.691 to 0.938 across M1–M6 (Supplementary Table C5). Because the sample is positive-oversampled, it validates assistant-label boundaries rather than estimating support-form prevalence.

The validation evidence for turn-level engagement and support labels is therefore layered: human–human agreement supports codebook operability, targeted production audits directly check final user engagement and assistant scaffolding/support-form production labels across all six settings and post-update audits check known residual errors. Cross-corpus use of the labels is supported by applying the same schema, parser checks and post-processing rules to all three corpora and by reporting setting-level robustness checks in the Results and Supplementary Tables.

## A.5 Derived variables and denominator rules

This section defines how turn labels become the quantities plotted or modelled in the Results. At the turn level, the focal user-side outcome is a binary indicator for constructive engagement. At the conversation level, we use both the count of constructive user turns and an indicator for whether at least one constructive user turn occurs. Constructive ratios divide constructive user turns by total user turns; group-level constructive-ratio contrasts are turn-weighted ratios computed from total constructive user turns over total user turns within each group. Cognitive engagement ratios include passive, active and constructive cognitive engagement. Unless otherwise noted, constructive ratios use all user turns in the corresponding conversation or group as the denominator.

At the assistant side, scaffolding is operationalized as scaffolded support. Conversation-level scaffolding presence indicates whether a conversation contains at least one scaffolded assistant turn. Support-form analyses condition on conversations or turns with scaffolded support and then compare the presence versus absence of each support-form label. These comparisons are descriptive because support-form labels can co-occur and because conversations are not randomly assigned to receive a support form.

Behavioural depth is computed from turn counts and turn ordering. Post-answer depth counts turns after the first assistant response. Persistence after failure is analysed only for coding-oriented conversations because debugging failures, error messages and failed-execution contexts are more directly observable in coding than in writing.

## A.6 Statistical estimands and model specifications

This section maps each inferential quantity in the Results to its denominator, bootstrap unit, model formula, exposure specification and sensitivity checks. Descriptive prevalence analyses report proportions within the six task settings. Context analyses compare user framing, task ecology and depth groups. Conversation-level support associations use Poisson generalized linear models for constructive-turn counts and logistic models for binary constructive-engagement outcomes. Adjacent-turn analyses use assistant-to-user and user-to-assistant turn pairs to estimate conditional probabilities at the next-turn scale. Supplementary Tables C6, C8 and C9 provide the compact estimand map.

For the Section 2.3 adjusted conversation-level models, the Poisson outcome is the raw count of constructive user turns in a conversation, not a rate. The primary model does not include a user-turn offset. Conversation length is adjusted directly through total turns, so the exponentiated Poisson coefficient should be read as a multiplicative association with expected constructive-turn count conditional on observed conversation length and other covariates. The generic setting-specific mean model is  $\mathbb{E}(C_i) = \exp(\alpha_s + \beta_s S_i + \gamma_s^\top X_i)$ , where  $C_i$  is constructive-turn count,  $S_i$  is scaffolded-support presence and  $X_i$  contains the observed conversation-level context listed in Supplementary Table C8. The companion logistic model predicts whether a conversation contains any constructive user turn using scaffolded-support presence and the available setting-specific context. The user-framing stratified Poisson models use the



same Poisson specification after stratifying by framing and omitting user framing from the covariate set.

Covariates are included when they correspond to observed compositional differences at the analysis level. Depending on the analysis, these include user framing, task indicators in pooled models, turn counts, emotional ratios, error-related variables and turn index. Percentage-point contrasts are reported as effect sizes. Confidence intervals and p values for these contrasts use 1,000 conversation-level bootstrap resamples; adjacent-turn contrasts use 1,000 conversation-cluster bootstrap resamples rather than individual turn-pair resampling to account for repeated pairs within conversations. The adjusted Poisson and logistic estimates in Fig. 3 use model-based standard errors. Poisson overdispersion was checked with the Pearson  $\chi^2$  statistic divided by residual degrees of freedom. Dispersion ranged from 1.16 to 2.06 across the six task settings; scaling the Poisson standard errors with a quasi-Poisson dispersion adjustment preserved positive scaffolded-support associations and statistical distinguishability in every setting. To check exposure specification, we also refitted the same setting-specific Poisson mean model with user turns represented as the exposure offset; Supplementary Table C11 reports quasi-Poisson scaled rate ratios. Negative-binomial models were not used as the primary specification; the sensitivity analyses focus on variance misspecification and exposure specification under the same mean model rather than changing the main estimand. Support-form bracket tests use Benjamini–Hochberg false-discovery-rate adjustment within each six-form comparison family.

The Section 2.2 constructive-rate sensitivity uses a grouped-binomial logistic model at the conversation level. The numerator is the number of constructive user turns and the denominator is total user turns in the conversation. The predictors match the any-constructive context model: user framing, task ecology, length bucket and dataset fixed effects. Standard errors use a conversation-robust sandwich estimator. This sensitivity is included because any-constructive-turn outcomes partly reflect the number of turns available for observing at least one constructive turn.

The integrated adjacent-turn regression combines three feature groups in one logistic model: user state before the assistant turn, assistant scaffolding features and adjacent-turn context. The pooled specification includes dataset fixed effects. A model/source sensitivity uses exact WildChat model labels, LMSYS Chat model names recovered from conversation identifiers and ShareChat assistant/source families recovered from conversation identifiers. These adjusted estimates are interpreted as measured-covariate-adjusted associations. They do not remove unobserved user motivation, prior expertise, task difficulty or time-varying conversational state.

## A.7 Temporal estimands and interpretive boundary

The temporal appendix clarifies the scale of the temporal claim. The logs identify within-conversation ordering, but they do not randomize assistant support or observe durable learning outcomes. We therefore distinguish adjacent-turn estimands from coarser within-conversation summaries. The adjacent-turn analyses ask whether the immediately following user turn is more likely to be constructive after scaffolded support than after a non-scaffolded reference response, whether that contrast depends on the prior user state and whether support-form signals differ across prior user states.

The reverse adjacent-turn analysis asks whether the assistant is more likely to provide scaffolded support after constructive, active or passive user turns.

Coarser before-after summaries ask a different question: whether constructive engagement is higher after the first scaffolded assistant turn than before it. These are retained only as exploratory boundary checks, not as treatment-effect or accumulation estimates. The first scaffolded turn may occur precisely because the user is struggling, because a difficult task has emerged or because the exchange is already developing. The main temporal claim is therefore scale-specific: the clearest evidence concerns adjacent-turn, state- and form-conditioned coupling rather than a demonstrated global trajectory or cumulative learning gain.

## Appendix B Robustness checks and interpretive boundaries

### B.1 Cross-dataset consistency

The primary robustness checks are cross-dataset checks, not WildChat-only checks. Whenever an estimand can be computed from the shared annotation schema, it is reported across all six WildChat, LMSYS Chat and ShareChat task settings. These six-setting checks include the main prevalence summaries, constructive-ratio contrasts, post-answer depth contrasts, adjacent-turn lifts, support-form supply profiles and setting-level adjacent-turn regressions (Supplementary Tables C10, C14 and C19; Supplementary Figures D1–D3).

The Results claims are checked at the level appropriate to each estimand. Section 2.1 is a descriptive prevalence claim: cognitive engagement, constructive engagement and scaffolded support were observable in all six settings, with constructive engagement consistently the highest-level and less frequent user-side signal. Section 2.2 is primarily descriptive: user framing, task ecology and interaction depth organize where constructive engagement appears, with the same directional depth gradient across settings. Formal uncertainty estimates are reported where the claim involves a contrast or model coefficient. Section 2.3 uses conversation-level bootstrap confidence intervals and p values for turn-weighted scaffolded versus non-scaffolded constructive-ratio contrasts, adjusted Poisson/logistic model estimates and an offset-rate Poisson sensitivity. Section 2.4 reports support-form confidence intervals and between-stratum FDR-adjusted q values. Section 2.5 reports conversation-cluster bootstrap confidence intervals for adjacent-turn lifts and cluster-robust regression estimates for the integrated adjacent-turn models.

The main percentage-point contrasts were directionally consistent across the six settings. Constructive-ratio contrasts comparing conversations with versus without scaffolded support were positive and statistically distinguishable from zero in all six settings, with effect sizes ranging from +0.99 to +7.34 percentage points. Adjacent-turn constructive lifts were also positive in all six settings, with five of six contrasts statistically distinguishable from zero; LMSYS writing was positive but weaker (+0.27 percentage points, 95% CI -0.01 to +0.56,  $p = .061$ ). Supplementary Table C10 reports these effect sizes, confidence intervals and p values.

## B.2 Model/source robustness

The only robustness displays restricted to WildChat are the model-snapshot figures in Supplementary Figures D5–D9. They are restricted to WildChat because WildChat exposes user-facing model metadata with sufficient coverage for repeated within-corpus snapshot plots. LMSYS Chat and ShareChat are still included in the cross-dataset and model/source-adjusted analyses, but they do not provide directly comparable model-snapshot structure for the same visual repetition. The WildChat model-snapshot analyses preserved the main directional patterns reported in the main text: coding generally showed higher cognitive engagement than writing; intentionality gaps were visible across most user-facing models; scaffolded conversations were generally more constructive; and adjacent-turn contrasts were more interpretable than coarser within-conversation contrasts.

Integrated adjacent-turn regressions were used as robustness checks rather than as the primary narrative claim. These models predicted whether the next user turn was constructive after combining user-state features, assistant-support features and adjacent-turn context. Because model metadata differ by corpus, we report both within-setting models and pooled specifications. The within-setting models show the dataset-by-task pattern directly, with recoverable model/source fixed effects added separately inside each setting (Supplementary Table C14). The pooled specification includes dataset fixed effects; a sensitivity replaces these with exact WildChat model labels, LMSYS model names recovered from conversation identifiers and ShareChat assistant/source families recovered from conversation identifiers. This is interpreted as model/source adjustment rather than as a causal comparison among model versions. Model/source heterogeneity was detectable, but broad scaffolded support still improved prediction after user-state, context and model/source controls (LR  $\chi^2_1 = 509.7$ ,  $p < .001$ ). Because M1–M6 are support forms nested within scaffolded support rather than independent covariates, we interpret the decomposed models as asking which forms carry additional signal among scaffolded turns, not as a test of scaffolded support after removing its own components. Adding M1–M6 descriptors within scaffolded support improved fit (LR  $\chi^2_6 = 829.8$ ,  $p < .001$ ), and the full scaffolding block remained strongly informative after controls (LR  $\chi^2_7 = 1339.6$ ,  $p < .001$ ; Supplementary Table C16). Adding prior-state  $\times$  support-form interactions also improved fit (pooled model/source fixed effects: LR  $\chi^2_{18} = 297.3$ ,  $p < .001$ ; Supplementary Table C17). Fully decomposed coefficients clarify the source of this form-level signal: prior user state is the largest predictor of immediate constructive follow-up, and M4 explaining is the most stable positive support-form coefficient (Supplementary Table C15).

All model/source checks are used as robustness analyses rather than model-causal comparisons. Model labels are not experimentally assigned. They may be confounded with calendar time, user population, task mix, interface changes and logging context. For this reason, the manuscript does not claim that a specific model family caused the observed engagement differences.

### B.3 Label uncertainty

The label scheme intentionally separates top-level constructs from finer support forms. Top-level labels such as constructive engagement and scaffolded support have clear behavioural anchors and the strongest validation evidence; they support the main prevalence, association and temporal claims. Finer support-form labels are interpreted as aggregate form-pattern signals rather than isolated item-level judgments. This is why the main claims rely most heavily on constructive engagement, scaffolded support presence, support-form patterns that are consistent across settings and temporal contrasts that can be expressed with simple conditional probabilities.

Labels with lower prevalence or more ambiguous boundaries are treated more cautiously. Visible emotional expression is reported descriptively because it was rare. Active engagement helps describe the engagement ecology, but the central outcome is constructive engagement. Support-form labels M1–M6 can co-occur and are not interpreted as mutually exclusive exposures. The human–human and human–LLM validation tables therefore support the use of the schema at the aggregate level, while the manuscript avoids claims that depend on adjudicating a single ambiguous turn in isolation.

### B.4 Selection into support

The most important threat to causal interpretation is selection into support. More complex tasks, more motivated users, better-specified questions or already-developing exchanges may both elicit scaffolded support and produce constructive follow-up. This possibility is part of the interactional process that the paper studies. The empirical question is therefore whether scaffolded support remains associated with constructive participation after observed context and immediate prior user state are taken into account, not whether support was randomly assigned.

Conversation-level adjusted models reduce measured compositional differences but cannot remove unobserved selection. Adjacent-turn analyses narrow the temporal window and condition on the immediately preceding user state, which helps distinguish a broad conversation-level association from a state-conditioned coupling pattern. Reverse adjacent-turn summaries also show that user state predicts subsequent assistant scaffolding, consistent with a coupled process rather than a one-directional exposure. The safest interpretation is therefore state-dependent coupling between assistant support and user engagement, not a claim that scaffolded support alone caused learning.

### B.5 Privacy-preserving reporting

The study analyses naturalistic conversations, so the manuscript avoids reproducing verbatim raw conversation logs. Operational examples are de-identified and paraphrased or lightly edited to illustrate decision rules without exposing user content. Aggregate statistics, model summaries and validation metrics are reported in the manuscript and supplementary tables. The analyses do not link conversations into user histories or estimate longitudinal user-level trajectories; temporal claims are restricted to within-conversation turn ordering.

**Table C1 Operational examples for turn-level labels.** Examples are de-identified and paraphrased from validation review patterns; they illustrate the decision rule rather than reproducing raw conversation text.

| Construct                      | Paraphrased turn                                                                                             | Annotation rationale                                                                                         |
|--------------------------------|--------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| P passive user engagement      | "Thanks, that makes sense."                                                                                  | Minimal uptake or acknowledgement without new reasoning, evaluation or task transformation.                  |
| A active user engagement       | "Can you convert that answer into a command I can run in my setup?"                                          | Task-relevant follow-up or manipulation of supplied information, but without elaborating or testing an idea. |
| C constructive user engagement | "If the cache expires before the asynchronous callback returns, would the same race condition still happen?" | The user tests an implication, boundary condition or transfer of the prior explanation.                      |
| Scaffolding                    | "First isolate the failing step; that will show whether the issue is parsing or state management."           | The assistant structures the user's process and leaves diagnostic or revision work with the user.            |
| M1 feedback                    | "Your diagnosis is close; the bug is in the index offset, not the loop boundary."                            | Evaluative feedback on the user's prior attempt or understanding.                                            |
| M2 hinting                     | "Check the parser log around the first malformed token before changing the model."                           | A cue or partial path that helps the user proceed without fully resolving the task.                          |
| M3 instructing                 | "First split the condition into two branches, then add a regression test."                                   | Explicit steps or strategy that organizes the user's work.                                                   |
| M4 explaining                  | "The exception occurs because the event loop is already running, so a nested call cannot acquire it."        | A rationale or mechanism that clarifies why the issue occurs.                                                |
| M5 modeling                    | "A small pattern you can adapt is: validate input, normalize it, then pass it to the writer."                | A sample approach intended to be emulated or adapted rather than merely copied.                              |
| M6 questioning                 | "Which constraint matters more here: preserving tone or shortening the paragraph?"                           | A question that advances problem solving or reflection.                                                      |

## Appendix C Supporting tables

**Table C2 Complete de-identified coding case with turn-level labels.** The illustrative transcript is derived from a retained conversation and lightly edited while preserving turn sequence, task structure and annotation decisions. Colored badges encode role, user engagement level and assistant support category.

| <b>Coding case</b> SQL transaction handling<br>WildChat coding; intentional framing; four turns. Constructive user follow-up probes a programming mechanism; scaffolded support provides explanation, procedure and modelling. |                                                                                                                                                                                                                                                                             |                                                                                                                              |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------|
| <b>Encoding</b> ↑ C constructive → A active * Scaffolded · Reference                                                                                                                                                           |                                                                                                                                                                                                                                                                             |                                                                                                                              |
| <b>Codes</b> I2 support intent; M support-form labels may co-occur.                                                                                                                                                            |                                                                                                                                                                                                                                                                             |                                                                                                                              |
| Turn                                                                                                                                                                                                                           | De-identified transcript                                                                                                                                                                                                                                                    | Label and annotation cue                                                                                                     |
| ▷ U1<br>User                                                                                                                                                                                                                   | “How does XACT_ABORT work in a SQL Server script? If my delete statements are inside a transaction so failed deletes should roll back, should XACT_ABORT be on or off? Does it apply only to the transaction or to the whole script?”                                       | ↑ C constructive<br>The user probes a <b>mechanism</b> , asks about <b>scope</b> and links the answer to rollback behaviour. |
| ◊ A1<br>Assistant                                                                                                                                                                                                              | The assistant explains that XACT_ABORT ON makes many run-time errors abort and roll back the current transaction, distinguishes transaction-level rollback from statements outside the transaction and explains why the setting is commonly used for update/delete scripts. | * Scaffolded I2 M4<br>The reply <b>explains the mechanism</b> connecting errors, transactions and rollback.                  |
| ▷ U2<br>User                                                                                                                                                                                                                   | “Do I need to set XACT_ABORT OFF at the end because of scope? What about using TRY/CATCH outside a transaction instead?”                                                                                                                                                    | ↑ C constructive<br>The user extends the prior answer by comparing <b>scope</b> and an <b>alternative strategy</b> .         |
| ◊ A2<br>Assistant                                                                                                                                                                                                              | The assistant explains the session-scoped setting, gives a cautious cleanup pattern, contrasts XACT_ABORT with TRY/CATCH and sketches a safer transaction pattern with BEGIN TRY, COMMIT, BEGIN CATCH and ROLLBACK.                                                         | * Scaffolded I2 M3 M4 M5<br>The reply combines <b>procedure, explanation</b> and an adaptable model.                         |

**Table C3 Complete de-identified writing case with turn-level labels.** The illustrative transcript is derived from a retained conversation and lightly edited while preserving turn sequence, task structure and annotation decisions. Colored badges encode role, user engagement level and assistant support category.

| <p><b>Writing case</b> <b>Formal report justification</b><br/> WildChat writing; intentional framing; six turns. The exchange moves from direct explanation use to constructive reframing, then to direct drafting under a new style constraint.</p> |                                                                                                                                                                                                                                            |                                                                                                                                                                          |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <p><b>Encoding</b> <span>↑ C constructive</span> <span>→ A active</span> <span>* Scaffolded</span> <span>· Reference</span></p> <p><b>Codes</b> <span>I2</span> support intent; <span>M</span> support-form labels may co-occur.</p>                 |                                                                                                                                                                                                                                            |                                                                                                                                                                          |
| Turn                                                                                                                                                                                                                                                 | De-identified transcript                                                                                                                                                                                                                   | Label and annotation cue                                                                                                                                                 |
| <span>▷ U1</span><br>User                                                                                                                                                                                                                            | “Help me explain why a machine-learning dataset should include difficult samples and not just clean samples, so that the trained model will be more robust. Give me a more detailed explanation.”                                          | <span>→ A active</span><br>The user requests detail but does not yet evaluate, revise or extend an idea.                                                                 |
| <span>◊ A1</span><br>Assistant                                                                                                                                                                                                                       | The assistant gives a direct explanation: diverse and difficult samples expose the model to real-world variation, reduce dependence on overly clean patterns and improve generalization to noisier cases.                                  | <span>· Reference</span><br>The reply supplies requested content without structuring the user’s next reasoning or writing process.                                       |
| <span>▷ U2</span><br>User                                                                                                                                                                                                                            | “Suppose I am explaining this to someone who does not believe that a model trained only on clean data generalizes worse than one trained with diverse, difficult data. Help me explain why this is the case in more detail.”               | <span>↑ C constructive</span><br>The user <b>reframes the task</b> around persuasion and justification, extending the explanation into a more demanding context.         |
| <span>◊ A2</span><br>Assistant                                                                                                                                                                                                                       | The assistant explains the mechanism with an analogy and rationale: clean-only training narrows the problem space, whereas difficult examples expose edge cases, noise and variation that make the learned decision boundary less brittle. | <span>* Scaffolded</span> <span>I2</span> <span>M4</span><br>The reply teaches the <b>conceptual mechanism</b> behind robustness rather than merely restating the claim. |
| <span>▷ U3</span><br>User                                                                                                                                                                                                                            | “I need a different type of explanation for a formal report. I should justify the need for a diverse dataset without relying on analogies. Can you help me do that?”                                                                       | <span>→ A active</span><br>The user specifies a <b>new writing constraint</b> and asks for a format shift, mainly as directive refinement.                               |
| <span>◊ A3</span><br>Assistant                                                                                                                                                                                                                       | The assistant provides a formal report-style explanation linking sample diversity to generalization, deployment variation, robustness and evaluation risk.                                                                                 | <span>· Reference</span><br>The reply directly supplies formal wording rather than adding process guidance or feedback on the user’s own draft.                          |

**Table C4 Domain-stratified human–human codebook agreement for labels used in the main analyses.** Coding uses the latest three-round validation rerun over 658 task turns, 643 after deduplication; writing uses the final writing-210 validation labels. These validation sets assess whether human reviewers could apply the coding and writing codebook consistently; they are not, by themselves, direct estimates of production LLM-label accuracy and are not separately powered within each public corpus. F1, Matthews correlation coefficient (MCC) and Gwet’s AC1 are reported per binary label. Coding MCC and AC1 are recomputed from the exported support and positive-count fields; writing metrics are from the final writing agreement report. Rare degenerate labels with no positive cases in a validation set are omitted. Production-label checks are reported in Supplementary Table C5, and post-update audits are described in Supplementary Section A.4.

| Domain  | Role      | Label            | F1   | MCC  | Gwet AC1 | Support |
|---------|-----------|------------------|------|------|----------|---------|
| Coding  | User      | Active (A)       | 0.75 | 0.72 | 0.87     | 315     |
| Coding  | User      | Constructive (C) | 0.85 | 0.82 | 0.90     | 315     |
| Coding  | User      | Passive (P)      | 0.88 | 0.87 | 0.96     | 315     |
| Coding  | User      | Emotional (E)    | 0.83 | 0.82 | 0.97     | 315     |
| Coding  | Assistant | Scaffolding      | 0.87 | 0.69 | 0.70     | 328     |
| Coding  | Assistant | Feedback (M1)    | 0.90 | 0.89 | 0.95     | 328     |
| Coding  | Assistant | Hinting (M2)     | 0.79 | 0.75 | 0.89     | 328     |
| Coding  | Assistant | Instructing (M3) | 0.81 | 0.77 | 0.92     | 328     |
| Coding  | Assistant | Explaining (M4)  | 0.86 | 0.79 | 0.84     | 328     |
| Coding  | Assistant | Modeling (M5)    | 0.84 | 0.79 | 0.87     | 328     |
| Coding  | Assistant | Questioning (M6) | 0.82 | 0.81 | 0.95     | 328     |
| Writing | User      | Active (A)       | 0.73 | 0.68 | 0.81     | 90      |
| Writing | User      | Constructive (C) | 0.89 | 0.83 | 0.86     | 90      |
| Writing | User      | Passive (P)      | 0.80 | 0.81 | 0.99     | 90      |
| Writing | User      | Emotional (E)    | 0.80 | 0.79 | 0.98     | 90      |
| Writing | Assistant | Scaffolding      | 0.97 | 0.80 | 0.93     | 120     |
| Writing | Assistant | Feedback (M1)    | 0.89 | 0.80 | 0.80     | 120     |
| Writing | Assistant | Hinting (M2)     | 0.82 | 0.70 | 0.71     | 120     |
| Writing | Assistant | Instructing (M3) | 0.71 | 0.67 | 0.86     | 120     |
| Writing | Assistant | Explaining (M4)  | 0.84 | 0.70 | 0.70     | 120     |
| Writing | Assistant | Modeling (M5)    | 0.94 | 0.91 | 0.95     | 120     |
| Writing | Assistant | Questioning (M6) | 0.96 | 0.96 | 0.99     | 120     |



**Table C5 Consolidated human–LLM production-label agreement.** User-framing audit sampled 450 first user turns across the six corpus-by-task settings, with production-positive intentional-framing cases deliberately oversampled. User engagement audit sampled 600 user turns, with 100 cases per setting and production constructive cases deliberately oversampled. Assistant audit sampled 180 assistant turns, with 30 cases per setting and scaffolded-support and support-form positives deliberately oversampled. Final human labels were confirmed after item-level review. Metrics compare final LLM-assisted production labels with human-confirmed labels using one-vs-rest binary coding for each label. Because the samples are positive-oversampled, these metrics validate label boundaries rather than estimating population prevalence.

| Role                                           | Label               | Audit N | Prod. + | Human + | Precision | Recall | F1    | MCC   | Gwet AC1 | Accuracy |
|------------------------------------------------|---------------------|---------|---------|---------|-----------|--------|-------|-------|----------|----------|
| <i>Conversation-level framing</i>              |                     |         |         |         |           |        |       |       |          |          |
| Conversation                                   | Intentional framing | 450     | 300     | 239     | 0.770     | 0.967  | 0.857 | 0.677 | 0.671    | 0.829    |
| <i>User engagement</i>                         |                     |         |         |         |           |        |       |       |          |          |
| User                                           | Constructive (C)    | 600     | 300     | 327     | 0.963     | 0.884  | 0.922 | 0.840 | 0.837    | 0.912    |
| User                                           | Active (A)          | 600     | 120     | 111     | 0.850     | 0.919  | 0.883 | 0.856 | 0.935    | 0.958    |
| User                                           | Passive (P)         | 600     | 90      | 80      | 0.889     | 1.000  | 0.941 | 0.934 | 0.978    | 0.983    |
| User                                           | None                | 600     | 90      | 82      | 0.889     | 0.976  | 0.930 | 0.920 | 0.974    | 0.980    |
| <i>Assistant scaffolding and support forms</i> |                     |         |         |         |           |        |       |       |          |          |
| Assistant                                      | Scaffolding         | 180     | 144     | 139     | 0.896     | 0.928  | 0.912 | 0.590 | 0.791    | 0.861    |
| Assistant                                      | Feedback (M1)       | 180     | 45      | 51      | 1.000     | 0.882  | 0.938 | 0.918 | 0.945    | 0.967    |
| Assistant                                      | Hinting (M2)        | 180     | 43      | 45      | 0.930     | 0.889  | 0.909 | 0.880 | 0.930    | 0.956    |
| Assistant                                      | Instructing (M3)    | 180     | 51      | 55      | 0.745     | 0.691  | 0.717 | 0.600 | 0.715    | 0.833    |
| Assistant                                      | Explaining (M4)     | 180     | 87      | 110     | 1.000     | 0.791  | 0.883 | 0.772 | 0.747    | 0.872    |
| Assistant                                      | Modelling (M5)      | 180     | 48      | 79      | 0.958     | 0.582  | 0.724 | 0.631 | 0.642    | 0.806    |
| Assistant                                      | Questioning (M6)    | 180     | 35      | 20      | 0.543     | 0.950  | 0.691 | 0.675 | 0.873    | 0.906    |

**Table C6 Analytic model families used in the main analyses.** The table summarizes descriptive and conversation-level model families. Adjusted models are interpreted as covariate-adjusted associations in naturalistic data, not as causal treatment-effect estimates.

| Analysis family        | Unit and outcome                                                                           | Main contrast or model                                                                                                                                                                                                                                                                                                          | Scope                                                                                                                                                                                                                                                                    |
|------------------------|--------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Engagement prevalence  | User turns and conversations; cognitive, constructive and emotional engagement indicators. | Descriptive proportions by task setting, task ecology and user framing.                                                                                                                                                                                                                                                         | Establishes where learning-oriented engagement is visible; no causal contrast is implied.                                                                                                                                                                                |
| Engagement composition | Cognitively engaged user turns; passive, active and constructive depth levels.             | Composition of cognitive engagement within each task setting.                                                                                                                                                                                                                                                                   | Motivates constructive engagement as the focal high-level outcome while preserving the broader cognitive-engagement hierarchy.                                                                                                                                           |
| Interaction depth      | Conversations; at least one constructive user turn.                                        | Probability of constructive engagement across behavioural-depth buckets.                                                                                                                                                                                                                                                        | Depth is interpreted as interactional opportunity and task development, not as a manipulated exposure.                                                                                                                                                                   |
| Scaffolding presence   | Conversations; constructive-turn count and any constructive turn.                          | Setting-specific Poisson GLM for constructive-turn count with scaffolded-support presence and observed conversation-level context; no user-turn offset in the primary count model.<br>Setting-specific logistic model for whether any constructive user turn occurs. Estimable covariates are listed in Supplementary Table C8. | Adjusted estimates reduce measured compositional differences but remain associational because support is selected into by task difficulty and user state. Pearson overdispersion was checked and quasi-Poisson scaling preserved the direction of the Poisson estimates. |

**Table C7 Additional analytic model families and robustness checks.** These analyses qualify the main associations by support form, temporal scale and available model or source metadata.

| Analysis family                     | Unit and outcome                                                                                                     | Main contrast or model                                                                                                                                                                                      | Scope                                                                                                                                                |
|-------------------------------------|----------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| Support-form associations           | Scaffolded conversations; constructive engagement ratios.                                                            | Mean differences for conversations containing a given support-means label versus scaffolded conversations without that label, summarized by user framing and task ecology.                                  | Descriptive because means labels are non-exclusive and can co-occur in the same assistant turn.                                                      |
| Local temporal coupling             | Adjacent assistant–user or user–assistant turn pairs; next constructive user turn or next scaffolded assistant turn. | Conditional probability contrasts around adjacent turns; detailed estimands are listed in Supplementary Table C9.                                                                                           | Interpreted as local ordering and state-dependent coupling, not as randomized support effects.                                                       |
| Integrated adjacent-turn regression | Adjacent assistant–user turn pairs; next constructive user turn.                                                     | Logistic regression combining prior user state, scaffolded-support presence, support means, task/framing context, prior-state $\times$ support-form interactions and dataset or model/source fixed effects. | Tests whether the local coupling pattern depends jointly on user state and support form; standard errors are clustered by conversation.              |
| Coarse temporal boundary checks     | Conversations containing scaffolded support; constructive engagement before and after the first scaffolded turn.     | Mean after-minus-before contrast reported as an exploratory appendix check.                                                                                                                                 | Not used as evidence of accumulation because the first scaffolded turn may be selected by difficulty, uncertainty or an already developing exchange. |
| WildChat model robustness           | WildChat conversations or turns; main prevalence, association and temporal outcomes.                                 | Repetition across user-facing model snapshots or coarse model families.                                                                                                                                     | Robustness summaries only; model labels are observational and may be confounded with time, user mix and task mix.                                    |

**Table C8 Setting-specific covariate sets for the Section 2.3 adjusted models.** Each model was fitted separately within the six task settings. Task ecology is therefore fixed within a setting and is not interpreted as a within-setting coefficient. The table lists the estimable covariates used to adjust the scaffolded-support association; pooled context models use task ecology separately as shown in Supplementary Tables C12 and C13.

| Settings                                                                  | Primary Poisson count model                                                                                                                                             | Logistic any-constructive model                         | Sensitivities and notes                                                                                                                                                                                                                                                                                                                                                  |
|---------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| WC coding; LMSYS coding; SC coding; WC writing; LMSYS writing; SC writing | Scaffolded-support presence; user framing; total turns; emotional-expression ratio; observable error marker; persistence-after-failure marker; high-persistence marker. | Scaffolded-support presence; user framing; total turns. | The offset-rate sensitivity used the same Poisson mean model but represented user turns with an exposure offset. User-framing stratified Poisson models omitted user framing because it was fixed within stratum. Error and persistence markers are most interpretable in coding-oriented contexts but are retained as observed conversation-level context when present. |

**Table C9 Temporal estimands for adjacent-turn support–engagement coupling.** Main temporal analyses use adjacent turn pairs. Coarser before–after summaries are exploratory boundary checks because scaffolded turns are selected by task difficulty, user motivation and conversation trajectory.

| Estimand                            | Definition                                                                                                                                  | Reported in main text                 | Scope                                                                                                                       |
|-------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------|-----------------------------------------------------------------------------------------------------------------------------|
| Overall adjacent-turn lift          | Next-turn constructive probability after scaffolded support versus a non-scaffolded reference response, estimated within each task setting. | Figure 5a                             | Local association between support type and immediately following constructive engagement; not a randomized support effect.  |
| Prior-state adjacent contrast       | Same next-turn contrast, stratified by prior constructive, active and passive user states.                                                  | Figure 5b                             | Tests whether local coupling depends on the user's immediately preceding state.                                             |
| Reverse state-conditioned support   | Probability that the next assistant turn is scaffolded after constructive, active and passive user turns.                                   | Figure 5c                             | Describes how scaffolding supply varies with the user's current engagement state.                                           |
| Prior-state support-form regression | Adjacent-turn logistic regression for next-turn constructive engagement, adding prior-state $\times$ M1–M6 interactions.                    | Figure 5d and Supplementary Table C17 | Tests whether local coupling depends jointly on prior user state and support form.                                          |
| Around-first-scaffold contrast      | Mean constructive engagement after versus before the first scaffolded assistant turn, among conversations containing scaffolded support.    | Supplementary Figure D1               | Exploratory boundary check only; the first scaffolded turn may be selected by difficulty or an already developing exchange. |

**Table C10 Uncertainty estimates for key support–engagement contrasts.**

Constructive-ratio contrasts are turn-weighted conversation-bootstrap estimates, with ratios computed as total constructive user turns divided by total user turns within scaffolded versus non-scaffolded conversation groups. Post-answer depth contrasts use conversation-level bootstrap resampling. Adjacent-turn contrasts bootstrap conversations over assistant-to-user pairs. Percentage-point and turn differences are the effect sizes reported in the main figures.

| Setting       | Contrast                                  | Effect      | 95% CI         | p value |
|---------------|-------------------------------------------|-------------|----------------|---------|
| WC coding     | Scaffolded – reference constructive ratio | +4.56 pp    | [+4.15, +4.99] | < .001  |
| WC coding     | Scaffolded – reference post-answer depth  | +2.23 turns | [+2.11, +2.35] | < .001  |
| LMSYS coding  | Scaffolded – reference constructive ratio | +2.50 pp    | [+2.19, +2.83] | < .001  |
| LMSYS coding  | Scaffolded – reference post-answer depth  | +1.42 turns | [+1.34, +1.52] | < .001  |
| SC coding     | Scaffolded – reference constructive ratio | +7.34 pp    | [+5.29, +9.27] | < .001  |
| SC coding     | Scaffolded – reference post-answer depth  | +3.36 turns | [+2.69, +4.01] | < .001  |
| WC writing    | Scaffolded – reference constructive ratio | +0.99 pp    | [+0.82, +1.16] | < .001  |
| WC writing    | Scaffolded – reference post-answer depth  | +3.13 turns | [+2.96, +3.30] | < .001  |
| LMSYS writing | Scaffolded – reference constructive ratio | +1.09 pp    | [+0.87, +1.31] | < .001  |
| LMSYS writing | Scaffolded – reference post-answer depth  | +2.10 turns | [+1.91, +2.28] | < .001  |
| SC writing    | Scaffolded – reference constructive ratio | +3.20 pp    | [+1.55, +4.89] | < .001  |
| SC writing    | Scaffolded – reference post-answer depth  | +3.38 turns | [+1.92, +4.64] | < .001  |
| WC coding     | Adjacent scaffolded – reference lift      | +2.68 pp    | [+2.22, +3.09] | < .001  |
| LMSYS coding  | Adjacent scaffolded – reference lift      | +2.01 pp    | [+1.54, +2.50] | < .001  |
| SC coding     | Adjacent scaffolded – reference lift      | +6.21 pp    | [+4.35, +8.10] | < .001  |
| WC writing    | Adjacent scaffolded – reference lift      | +1.13 pp    | [+0.88, +1.41] | < .001  |
| LMSYS writing | Adjacent scaffolded – reference lift      | +0.27 pp    | [−0.01, +0.56] | .061    |
| SC writing    | Adjacent scaffolded – reference lift      | +3.35 pp    | [+0.80, +6.62] | .025    |

**Table C11 Offset-rate sensitivity for scaffolded-support Poisson models.** The primary Fig. 3b Poisson model uses raw constructive-turn counts with conversation length entered as a covariate. As a rate sensitivity, this table refits the same setting-specific mean model with `offset(log(user_turns))`. Cells report the scaffolded-support rate ratio with 95% confidence intervals and two-sided p values using quasi-Poisson scaled standard errors.

| Setting       | Offset rate ratio [95% CI], p value | Pearson dispersion | Conversations |
|---------------|-------------------------------------|--------------------|---------------|
| WC coding     | 1.57 [1.47, 1.68], < .001           | 1.29               | 31,878        |
| LMSYS coding  | 1.47 [1.38, 1.56], < .001           | 1.13               | 32,114        |
| SC coding     | 1.58 [1.31, 1.91], < .001           | 1.29               | 2,481         |
| WC writing    | 1.36 [1.26, 1.46], < .001           | 1.15               | 39,534        |
| LMSYS writing | 1.70 [1.50, 1.94], < .001           | 1.12               | 21,023        |
| SC writing    | 1.65 [1.27, 2.15], < .001           | 1.32               | 1,539         |

**Table C12 Conversation-level context model for constructive engagement.** The outcome is whether a conversation contains at least one constructive user turn. The logistic regression includes user framing, task ecology, conversation-length bucket and dataset fixed effects. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-robust standard errors. The reference category is unintentional framing, writing-oriented task, 2–3 user turns and WildChat.

| Predictor           | OR [95% CI], p value      |
|---------------------|---------------------------|
| Intentional framing | 3.03 [2.93, 3.13], < .001 |
| Coding task ecology | 1.89 [1.71, 2.08], < .001 |
| 4–6 user turns      | 2.07 [1.99, 2.15], < .001 |
| 7+ user turns       | 3.58 [3.42, 3.74], < .001 |
| LMSYS Chat          | 0.70 [0.68, 0.73], < .001 |
| ShareChat           | 2.29 [2.12, 2.46], < .001 |
| Conversations       | 128,569                   |

**Table C13 Conversation-level constructive-rate sensitivity model.** The outcome is the number of constructive user turns out of total user turns in each conversation, modelled with a grouped-binomial logistic regression. The model includes user framing, task ecology, conversation-length bucket and dataset fixed effects, using user-turn counts as denominators to reduce the mechanical opportunity effect in any-constructive-turn outcomes. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-robust standard errors. The reference category is unintentional framing, writing-oriented task, 2–3 user turns and WildChat.

| Predictor               | OR [95% CI], p value      |
|-------------------------|---------------------------|
| Intentional framing     | 2.78 [2.68, 2.88], < .001 |
| Coding task ecology     | 1.66 [1.47, 1.88], < .001 |
| 4–6 user turns          | 1.13 [1.09, 1.18], < .001 |
| 7+ user turns           | 1.06 [1.01, 1.11], .010   |
| LMSYS Chat              | 0.72 [0.69, 0.75], < .001 |
| ShareChat               | 2.30 [2.13, 2.50], < .001 |
| Conversations           | 128,569                   |
| User turns              | 491,685                   |
| Constructive user turns | 24,227                    |

**Table C14 Setting-level integrated adjacent-turn models.** Each column reports a separate within-setting logistic regression for whether the next user turn is constructive. The scaffolded-support row comes from a broad scaffolded-support model. Support-form rows come from a decomposed model that includes broad scaffolded-support presence and co-occurring M1–M6 forms, so form coefficients compare variation within scaffolded support rather than replacing the broad scaffolded-support contrast. All models include prior user state, intentional framing, assistant-turn index and recoverable model/source fixed effects where available. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors.

| Predictor               | WC coding                   | LMSYS coding                | SC coding                   | WC writing                   | LMSYS writing                 | SC writing                  |
|-------------------------|-----------------------------|-----------------------------|-----------------------------|------------------------------|-------------------------------|-----------------------------|
| A2U pairs               | 92,042                      | 71,910                      | 9,001                       | 124,214                      | 55,580                        | 5,825                       |
| Conversations           | 31,809                      | 31,606                      | 2,474                       | 39,455                       | 20,936                        | 1,537                       |
| Model/source FE count   | 13                          | 23                          | 4                           | 13                           | 23                            | 3                           |
| Scaffolded support      | 1.13 [1.08, 1.19]<br>< .001 | 1.25 [1.16, 1.34]<br>< .001 | 1.24 [1.09, 1.42]<br>.001   | 1.57 [1.38, 1.78]<br>< .001  | 1.04 [0.86, 1.26]<br>.679     | 1.34 [0.98, 1.85]<br>.068   |
| Prior user constructive | 5.20 [4.84, 5.59]<br>< .001 | 6.01 [5.43, 6.66]<br>< .001 | 4.13 [3.46, 4.93]<br>< .001 | 8.28 [6.73, 10.19]<br>< .001 | 11.89 [9.10, 15.54]<br>< .001 | 6.23 [4.43, 8.76]<br>< .001 |
| Prior user active       | 1.67 [1.57, 1.76]<br>< .001 | 1.58 [1.46, 1.71]<br>< .001 | 1.56 [1.32, 1.83]<br>< .001 | 1.91 [1.68, 2.18]<br>< .001  | 2.18 [1.80, 2.64]<br>< .001   | 1.47 [1.07, 2.01]<br>.017   |
| Prior user passive      | 2.19 [1.65, 2.92]<br>< .001 | 2.01 [1.38, 2.92]<br>< .001 | 1.82 [1.25, 2.66]<br>.002   | 2.49 [1.69, 3.65]<br>< .001  | 2.32 [1.29, 4.17]<br>.005     | 0.86 [0.40, 1.84]<br>.703   |
| Intentional framing     | 1.57 [1.49, 1.66]<br>< .001 | 1.46 [1.35, 1.58]<br>< .001 | 1.57 [1.34, 1.83]<br>< .001 | 1.02 [0.89, 1.16]<br>.798    | 0.86 [0.70, 1.06]<br>.150     | 1.37 [0.93, 2.02]<br>.107   |
| M1 feedback             | 1.19 [1.07, 1.31]<br>.001   | 1.21 [0.98, 1.50]<br>.079   | 1.28 [1.03, 1.60]<br>.025   | 0.69 [0.53, 0.91]<br>.008    | 0.81 [0.51, 1.31]<br>.397     | 1.74 [1.12, 2.71]<br>.013   |
| M2 hinting              | 1.04 [0.95, 1.14]<br>.429   | 0.90 [0.72, 1.11]<br>.332   | 0.85 [0.65, 1.10]<br>.209   | 1.52 [1.20, 1.93]<br>< .001  | 1.63 [1.09, 2.44]<br>.018     | 0.76 [0.45, 1.29]<br>.314   |
| M3 instructing          | 0.89 [0.83, 0.96]<br>.003   | 0.84 [0.69, 1.04]<br>.103   | 0.93 [0.77, 1.13]<br>.491   | 1.36 [1.04, 1.78]<br>.027    | 1.06 [0.71, 1.58]<br>.779     | 0.94 [0.45, 1.94]<br>.860   |
| M4 explaining           | 1.21 [1.09, 1.35]<br>< .001 | 1.36 [1.08, 1.72]<br>.008   | 1.77 [1.35, 2.33]<br>< .001 | 1.81 [1.40, 2.34]<br>< .001  | 1.12 [0.75, 1.68]<br>.570     | 1.09 [0.58, 2.06]<br>.789   |
| M5 modelling            | 0.86 [0.80, 0.93]<br>< .001 | 0.73 [0.63, 0.86]<br>< .001 | 0.76 [0.60, 0.96]<br>.021   | 0.69 [0.47, 1.01]<br>.055    | 1.04 [0.64, 1.68]<br>.872     | 1.57 [0.94, 2.61]<br>.084   |
| M6 questioning          | 0.73 [0.62, 0.86]<br>< .001 | 0.30 [0.21, 0.44]<br>< .001 | 0.70 [0.54, 0.90]<br>.006   | 1.13 [0.74, 1.72]<br>.567    | 0.66 [0.36, 1.21]<br>.179     | 1.66 [0.90, 3.08]<br>.106   |

**Table C15 Integrated adjacent-turn regression combining user state, assistant scaffolding and model controls.** The outcome is whether the next user turn is constructive. The scaffolded-support row is estimated in a broad scaffolded-support model; M1, M4 and M6 rows are estimated in a decomposed support-form model that includes broad scaffolded-support presence and co-occurring support forms. Estimates are odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors. The primary pooled model uses dataset fixed effects; the model/source sensitivity adds recoverable assistant model or source-family fixed effects.

| Predictor               | Pooled dataset FE         | Pooled model/source FE                | WildChat model FE                     |
|-------------------------|---------------------------|---------------------------------------|---------------------------------------|
| Scaffolded support      | 1.50 [1.44, 1.55], < .001 | 1.47 [1.41, 1.52], < .001             | 1.56 [1.49, 1.63], < .001             |
| Prior user constructive | 8.90 [8.45, 9.38], < .001 | 8.55 [8.11, 9.01], < .001             | 9.11 [8.53, 9.74], < .001             |
| Prior user active       | 2.62 [2.52, 2.72], < .001 | 2.57 [2.47, 2.67], < .001             | 2.85 [2.71, 3.00], < .001             |
| Prior user passive      | 2.08 [1.76, 2.44], < .001 | 2.07 [1.76, 2.44], < .001             | 2.32 [1.85, 2.90], < .001             |
| Intentional framing     | 1.40 [1.35, 1.46], < .001 | 1.42 [1.37, 1.48], < .001             | 1.40 [1.33, 1.47], < .001             |
| Coding task             | 1.37 [1.22, 1.54], < .001 | 1.37 [1.22, 1.54], < .001             | 1.21 [1.05, 1.40], .009               |
| M1 feedback             | 0.79 [0.73, 0.86], < .001 | 0.80 [0.74, 0.86], < .001             | 0.74 [0.67, 0.82], < .001             |
| M4 explaining           | 2.15 [1.99, 2.32], < .001 | 2.12 [1.96, 2.30], < .001             | 2.11 [1.92, 2.31], < .001             |
| M6 questioning          | 0.82 [0.72, 0.93], .002   | 0.84 [0.74, 0.96], .008               | 1.08 [0.93, 1.25], .340               |
| Model/source block      | Reference                 | LR $\chi^2_{41} = 714.4$ , $p < .001$ | LR $\chi^2_{13} = 553.2$ , $p < .001$ |

**Table C16 Incremental contribution of scaffolding features in integrated adjacent-turn models.** Likelihood-ratio block tests compare nested logistic regressions for whether the next user turn is constructive. All models include prior user state, user framing, task ecology, assistant-turn index and the indicated dataset or model/source fixed effects.

| Scope                              | Added feature block                                  | LR $\chi^2$ (df) | p value |
|------------------------------------|------------------------------------------------------|------------------|---------|
| six task settings, dataset FE      | Broad scaffolded support after user/context controls | 584.6 (1)        | < .001  |
| six task settings, dataset FE      | M1–M6 descriptors within scaffolded support          | 896.7 (6)        | < .001  |
| six task settings, dataset FE      | Full scaffolding block after user/context controls   | 1481.3 (7)       | < .001  |
| six task settings, model/source FE | Broad scaffolded support after user/context controls | 509.7 (1)        | < .001  |
| six task settings, model/source FE | M1–M6 descriptors within scaffolded support          | 829.8 (6)        | < .001  |
| six task settings, model/source FE | Full scaffolding block after user/context controls   | 1339.6 (7)       | < .001  |
| WildChat only, model FE            | Broad scaffolded support after user/context controls | 418.6 (1)        | < .001  |
| WildChat only, model FE            | M1–M6 descriptors within scaffolded support          | 519.4 (6)        | < .001  |
| WildChat only, model FE            | Full scaffolding block after user/context controls   | 937.9 (7)        | < .001  |

**Table C17 Prior user state and all six support forms jointly structure adjacent-turn constructive engagement.** The first panel reports likelihood-ratio tests adding prior-state  $\times$  M1–M6 interactions to the integrated adjacent-turn logistic model. The second panel reports pooled model/source fixed-effect estimates for each support form within each prior user state. Estimates are odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors.

| Panel                            | Scope / prior state                | Term or test               | Estimate                          |
|----------------------------------|------------------------------------|----------------------------|-----------------------------------|
| Interaction block                | six task settings, dataset FE      | Prior state $\times$ M1–M6 | LR $\chi^2_{18} = 305.4$ , < .001 |
| Interaction block                | six task settings, model/source FE | Prior state $\times$ M1–M6 | LR $\chi^2_{18} = 297.3$ , < .001 |
| Interaction block                | WildChat only, model FE            | Prior state $\times$ M1–M6 | LR $\chi^2_{18} = 186.0$ , < .001 |
| State-stratified model/source FE | prior constructive                 | M1 feedback                | 0.72 [0.64, 0.81], < .001         |
|                                  | prior constructive                 | M2 hinting                 | 0.91 [0.79, 1.04], .155           |
|                                  | prior constructive                 | M3 instructing             | 0.92 [0.81, 1.05], .236           |
|                                  | prior constructive                 | M4 explaining              | 1.72 [1.45, 2.04], < .001         |
|                                  | prior constructive                 | M5 modelling               | 0.81 [0.70, 0.93], .003           |
|                                  | prior constructive                 | M6 questioning             | 0.95 [0.72, 1.25], .694           |
|                                  | prior active                       | M1 feedback                | 0.83 [0.72, 0.95], .008           |
|                                  | prior active                       | M2 hinting                 | 1.14 [1.03, 1.27], .016           |
|                                  | prior active                       | M3 instructing             | 0.98 [0.89, 1.08], .722           |
|                                  | prior active                       | M4 explaining              | 2.13 [1.88, 2.41], < .001         |
|                                  | prior active                       | M5 modelling               | 0.83 [0.76, 0.91], < .001         |
|                                  | prior active                       | M6 questioning             | 0.75 [0.61, 0.93], .009           |
|                                  | prior passive                      | M1 feedback                | 2.16 [1.16, 4.04], .015           |
|                                  | prior passive                      | M2 hinting                 | 1.79 [0.95, 3.37], .071           |
|                                  | prior passive                      | M3 instructing             | 1.52 [0.67, 3.43], .312           |
|                                  | prior passive                      | M4 explaining              | 2.88 [1.55, 5.35], < .001         |
|                                  | prior passive                      | M5 modelling               | 0.46 [0.20, 1.04], .061           |
|                                  | prior passive                      | M6 questioning             | 2.19 [1.12, 4.27], .021           |



**Table C18 Support-intent labels are dominated by cognitive support.** Scaffolded assistant turns can receive metacognitive (I1), cognitive (I2) or affective (I3) intent labels. Scaffolded-turn prevalence reports the share of scaffolded assistant turns containing each intent label. Constructive ratios compare scaffolded conversations containing at least one assistant turn with the intent label with scaffolded conversations without that intent label, using turn-weighted constructive user-turn ratios. Intent labels can co-occur, and affective intent is sparse, so these contrasts are interpreted descriptively.


| Scope   | Intent label     | Scaffolded turns (%) | With intent (%) | Without intent (%) | Diff. (pp) | Conv. with / without |
|---------|------------------|----------------------|-----------------|--------------------|------------|----------------------|
| Pooled  | I1 metacognitive | 8.7                  | 5.0             | 6.9                | -2.0       | 10,924 / 55,892      |
| Pooled  | I2 cognitive     | 91.6                 | 6.8             | 2.8                | +4.0       | 61,942 / 4,874       |
| Pooled  | I3 affective     | 0.8                  | 12.4            | 6.4                | +6.0       | 1,052 / 65,764       |
| Coding  | I1 metacognitive | 8.2                  | 7.0             | 9.6                | -2.6       | 6,782 / 36,179       |
| Coding  | I2 cognitive     | 92.3                 | 9.4             | 3.9                | +5.6       | 40,279 / 2,682       |
| Coding  | I3 affective     | 1.0                  | 16.1            | 8.9                | +7.2       | 778 / 42,183         |
| Writing | I1 metacognitive | 9.6                  | 2.5             | 2.8                | -0.3       | 4,142 / 19,713       |
| Writing | I2 cognitive     | 90.4                 | 2.8             | 1.8                | +0.9       | 21,663 / 2,192       |
| Writing | I3 affective     | 0.6                  | 3.9             | 2.7                | +1.2       | 274 / 23,581         |

**Table C19 Support-supply profile within scaffolded assistant turns.** Values show the percentage of scaffolded assistant turns containing each scaffolding-means label. Support-form labels are not mutually exclusive, so rows do not sum to 100%. Coding-oriented scaffolded turns are dominated by explanation-heavy support, whereas writing-oriented scaffolded turns show a more mixed support profile.


| Setting           | M1<br>Feedback | M2<br>Hinting | M3<br>Instructing | M4<br>Explaining | M5<br>Modelling | M6<br>Questioning |
|-------------------|----------------|---------------|-------------------|------------------|-----------------|-------------------|
| WildChat coding   | 8.8            | 14.7          | 29.7              | 84.0             | 23.9            | 6.9               |
| LMSYS coding      | 4.4            | 7.6           | 9.7               | 76.7             | 16.5            | 13.9              |
| ShareChat coding  | 13.3           | 12.7          | 25.7              | 81.2             | 18.4            | 18.6              |
| WildChat writing  | 16.6           | 9.6           | 59.2              | 29.5             | 15.0            | 6.9               |
| LMSYS writing     | 11.3           | 10.9          | 52.3              | 26.9             | 10.3            | 13.4              |
| ShareChat writing | 24.2           | 23.2          | 22.9              | 45.2             | 23.5            | 38.1              |

## Appendix D Supporting figures

The supporting figures separate cross-dataset extensions from model-snapshot checks. Supplementary Figures D1 and D2 extend the main analyses across the six WildChat, LMSYS Chat and ShareChat task settings. Supplementary Figure D3 reports the support-form supply profile across the same six settings, and Supplementary Figure D4 provides paraphrased support-form cards for the six scaffolded-support forms. Supplementary Figures D5–D9 are a separate WildChat-only model-snapshot robustness layer. They are not substitutes for the six-setting checks; they ask whether the WildChat portion of the findings is concentrated in a particular user-facing model snapshot. LMSYS Chat and ShareChat remain represented in the six-setting figures and in the model/source-adjusted regressions, but they do not provide directly comparable user-facing model-snapshot structure for the same visual repetition. All model-snapshot figures are descriptive because model labels in naturalistic corpora are observational rather than experimentally assigned.



**Fig. D1 Coarse around-first-scaffold contrasts are exploratory boundary checks across the six task settings.** Bars show the difference in constructive engagement after versus before the first scaffolded assistant turn within conversations containing scaffolded support. Error bars show 95% conversation-bootstrap confidence intervals. These before–after contrasts are not used as treatment-effect or accumulation estimates because the first scaffolded turn may be selected by difficulty, uncertainty or an already developing exchange.



**Fig. D2 Scaffolded-support associations by user framing across the six task settings.** Stratified Poisson count-ratio estimates compare the association between scaffolded-support presence and constructive participation in intentional versus unintentional conversations. Error bars show 95% model-based confidence intervals. The figure is interpreted as descriptive moderation rather than evidence that framing causally changes support effectiveness, because user framing, task complexity and persistence are not randomly assigned.


| Support-form supply within scaffolded assistant turns (%) |      |      |      |      |      |      |
|-----------------------------------------------------------|------|------|------|------|------|------|
| WC coding                                                 | 8.8  | 14.7 | 29.7 | 84.0 | 23.9 | 6.9  |
| LMSYS coding                                              | 4.4  | 7.6  | 9.7  | 76.7 | 16.5 | 13.9 |
| SC coding                                                 | 13.3 | 12.7 | 25.7 | 81.2 | 18.4 | 18.6 |
| WC writing                                                | 16.6 | 9.6  | 59.2 | 29.5 | 15.0 | 6.9  |
| LMSYS writing                                             | 11.3 | 10.9 | 52.3 | 26.9 | 10.3 | 13.4 |
| SC writing                                                | 24.2 | 23.2 | 22.9 | 45.2 | 23.5 | 38.1 |
|                                                           | M1   | M2   | M3   | M4   | M5   | M6   |

WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat. Support-form labels are non-exclusive.


**Fig. D3 Support-form supply within scaffolded assistant turns across the six task settings.** Cells report the percentage of scaffolded assistant turns containing each non-exclusive support-form label. Rows do not sum to 100% because a single scaffolded turn can contain multiple support forms. The figure provides task-context background for the support-form association contrasts in Fig. 4.

|                                                                                                                                                                                                                                                                                                                                    |                                                                                                                                                                                                                                                                                                                                        |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <p><b>M1 feedback</b><br/><i>Evaluates a user's prior attempt</i></p> <p>"Your diagnosis is close: the race is not in the callback itself, but in when the cache key is reused. Check that key before changing the loop."</p> <p><b>Role:</b> Returns the user's attempt as something to repair, revise or extend.</p>             | <p><b>M2 hinting</b><br/><i>Offers a partial cue</i></p> <p>"Before rewriting the parser, inspect the first malformed token and compare it with the token immediately before it."</p> <p><b>Role:</b> Leaves solution work open while directing attention to a useful diagnostic path.</p>                                             |
| <p><b>M3 instructing</b><br/><i>Provides procedural steps</i></p> <p>"First isolate the failing branch, then add a regression test, and only then refactor the duplicated condition."</p> <p><b>Role:</b> Organizes the user's next actions through explicit sequencing or strategy.</p>                                           | <p><b>M4 explaining</b><br/><i>Makes a mechanism visible</i></p> <p>"The exception occurs because the event loop is already running; a nested call asks for a resource that cannot be acquired twice."</p> <p><b>Role:</b> Provides a rationale that can be tested, adapted or transferred.</p>                                        |
| <p><b>M5 modelling</b><br/><i>Shows an adaptable pattern</i></p> <p>"A reusable structure is: validate input, normalize fields, then pass the cleaned object to the writer. You can adapt this skeleton to your case."</p> <p><b>Role:</b> Supplies an example or pattern intended for adaptation rather than only completion.</p> | <p><b>M6 questioning</b><br/><i>Elicits a constraint or reflection</i></p> <p>"Which constraint matters more here: preserving the author's tone or shortening the paragraph? That choice should determine the revision strategy."</p> <p><b>Role:</b> Prompts the user to specify priorities, constraints or next reasoning steps.</p> |


**Fig. D4 Representative support-form cards for the six scaffolded-support forms.** Cards show de-identified, paraphrased assistant replies that illustrate the annotation decision rule for each support-form label rather than reproducing raw conversation text. Support-form labels are non-exclusive, so a single assistant turn can receive more than one form label.




**Fig. D5 WildChat model-snapshot prevalence profiles preserve the main domain contrast.** Across user-facing WildChat model snapshots, coding-oriented conversations generally show higher cognitive engagement and higher assistant scaffolding rates than writing-oriented conversations. This is a WildChat-only model-snapshot check; cross-dataset prevalence and association checks are reported separately across all six settings. The model labels are observational and may reflect time, user mix and traffic composition, so the figure is not a model-causal comparison.



**Fig. D6 Intentionality gaps are visible across WildChat model snapshots.** The figure reports model-specific differences between intentional and unintentional conversations in cognitive and constructive engagement, separately for WildChat coding and writing. It complements the six-setting user-framing analyses by checking whether the WildChat pattern is concentrated in a particular model snapshot.




**Fig. D7 Conversation-level support associations are broadly positive across WildChat model snapshots.** The figure repeats major conversation-level analyses by user-facing WildChat model snapshot, including constructive participation, adjusted support associations and exploratory coarse before–after contrasts around the first scaffolded assistant turn. It is a within-WildChat robustness check; cross-dataset scaffolded-support contrasts are reported across all six settings in the supporting tables and main figures.



**Fig. D8 Local temporal signals vary across WildChat model snapshots.** The figure reports adjacent-turn constructive lifts and emotion-to-support responsiveness by WildChat model snapshot. Adjacent-turn constructive lifts are often positive, especially in coding-oriented conversations, but the magnitude varies across snapshots; emotion-to-support responsiveness is more heterogeneous and is treated as exploratory.





**Fig. D9 Coarse WildChat model-family aggregation preserves the main directional patterns.** WildChat model snapshots are grouped into coarse model families to summarize cognitive engagement prevalence, adjusted support associations and adjacent-turn constructive lift. Coding remains higher than writing across families, and the support-association and adjacent-turn lift patterns remain directionally consistent on average.