

Informal learning in everyday human–LLM interaction

Zixin Chen¹, Haotian Li^{2*}, Ziang Xiao³, Huamin Qu¹, Xing Xie²

¹The Hong Kong University of Science and Technology, Hong Kong, China.

²Microsoft Research Asia, Beijing, China.

³Johns Hopkins University, Baltimore, 21218, Maryland, United States.

*Corresponding author(s). E-mail(s): haotian.li@microsoft.com;
Contributing authors: zchendf@connect.ust.hk; ziang.xiao@jhu.edu;
huamin@cse.ust.hk; xing.xie@microsoft.com;

Abstract

As LLMs become increasingly capable of completing tasks for users, a central concern is that working with them may erode the opportunities through which people develop their own capabilities. We analyse large scale human–LLM conversations to surface how users engage with LLMs in learning-oriented ways, revealing the prevalence and interactional mechanisms through which informal learning may occur in everyday AI use. Across 128,569 naturalistic conversations, we translated learning-science constructs into turn-level behavioural signatures. Cognitive engagement—visible uptake of assistant output—appeared in 31.9% of 491,685 user turns, whereas constructive engagement, the deepest observable form of learning-oriented engagement, appeared in 4.9%, showing that deeper sense-making was recurrent but selective. Our study further identifies factors associated with these forms of engagement. Scaffolded assistant support, present in 31.7% of assistant turns, consistently marked richer constructive participation, with associations varying by user framing, task ecology, support form, timing and prior user state. These findings make behavioural signatures of informal learning observable at scale and shift AI evaluation: from answer-delivery efficiency toward the preservation of cognitive opportunities for users to reason, test ideas and construct understanding in the course of everyday problem-solving.

Keywords: human–AI interaction, informal learning, large language models, scaffolding, learner engagement

1 Introduction

Large language models (LLMs) are becoming embedded in the cognitive work of everyday life: people increasingly turn to them to write, code, search and reason through unfamiliar problems [1–3]. Much of human competence, however, is acquired precisely by doing such work: reasoning through a problem, attempting a solution, reworking it until it holds [4, 5]. A growing body of evidence warns that delegating that effort to machines can erode the very capabilities the work would otherwise build [6, 7]. Whether working with an assistant diminishes *users’ own cognitive work* or instead becomes an occasion for thought depends not on the technology but on how people engage with it: they may take an answer and move on, or use it to test their own ideas, seek out explanations and pursue a question to its root [8, 9]. Everyday human–LLM interaction has thus become a vast and mostly undesigned setting for informal learning—and, as AI does more of this work, an increasingly consequential one for whether people keep learning and thinking for themselves [10–12].

Yet this informal learning has gone largely unexamined. Existing research on LLMs and learning has focused instead on settings where learning support is deliberately designed, such as automated feedback, question generation, retrieval practice and tutoring systems [13–17]. These studies show that LLMs can enact instructional strategies when learning is the explicit design target [18, 19]; they tell us less about general-purpose LLM use, where users usually seek task progress rather than instruction. Everyday use is also harder to study: it follows no curriculum and yields no test score, so the learning at stake is easy to suspect but hard to see. Learning science offers a way to detect it: deeper forms of overt engagement are associated with stronger learning outcomes, so engagement can serve as an observable process indicator in informal, user-generated interactions whose goals and endpoints are not pre-specified by curricula or study protocols [20–22]. We therefore ask whether learning-oriented engagement is visible in everyday human–LLM interaction, where its deepest constructive form appears, and how assistant scaffolding is coupled with further constructive participation.

We analyse 128,569 naturalistic conversations from three public large-scale conversational corpora: WildChat, LMSYS Chat and ShareChat [23–25]. Across these corpora, we focus on coding and writing as two common task ecologies in which LLMs are used to produce, revise and troubleshoot knowledge work [1, 8, 26], yielding 491,685 user turns and 489,785 assistant turns. Our analyses move from engagement prevalence, to the contextual organization of constructive engagement, to how assistant scaffolding is associated with user participation at conversation-level, support-form and adjacent-turn scales. To make these questions measurable in naturalistic logs, we translate theories of engagement and scaffolding into role-specific, turn-level behavioural signature. User turns are coded by the depth of visible uptake of assistant output: *passive* receipt, *active* use or follow-up, and *constructive* engagement, where users elaborate, test, revise or extend ideas in the exchange [20, 21]. Assistant turns are coded by whether they primarily deliver the requested content or provide scaffolded support: contingent assistance that diagnoses difficulties, decomposes tasks, explains rationales, gives feedback or hints, or prompts reflection while leaving the user to continue the work [27–30]. Together, these signatures allow us to

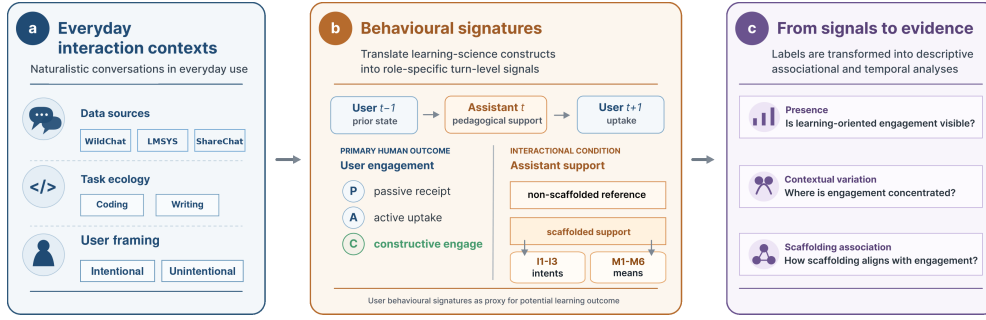


Fig. 1 Analytic framework for studying informal learning in everyday human-LLM interaction. **a**, Public conversations are situated by data source, task ecology and user framing. **b**, Engagement and scaffolding constructs are translated into role-specific turn-level labels for user engagement and assistant support, with assistant support further characterized by support intent (I1-I3: metacognitive, cognitive and affective) and support form (M1-M6: feedback, hinting, instructing, explaining, modelling and questioning); Supplementary Table C1 provides plain-language examples of each label. **c**, These labels support analyses of engagement presence, contextual organization and support-engagement coupling at conversation-level, support-form and adjacent-turn scales. Labels indicate observable learning-oriented participation, not durable learning outcomes or causal instructional effects.

examine where users remain actively involved in reasoning with generated output, and how constructive participation is organized by task context, assistant support and local turn ordering.

Across the three corpora, everyday human-LLM interaction emerged as a genuine but selective and conditional site of learning. First, learning-oriented engagement was recurrent in everyday human-LLM interaction, but its deepest constructive form was selective: cognitive engagement appeared in nearly one third of user turns, whereas constructive engagement appeared in about five percent. Second, constructive engagement was ecologically organized: it was more common under explicit learning-oriented user framing, varied across task ecologies and became most visible in sustained exchanges where users returned to, interrogated and reshaped the assistant’s response. Third, assistant scaffolding marked richer constructive participation, but not as a uniform exposure. Its association depended on how support was positioned within the exchange: feedback-like and explanatory forms aligned most clearly with constructive engagement at the conversation level, while adjacent-turn analyses showed that support-engagement coupling in local sequences was shaped by prior user state and support form. Together, these findings locate everyday human-LLM interaction in a middle space between ordinary tool use and designed instruction. These results offer an empirical counterpoint to fears that AI use simply erodes human thinking: people have not ceded their thinking to the model but continue to engage constructively, if selectively. They suggest that the educational value of future AI assistants lies not in making every interaction tutor-like, but in calibrating when to answer, explain, give feedback or step back, preserving the cognitive agency through which people question, test, revise and build judgment while solving real problems.

Table 1 Dataset and behavioural overview. Summary statistics for the WildChat, LMSYS Chat and ShareChat task settings analysed in the study. Constructive engagement is reported as the deepest observable level within cognitive engagement and serves as the focal user-side measure in subsequent analyses.

Setting	Sample size			User framing	Assistant support	Cognitive engagement		Emotional engagement
	Conversations	User turns	Assistant turns	Intentional conv. (%)	Scaffolded turns (%)	Overall (%)	Constructive level (%)	Emotional (%)
WildChat coding	31,878	124,073	124,623	37.13	51.03	50.21	9.23	0.95
WildChat writing	39,534	163,705	164,824	25.80	23.53	16.34	2.22	0.57
LMSYS coding	32,114	107,065	105,174	41.49	29.07	45.26	5.44	0.91
LMSYS writing	21,023	77,906	76,279	20.64	19.43	15.30	1.53	0.21
ShareChat coding	2,481	11,541	11,522	45.63	50.67	45.00	15.17	3.38
ShareChat writing	1,539	7,395	7,363	29.69	22.42	33.05	5.18	0.34
Total / pooled	128,569	491,685	489,785	32.12	31.70	31.94	4.93	0.74

Note. User framing is reported as the percentage of conversations with explicit learning-oriented framing. Cognitive, constructive and emotional engagement are percentages of user turns; scaffolded support is computed over assistant turns. The constructive column uses all user turns as the denominator, whereas Fig. 2a decomposes cognitively engaged user turns into passive, active and constructive levels.

2 Results

The results follow the analytic sequence in Fig. 1: engagement prevalence, contextual organization and support–engagement coupling. We first quantify cognitive and constructive engagement in everyday human–LLM conversations, then locate constructive engagement across user framing, task ecology and interaction depth. We then turn to assistant support, asking whether scaffolded conversations contain richer constructive participation, which support forms carry this association and how support is ordered with prior and subsequent user engagement at the adjacent-turn scale.

2.1 Everyday human–LLM conversations contain learning-oriented engagement

Learning-oriented engagement was visible across all six task settings. Across 491,685 user turns, 31.9% showed cognitive engagement and 4.9% reached constructive engagement, the deepest observable user-side signal in the annotation scheme (Table 1). **Explicit affective expression was much less visible, (0.74% of user turns), consistent with the coding and writing focus of the sample.** Thus, everyday human–LLM conversations often contained observable user uptake of assistant responses, but deeper constructive sense-making remained selective rather than routine.

Figure 2a shows how cognitively engaged turns were composed. Active uptake was the dominant form in every setting, accounting for 62.0–86.4% of cognitively engaged turns, whereas constructive engagement accounted for 10.0–33.7% and passive receipt for 1.2–14.6%. This composition separates broad engagement from deeper sense-making. Cognitive engagement shows that everyday LLM use often involves users working with assistant output; active uptake captures use, application or follow-up; and constructive engagement identifies the cases in which users elaborate, test, revise or extend ideas. The remaining analyses therefore focus on where constructive engagement appears and how it is associated with user framing, task ecology, interaction depth and assistant support.

2.2 Constructive engagement is patterned by user framing, task ecology and sustained exchange

Having established that everyday human–LLM conversations contain learning-oriented engagement, we next examined where its constructive form became visible. Explicit learning-oriented framing provided the first contrast (Fig. 2b). In every setting, conversations framed around learning, understanding or exploration had higher cognitive and constructive user-turn rates than conversations without such framing. The cognitive-engagement lift ranged from +17.7 to +47.3 percentage points across settings, and the constructive-engagement lift ranged from +2.5 to +12.8 points. Yet constructive participation was also present in conversations without explicit learning-oriented framing. User framing therefore amplified constructive engagement, but did not define it: learning-relevant participation also emerged in task-oriented exchanges where learning was incidental to solving the problem at hand.

Task ecology provided a second source of variation. Across all three corpora, coding-oriented conversations showed higher cognitive and constructive engagement than writing-oriented conversations. Cognitive engagement ranged from 45.00–50.21% in coding conversations, compared with 15.30–33.05% in writing conversations; constructive engagement ranged from 5.44–15.17% in coding and 1.53–5.18% in writing. A conversation-level context model showed the same direction after including user framing, conversation-length bucket and dataset fixed effects, with coding task ecology positively associated with the presence of at least one constructive user turn (OR=1.89, 95% CI 1.71–2.08, $p < .001$; Supplementary Table C12).

The coding–writing contrast points to task affordances that organize baseline opportunities for visible constructive participation. Coding tasks often externalize uncertainty through errors, constraints and executable tests, creating occasions for users to diagnose failures, refine assumptions, compare alternatives and test candidate solutions. Writing tasks also involve substantial revision and evaluation, but the interaction can often proceed through accepting, adapting or requesting a draft revision without requiring the user to articulate the rationale for the change [31–33]. The observed task difference is therefore consistent with the ecological interpretation: constructive engagement is more likely to appear when the task makes uncertainty, evaluation and revision explicit in the exchange.

The distinction between broad cognitive engagement and constructive engagement reinforces this task-ecology account. LMSYS Chat coding illustrates this distinction: it had high cognitive engagement, yet its cognitively engaged turns were dominated by active uptake rather than constructive elaboration (Fig. 2a). This pattern illustrates that users can work actively with assistant output—applying code, requesting repairs or asking for the next step—without making their reasoning, tests or revisions explicit enough to surface as constructive engagement. Taken together with the coding–writing contrast, this shows that task ecology patterned not only how often users worked with assistant output, but also whether that work became visible as articulated questioning, testing, refinement or extension.

Sustained exchange provided the third context in which constructive participation became visible (Fig. 2c). We operationalized sustained exchange with conversation-length buckets, asking whether a conversation contained at least one constructive

user turn. Across all six settings, conversations with more user turns were more likely to contain at least one constructive user turn. The gradient was steeper in coding-oriented settings, but the same direction appeared in writing-oriented settings. This pattern indicates that constructive engagement is most visible when everyday LLM use develops beyond a one-off request for an answer. In such exchanges, users have room to return to the assistant’s response, introduce new constraints, test implications, request revisions and reshape their understanding. Because this panel uses an any-event outcome, we treat the length gradient as a visibility pattern; a grouped-binomial constructive-rate sensitivity attenuated but preserved the positive length association (4–6 user turns: OR=1.13; 7+ user turns: OR=1.06; Supplementary Table C13).

Together, these results show that constructive engagement was ecologically organized rather than evenly distributed across everyday use. It was amplified by explicit learning-oriented framing, patterned by task affordances that create occasions for testing and revision, and most visible in sustained exchanges. A conversation-level context model with user framing, task ecology, conversation-length bucket and dataset fixed effects supported the same organization, and the rate sensitivity showed that the length association was attenuated but preserved after accounting for user-turn denominators (Supplementary Tables C12 and C13). Having located these user-side and contextual patterns, we next ask whether assistant scaffolding marks conversations with richer constructive participation.

2.3 Scaffolded support marks richer constructive participation

Having located the user-side and contextual conditions under which constructive engagement became visible, we next turned to assistant-side scaffolding: support that structures the user’s process rather than simply returning a deliverable, leaving room for diagnosis, repair or evaluation work [27–30].

Scaffolded support appeared in 31.7% of assistant turns (Table 1), making it a common interactional condition rather than a rare pedagogical exception. At the conversation level, conversations containing at least one scaffolded assistant turn had higher turn-weighted constructive user-turn ratios than reference conversations without scaffolded support in all six task settings (Fig. 3a). The constructive-ratio differences ranged from +0.99 to +7.34 percentage points across WildChat, LMSYS Chat and ShareChat, and all six contrasts were distinguished from zero in conversation-level bootstrap tests ($p < .001$; Supplementary Table C10). Scaffolded conversations also showed more turns after the first assistant response in every setting (+1.42 to +3.38 turns; Fig. 3d), indicating that scaffolded support marked exchanges with more sustained participation. Together, these descriptive contrasts show that scaffolded support co-occurred with richer constructive participation across corpora and task settings.

The association remained positive after accounting for observed user framing and interactional context in covariate-adjusted models. In setting-specific models, scaffolded-support presence was associated with higher constructive-turn counts, with Poisson count ratios ranging from 1.569 to 2.491, and with higher odds that a conversation contained at least one constructive user turn, with logistic odds ratios ranging from 1.437 to 2.140 (Fig. 3b). The same positive direction appeared within both




Fig. 2 Learning-oriented engagement varies by engagement form, user framing and conversation length. **a**, Composition of cognitively engaged user turns by passive, active and constructive levels across the six task settings. **b**, Cognitive and constructive user-turn rates in conversations with versus without explicit learning-oriented framing; annotations show the framed-minus-unframed difference in percentage points. **c**, Share of conversations containing at least one constructive user turn by conversation-length bucket. Panel a uses cognitively engaged user turns as the denominator; panel b reports percentages of user turns; panel c reports conversation-level shares. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.

explicitly learning-oriented and non-explicitly framed conversations (Fig. 3c), indicating that the support–engagement association was not limited to interactions users announced as learning episodes. A rate sensitivity treating user turns as the exposure preserved positive scaffolded-support associations in all six settings (rate ratios 1.356–1.703; all $p < .001$; Supplementary Table C11). Across these checks, scaffolded support remained a consistent conversation-level factor of richer constructive participation. This conversation-level signal motivates a finer analysis of the support forms through which scaffolded support was expressed.

2.4 Support form differentiates the scaffolding–engagement association

The conversation-level analyses established scaffolded support as a consistent factor of richer constructive participation. We next asked which features of scaffolded support carried this signal. For each scaffolded assistant turn, the annotation scheme recorded two parallel descriptors: a support-intent label indicating whether the turn primarily targeted metacognitive, cognitive or affective support (I1–I3), and one or more support-form labels indicating how the support was delivered—feedback, hinting, instructing, explaining, modelling or questioning (M1–M6). This two-axis scheme




Fig. 3 Scaffolded support marks richer constructive participation across task settings. **a**, Turn-weighted constructive user-turn ratios in conversations containing scaffolded support versus reference conversations without scaffolded support. **b**, Covariate-adjusted within-setting associations between scaffolded-support presence and constructive participation, reported as Poisson count ratios and logistic odds ratios. **c**, Poisson count ratios for scaffolded-support presence stratified by explicit versus non-explicit learning-oriented framing. **d**, Difference in the number of turns after the first assistant response between scaffolded and reference conversations. Numbers indicate scaffolded-minus-reference differences, except panels b and c, which report model estimates. Error bars in panels b and c show model-based 95% confidence intervals; bootstrap confidence intervals for the descriptive contrasts in panels a and d are reported in Supplementary Table C10. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.


follows scaffolding accounts that distinguish the purpose of assistance from the means through which assistance is provided [27–30]. Support intent offered limited contrast because scaffolded turns were predominantly cognitive in purpose (Supplementary Table C18). We therefore focused on the form axis, comparing scaffolded conversations containing each support form with scaffolded conversations without that form and reporting turn-weighted constructive user-turn ratio differences (Fig. 4a,b).

Support form sharply differentiated the scaffolded-support association. Feedback-like support (M1) showed the largest constructive-engagement contrast in both explicitly framed and non-explicitly framed conversations, with the association strongest under explicit learning-oriented framing (+14.7 versus +6.7 percentage points; framing difference, +8.0 points; Fig. 4a). Explaining (M4) showed the same direction (+6.7 versus +3.0 points; framing difference, +3.6 points), and hinting (M2) was positive but smaller. Instructing (M3) and questioning (M6), by contrast, were lower or negative, especially in explicitly framed conversations. Task-stratified estimates gave a complementary view (Fig. 4b): feedback-like support (M1) and explaining (M4) were

Table 2 Summary of key support–engagement associations across the three corpora. Constructive-ratio contrasts compare conversations with versus without scaffolded support using turn-weighted constructive user-turn ratios. Adjusted models report the association of scaffolded-support presence with constructive-turn count (Poisson count ratio) and with having at least one constructive turn (logit odds ratio). Adjacent-turn lift is the difference in the probability that the next user turn is constructive after scaffolded versus non-scaffolded reference assistant turns.

Setting	Constructive ratio Has Scaf. / No Scaf.	Difference (pp)	Depth diff. (turns)	Poisson count ratio	Logit OR	Adjacent lift (pp)
WildChat coding	0.101 / 0.056	+4.6	+2.23	1.852	1.761	+2.68
LMSYS coding	0.065 / 0.040	+2.5	+1.42	1.622	1.544	+2.01
ShareChat coding	0.163 / 0.090	+7.3	+3.36	1.893	2.140	+6.21
WildChat writing	0.027 / 0.017	+1.0	+3.13	1.569	1.437	+1.13
LMSYS writing	0.022 / 0.011	+1.1	+2.10	1.985	1.764	+0.27
ShareChat writing	0.066 / 0.034	+3.2	+3.38	2.491	1.757	+3.35

positive in both coding- and writing-oriented task ecologies, whereas instructing (M3), modelling (M5) and questioning (M6) varied more strongly by task.



Support-form labels are non-exclusive; q values are Benjamini-Hochberg FDR-adjusted.

Fig. 4 Support forms differentiate scaffolded interaction. **a**, Constructive-engagement differences associated with each support form among scaffolded conversations, stratified by intentional versus unintentional user framing. **b**, The same support-form contrasts stratified by coding- versus writing-oriented task ecology. Bars are turn-weighted constructive user-turn ratio differences, comparing scaffolded conversations containing a support form with scaffolded conversations without that form; error bars show 95% confidence intervals, and bracket annotations show between-stratum differences and Benjamini–Hochberg FDR-adjusted q values within each panel family. Support-form labels correspond to M1–M6 and are not mutually exclusive.

The form-level pattern is distinctive because these support moves occurred inside general-purpose task assistance rather than within a designed lesson. In formal instruction, feedback, explanation and direct guidance are usually embedded in planned sequences with explicit learning goals. In everyday LLM use, the same forms are interleaved with task completion, so an assistant response can either close the problem by supplying the next step or leave part of the problem available for the user

to inspect, revise or test. Feedback-like support (M1) was strongly aligned with constructive engagement because it takes the user’s prior attempt as the object of the response, making a mismatch, partial success or error available for evaluation, repair or extension [34–36]. Explaining (M4) similarly opens the assistant’s answer by exposing mechanisms or rationales that users can test, adapt or connect to their own understanding [37, 38]. Hinting (M2) can leave solution work open, but its association was smaller. Instructing (M3), modelling (M5) and questioning (M6) can be useful for task progress, examples or clarification; in these logs, however, they were less consistently aligned with constructive elaboration. This is consistent with assistance-dilemma and scaffolding accounts in which highly executable steps, models to imitate or information-management moves can reduce the need for learners to generate, evaluate or revise the next idea [22, 39, 40]. Thus, in everyday human–LLM interaction, support form is informative not only as a pedagogical category but as an interactional position: it indicates whether assistance turns the answer into material for the user’s next cognitive move or completes the deliverable with little further reasoning visible.

The support-form supply profile adds task context to these associations (Supplementary Fig. D3). Explaining (M4) dominated scaffolded coding turns across the three corpora, whereas writing-oriented support drew more evenly on instructing (M3), explaining (M4) and questioning (M6). This support ecology clarifies the interactional environment in which the form-level associations arose: coding and writing differed both in which support forms aligned with constructive engagement and in which scaffolding repertoires users typically encountered.

Together, these analyses refine the conversation-level support result. Scaffolded support marked richer constructive participation, and this association was carried most clearly by particular support forms within task-specific scaffolding repertoires. The form-level result makes the temporal question sharper: if support matters partly through how it responds to the user’s ongoing work, then its adjacent-turn association with constructive engagement should depend on the user’s immediately preceding state. We therefore next examine support–engagement coupling at the adjacent-turn scale.

2.5 Support–engagement coupling is state- and form-dependent

The preceding analyses established two things: scaffolded conversations contained richer constructive participation, and this association was carried unevenly by support form. These results identify where assistant support is associated with constructive engagement, but not the interactional scale at which the association appears. We therefore turned to adjacent-turn ordering, analysing how the user’s current state, the assistant’s support form and the immediately following user turn were linked.

At this adjacent-turn scale, scaffolded assistant turns were more likely than reference assistant turns to be followed by constructive engagement in all six task settings (Fig. 5a). The next-turn constructive lift was positive in every setting, ranging from +0.27 to +6.21 percentage points. Five of the six adjacent-turn lifts were statistically distinguishable from zero in conversation-cluster bootstrap tests; LMSYS Chat writing was positive but weaker (+0.27 percentage points, 95% CI -0.01 to +0.56, $p = .061$;




Fig. 5 Support–engagement coupling is state- and form-dependent across task settings. **a**, Next-turn constructive lift after scaffolded versus reference assistant turns. **b**, Probability of a constructive next user turn by prior user state and assistant support condition; annotations show scaffolded-minus-reference differences. **c**, Probability that the next assistant turn contains scaffolded support by prior user state. **d**, Pooled model/source fixed-effect odds ratios for focal support forms within prior user states. M1 feedback and M4 explaining are shown because Section 2.4 identified them as focal positive conversation-level forms; full M1–M6 interaction estimates and representative support-form cards are reported in Supplementary Table C17 and Supplementary Fig. D4. Error bars show 95% confidence intervals where displayed. Panel d reports adjacent-turn logistic models with conversation-cluster robust standard errors. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.

Supplementary Table C10). The temporal signal was therefore consistently positive across settings, but heterogeneous in magnitude.

This adjacent-turn coupling was organized by the user’s immediately preceding state. State-conditioned contrasts were generally positive but varied across prior constructive, active and passive turns (Fig. 5b). The user-to-assistant direction showed the same contingency: constructive and active user turns were more often followed by scaffolded support than passive turns, especially in coding-oriented conversations (Fig. 5c). These two directions suggest a coupled sequence rather than a one-way support pattern: engaged user states more often preceded scaffolded support, and scaffolded support was then more often followed by further constructive participation.

We then asked whether the form of support shaped this adjacent-turn coupling once prior user state and adjacent-turn context were modelled jointly. We fitted integrated logistic regressions predicting whether the next user turn was constructive,

combining prior user state, assistant-scaffolding descriptors, task and framing variables, turn position and model/source fixed effects. The scaffolding block improved fit beyond user state and context (full scaffolding block LR $\chi^2_7 = 1339.6$, $p < .001$; Supplementary Table C16), and prior-state $\times$ support-form interactions improved fit further (pooled model/source fixed effects: LR $\chi^2_{18} = 297.3$, $p < .001$; Supplementary Table C17). These tests show that the immediate support–engagement association was not reducible to the user’s preceding state alone; it depended jointly on what the user had just made visible and how the assistant supported the next step.

Figure 5d highlights feedback-like support and explaining, the two support forms that carried the clearest positive signal in the conversation-level analysis; complete M1–M6 interaction estimates and paraphrased support-form cards are reported in the Appendix (Supplementary Table C17; Supplementary Fig. D4). Explaining provided the more stable adjacent-turn signal, predicting constructive follow-up after prior constructive, active and passive user turns (ORs 1.72, 2.12 and 2.88, respectively; all $p < .001$). Feedback showed a different immediate role: its positive next-turn association was concentrated after prior passive turns, rather than after turns in which users were already active or constructive. This divergence clarifies the conversation-level result: feedback marked exchanges in which user attempts were available for evaluation and repair, whereas explanation more consistently provided material for the user’s immediate next move. Support form therefore mattered through its position in the unfolding user–assistant sequence, not as a context-free pedagogical category.

Taken together, the temporal analyses locate support–engagement coupling in adjacent-turn sequences. Constructive participation was most clearly organized where the user’s visible state, the assistant’s support condition and the form of support were considered together. Learning-oriented engagement in everyday human–LLM interaction is therefore assembled in situated next-turn sequences in which the user’s visible work and the assistant’s support form jointly shape whether constructive participation continues.

3 Discussion

Everyday human–LLM interaction as an informal learning ecology

By tracing engagement in large-scale public logs, this study shows that everyday human–LLM interaction can be analysed as an empirical ecology of informal learning. These conversations are organized around users’ own tasks, yet they contain behavioural traces of learning-oriented work. Section 2.1 establishes the depth structure of this participation: cognitive engagement appeared in nearly one third of user turns, whereas constructive engagement appeared in about five percent. This distinction matters because broad uptake and constructive sense-making mark different relationships to AI output. Uptake shows that users incorporate, apply or extend an assistant response in the task at hand; constructive engagement marks the more selective moments in which users make their reasoning visible by articulating constraints, testing implications, revising assumptions or extending explanations. In this

way, informal learning becomes observable not as a declared intention, but as a pattern of participation within ordinary interaction.

The contextual results in Section 2.2 show how this participation was organized. Explicit learning-oriented framing made constructive engagement more common, but task-oriented exchanges also contained constructive turns. Task ecology shaped which forms of sense-making became visible: coding conversations often exposed failure, constraints and tests, whereas writing conversations made revision, adaptation and evaluative choice visible when users brought those concerns into the exchange. Sustained conversations added temporal opportunity, allowing earlier responses to become points of return, comparison and revision. Together, these findings define everyday human–LLM interaction as an informal learning ecology: user orientation, task affordances and interaction depth shape when ordinary problem solving becomes an occasion to build, test and extend understanding.

Scaffolding as situated support

The assistant-side findings show how support enters this ecology. As shown in Section 2.3, conversations containing scaffolded support had richer constructive participation across all six task settings, and this association remained positive after adjustment for observed conversation-level context. This pattern is most informative when scaffolded support is read as a relation between the assistant response and the user’s next possible action. In everyday task-oriented LLM use, an answer can do more than provide a usable deliverable: it can expose a rationale, error, mechanism, constraint or partial solution that remains available for inspection. This view extends scaffolding theory, which emphasizes contingent assistance and the gradual transfer of responsibility [27, 28, 40], into general-purpose AI interaction, where task completion and opportunities for sense-making are interleaved within the same exchange.

The support-form results specify what this situated support looked like in the logs. As shown in Section 2.4, feedback-like support and explanation were most consistently aligned with constructive engagement. Feedback makes a prior attempt evaluable: it turns a user’s partial solution, diagnosis or draft into something that can be repaired, refined or extended [34, 35]. Explanation makes reasons and mechanisms inspectable, giving users material for comparison, self-explanation and transfer [37, 38]. Other forms, including instruction, modelling and questioning, supported task progress, examples and clarification, with more task-contingent alignment to visible constructive elaboration. The form-level contrasts point to the warrant that support attaches to generated content: some forms organize action, whereas others make visible the grounds on which an answer, attempt or mechanism can be used. For general-purpose AI assistance, this shifts scaffolding from the surface form of a response to the epistemic status of the output it produces: fluent answers become learning-relevant when they are also judgeable answers, with grounds, limits or failure points that can be inspected and adapted within the task.

Learning continuity across formal and everyday contexts

The temporal findings clarify what everyday AI assistance can add to a broader learning system. As shown in Section 2.5, support–engagement coupling was most interpretable in adjacent user–assistant–user sequences, where the user’s preceding state and the assistant’s support form shaped whether constructive participation continued. This scale gives everyday AI use a role that differs from formal instruction. Formal learning environments organize progression through curricula, sequencing, assessment, shared standards and opportunities to abstract beyond the immediate task. Everyday AI assistance organizes re-entry: it brings explanations, feedback and prior reasoning back into the moments where people apply knowledge under practical constraints.

This distinction matters because much learning depends on whether knowledge can be reactivated when it is needed. A concept explained in a course, tutorial or earlier exchange becomes educationally consequential when it can be recognized again in a new problem, adapted to a different constraint or used to diagnose a breakdown. Human–LLM exchanges provide occasions for that re-entry. A user encounters a concrete difficulty; the assistant supplies a rationale, comparison or correction; the user can then relate that material to the task at hand. The value of such episodes lies in making knowledge usable within action, where understanding is tested against the demands of a real problem.

This complementarity suggests a different role for personal AI assistants in education. Schools, courses and curricula provide progression, accountability and shared criteria. AI assistants can provide continuity across contexts: carrying explanations from one task to another, preserving traces of prior attempts, helping users notice recurring uncertainties and supporting consolidation after repeated encounters with a concept. In this role, lifelong AI assistance can function as a connective layer between structured learning and everyday practice, giving memory and personalization to the informal learning that emerges while people work.

AI-mediated work and the social organization of expertise

The broader implication is that AI-mediated work may reshape the pathways through which expertise is formed. Much professional learning develops through participation in the practices of a domain: proposing first attempts, encountering error, receiving feedback, revising choices, explaining decisions and seeing consequences unfold [41, 42]. These activities are often treated as parts of production, yet they also carry an apprenticeship function. The results provide an empirical entry point into this process in AI-mediated settings: learning-oriented participation was measurable, selective and structured by the conditions of interaction. As LLMs become embedded in writing, coding, search and problem solving, the formation of expertise becomes tied to how these environments organize participation in the practices through which judgment is built.

AI also changes the repertoire of capacities through which people enter expert practice. When systems can generate fluent drafts, code, explanations and recommendations, expertise may depend increasingly on formulating good problems, recognizing

relevant principles, interrogating generated solutions, judging evidence, revising under constraints and remaining accountable for consequences. These capacities extend beyond generic AI literacy; they are forms of situated judgment that develop through repeated participation in domain work. Schools, workplaces and platforms will therefore shape future expertise through the routines they build around AI: what novices practise, where feedback enters, how responsibility is assigned and how mistakes become occasions for development.

A science of everyday AI learning should therefore study cognitive apprenticeship as a distributed property of sociotechnical systems. The outcomes of interest extend beyond immediate productivity or single-session engagement to the formation of future practitioners: whether repeated AI-mediated routines cultivate judgment, autonomy and domain competence over time. The social stakes are distributional. Intelligent assistance can broaden access to expert-like feedback, explanation and revision, but the developmental value of that access will depend on institutional design: who participates in diagnosis, who practises revision, who observes consequences and who is asked to justify decisions. The long-term educational significance of AI will depend on how societies organize these pathways into expertise, and on whether AI-mediated work sustains broad participation in the practices through which judgment, autonomy and domain competence are formed.

Limitations

Several boundaries define the interpretation of these findings. The analyses use public conversational logs and characterize population-level interactional patterns rather than longitudinal learner histories. They do not observe users' prior knowledge, motivation, goals outside the conversation or later learning outcomes. The engagement labels capture behavioural manifestations of learning-oriented participation, not direct evidence of retention, transfer or skill development. Assistant support was also observed rather than assigned, so the associations reported here should be interpreted as structured patterns in naturalistic interaction rather than causal effects of scaffolding. Finally, the analysis focuses on English-language coding and writing conversations from three public corpora; other domains, languages, interfaces and institutional settings may organize different opportunities for engagement and support. These boundaries motivate the next stage of evidence: connecting the behavioural signatures identified here to independent measures of learning, autonomy and transfer in controlled and longitudinal settings.

Conclusion

As LLMs become infrastructure for everyday work and study, the central educational question extends beyond how individuals learn from intelligent systems to how learning opportunities are organized in AI-mediated societies. This study shows that ordinary human–LLM interaction already contains measurable behavioural signatures of informal learning, and that deeper constructive engagement emerges selectively through task ecologies, support forms and short interactional sequences. The long-term significance of AI may therefore depend not only on the knowledge or productivity it

delivers, but also on how societies organize learning opportunities: who encounters them, how often they arise and whether institutions sustain them over time.

4 Methods

4.1 Study design and analytic scope

This study analyses naturalistic public human–LLM conversations as observational digital traces of everyday work and study. We use three public conversational corpora: WildChat-4.8M (public non-toxic release) [43] (reported below as WildChat), LMSYS Chat-1M [24] (reported below as LMSYS Chat) and the ShareChat [25] strict-English subset (reported below as ShareChat). Within these corpora, we focus on two common task ecologies, coding and writing, because both are frequent everyday LLM uses but expose different opportunities for visible learning-oriented behaviour. Coding conversations often contain debugging, error diagnosis and iterative testing; writing conversations more often contain drafting, revision and directive editing [31–33].

The design is descriptive and associational, following the logic of large-scale digital trace and computational social-science analyses [44, 45]. The empirical sequence mirrors the Results section. We first describe whether learning-oriented engagement is visible in ordinary use, then ask where constructive engagement is concentrated, whether scaffolded assistant support marks richer constructive participation, which support forms carry that association and how support and engagement are ordered at the adjacent-turn scale. These analyses support claims about observable behavioural signatures in conversation logs. They do not estimate randomized exposure effects, durable learning outcomes, changes in users’ knowledge or stable learner traits.

4.2 Corpus construction and task settings

The analytic corpus contains 128,569 conversations and 981,470 total turns, including 491,685 user turns and 489,785 assistant turns. The six task settings are WildChat coding, WildChat writing, LMSYS Chat coding, LMSYS Chat writing, ShareChat coding and ShareChat writing. WildChat and LMSYS Chat provide large public human–LLM interaction logs from deployed chat systems, whereas ShareChat contributes a smaller strict-English analytic subset after our language screen. The settings are complementary naturalistic environments rather than experimentally matched platform conditions [23–25].

Corpus construction followed the same high-level sequence across datasets. Public source records were converted to role-ordered user–assistant conversations; conversations were restricted to English-oriented exchanges with at least four turns; coding and writing task ecologies were assigned with LLM-assisted semantic filters; and retained conversations were passed through the shared annotation pipeline. When keyword pre-filters were used, the LLM-assisted semantic task filter retained conversations with a score of at least 0.70. ShareChat additionally underwent a strict-English screen after task filtering, retaining only conversations whose final conversation-level language flag was English. Supplementary Tables A1 and A2 provide the raw-to-final provenance ledger, classifier sources and validation evidence for these filtering decisions.

4.3 Operational constructs

4.3.1 Task ecology, user framing and depth

Each conversation is represented as an ordered sequence of user and assistant turns together with available metadata. The main contextual variables are data source, task ecology, user framing, topic and language. Task ecology and user framing are analytic variables generated by the filtering and annotation pipeline; topic, language and model/source metadata are retained where available. User framing is a conversation-level LLM-assisted label for whether the opening user request or early exchange explicitly signals understanding, exploration or learning, as opposed to a request without such explicit learning-oriented framing. We interpret this as an observed textual signal, not as a stable measure of motivation or learner identity.

Interaction depth is measured from the number and ordering of turns in the conversation. It is used descriptively, as a factor of interactional opportunity and sustained exchange, not as an independent manipulation. Post-answer depth is measured as the number of turns after the first assistant response. Because any-event outcomes are easier to observe in longer conversations, the supplementary analyses include rate-based sensitivity checks for constructive engagement.

4.3.2 User engagement

The unit of annotation is the conversational turn. User turns are labelled for behavioural, emotional and cognitive engagement, drawing on engagement theory and the ICAP distinction between passive, active and constructive participation [20, 21]. Passive engagement (P) captures reception or acknowledgement without substantial transformation. Active engagement (A) captures task-relevant action, follow-up or manipulation of supplied information. Constructive engagement (C) captures deeper work with ideas, such as testing an implication, revising a hypothesis, articulating a constraint, transferring a mechanism or extending an explanation. Visible emotional expression is labelled separately and reported descriptively because it was rare.

The primary user-side outcome is constructive engagement, treated as the highest-level behavioural proxy for learning-oriented participation available in these logs. At the turn level, constructive engagement is a binary indicator for whether the user turn is labelled C. At the conversation level, we use both the count of constructive user turns and an indicator for whether a conversation contains at least one constructive user turn. Constructive ratios are defined as constructive user turns divided by user turns.

4.3.3 Assistant support

Assistant turns are labelled for scaffolding. Scaffolded support denotes assistant behaviour that structures, regulates, diagnoses or redirects the user’s work while leaving some responsibility with the user; the comparison category is a non-scaffolded reference response, including straightforward task fulfilment or stand-alone information provision. This definition follows the scaffolding tradition, especially contingency, transfer of responsibility and potential fading [27, 28, 30].

The primary assistant-side exposure is scaffolding, operationalized as scaffolded support. At the turn level, this is the distinction between scaffolded support and a non-scaffolded reference response. At the conversation level, we use an indicator for whether a conversation contains at least one scaffolded assistant turn. Scaffolded turns can also receive support-intent labels (I1 metacognitive, I2 cognitive and I3 affective support) and non-mutually exclusive support-form labels (M1 feedback, M2 hinting, M3 instructing, M4 explaining, M5 modelling and M6 questioning). Support-form analyses condition on scaffolded turns and compare conversations containing each support-form label with scaffolded conversations that do not contain that label. Operational definitions are provided in Supplementary Table A3; paraphrased examples and two complete de-identified case transcripts are provided in Supplementary Tables C1–C3.

4.4 Annotation workflow and validation

The annotation workflow combined theory-grounded codebook development, prompt iteration, parser checks, targeted post-processing, human review and production-scale LLM-assisted labelling. We used LLM-assisted annotation as a scalable text-analysis workflow while retaining human verification because prior work shows both the promise and the need for validation when language models are used as annotators [46, 47]. After the workflow stabilized, the same construct definitions, schema checks and post-processing logic were applied to the WildChat, LMSYS Chat and ShareChat analytic corpora. The archived annotation metadata record the provider, deployment family, API version, prompt version, retry policy, parser checks and rule-based post-processing rules used for production labelling; Supplementary Section A.3 gives the reproducibility details.

Human verification is reported in layers. Human–human agreement assesses whether trained reviewers could apply the codebook consistently. Human–LLM production-label checks compare final or post-processed production labels with human-confirmed labels. Targeted post-production audits check whether narrow rule-based reviews reduced known residual false positives. Agreement is summarized with per-label F1, Matthews correlation coefficient and Gwet’s AC1 because the labels are sparse, prevalence-imbalanced and sometimes non-mutually exclusive [48, 49]. The consolidated production-label audits indicate acceptable or better reliability across the labels used for the main claims: intentional framing reached F1=0.857, constructive-versus-non-constructive user engagement reached F1=0.922, broad scaffolded support reached F1=0.912 and support-form labels ranged from F1=0.691 to 0.938. Supplementary Tables C4 and C5 report the complete validation metrics and audit designs.

4.5 Statistical analyses

Descriptive analyses estimate engagement prevalence, scaffolding prevalence, passive/active/constructive composition, visible emotional-expression rates and behavioural-depth summaries within each task setting. Context analyses compare observed user framing, task ecology and depth groups. Because the probability that a conversation contains at least one constructive turn partly reflects the number of turns

available for observing one, we also fit a grouped-binomial constructive-rate sensitivity model in which the outcome is the constructive user-turn count out of total user turns, with user framing, task ecology, length bucket and dataset fixed effects as predictors.

Conversation-level support associations compare conversations with and without at least one scaffolded assistant turn. Constructive-ratio contrasts use turn-weighted ratios within each comparison group, with uncertainty from 1,000 conversation-level bootstrap resamples. Adjusted Section 2.3 models use setting-specific Poisson generalized linear models for raw constructive-turn counts and companion logistic models for whether a conversation contains at least one constructive user turn. The primary Poisson model treats conversation length as a covariate rather than using a user-turn offset; offset-rate and quasi-Poisson sensitivities are reported in the Supplement. These models are interpreted as adjusted associations, not causal effects.

Support-form contrasts compare scaffolded conversations containing a given support form with scaffolded conversations not containing that form. Because support-form labels are non-mutually exclusive and nested within scaffolded support, these contrasts are descriptive form-pattern analyses rather than mutually exclusive treatment comparisons. Bracket tests in Fig. 4 use Benjamini–Hochberg false-discovery-rate adjustment within each displayed family of six support-form comparisons.

Temporal analyses focus on adjacent-turn ordering. We compare the probability that the next user turn is constructive after scaffolded versus non-scaffolded assistant turns, estimate the same contrast within prior user states and estimate the probability that the next assistant turn is scaffolded after different prior user states. Adjacent-turn contrasts use 1,000 conversation-cluster bootstrap resamples because multiple turn pairs can come from the same conversation. Integrated adjacent-turn regressions combine user-state features, assistant-scaffolding features and adjacent-turn context in one model, with conversation-clustered standard errors. Supplementary Tables C6–C9 and C8 give the complete estimand map, covariate sets, model specifications and sensitivity checks.

4.6 LLM use and reporting boundaries

LLMs were used as objects of study and as annotation tools. The authors also used LLM-based tools to assist with code drafting, debugging, analysis workflow organization and language editing. LLM-generated assistance was not used as a substitute for human verification of the validation samples or for interpretation of the statistical results.

All reported results are aggregate turn-level or conversation-level analyses. The study does not link conversations into user histories, estimate user-level longitudinal trajectories or redistribute verbatim raw conversation logs. Operational examples are de-identified and paraphrased or lightly edited, and temporal analyses are confined to ordering within a conversation. Model/source checks are used as robustness analyses rather than model-causal comparisons, because model labels in public logs are not experimentally assigned and may be confounded with time, user population, task mix and logging context.

Data availability

This study analyses public human–LLM conversation corpora released by their original data providers: WildChat, LMSYS Chat-1M and ShareChat. Readers should obtain the raw public conversation data from the original releases and comply with the corresponding dataset licences, terms of use and redistribution restrictions. The manuscript does not redistribute verbatim raw conversation logs, user identifiers or linked user histories. The Supplementary Information includes two complete de-identified illustrative transcripts derived from retained conversations and lightly edited to preserve turn sequence, task structure and annotation decisions while reducing privacy and licence risk.

Derived data sufficient to verify the reported aggregate figures, tables and statistical summaries are archived on Zenodo at <https://doi.org/10.5281/zenodo.20995946>. The same release is also available in the GitHub repository at <https://github.com/CinderD/Informal-learning-in-everyday-human-LLM-interaction> (tag v0.1.0, commit 31e10cfea930). The release includes the numeric values underlying the main figures, de-identified conversation-level and adjacent-turn analytic labels without raw message text, table source files, confidence intervals, p values, false-discovery-rate q values, model-output summaries, provenance documentation and SHA-256 checksums. Any request for additional derived label files that are not included in the public release is assessed against the original dataset licences and privacy constraints.

Code availability

Custom code used to export source data and de-identified label tables, compute statistical analyses and generate figures and tables is archived with the derived-data release on Zenodo at <https://doi.org/10.5281/zenodo.20995946> and is also available in the GitHub repository at <https://github.com/CinderD/Informal-learning-in-everyday-human-LLM-interaction> (tag v0.1.0, commit 31e10cfea930). The repository excludes API keys, raw message text and files that cannot be redistributed under the original corpus terms. The archived release also includes annotation prompt templates, output schemas, parser checks, post-processing scripts and production manifests, excluding API keys and raw message text. Reproducing the full annotation pipeline from raw conversations requires user-provided API credentials and access to the raw public corpora from their original sources.

Ethics statement

The study is a secondary analysis of public, naturalistic human–LLM conversation datasets. The authors did not recruit participants, intervene in conversations or attempt to identify users. Analyses are reported at aggregate turn or conversation level. Operational examples are de-identified and paraphrased or lightly edited; the two complete illustrative transcripts preserve the original turn structure and annotation decisions but are not redistributed verbatim raw logs. Human validation was used to assess annotation quality on reviewed examples and did not involve interaction with the original dataset users.

Competing interests

H.L. and X.X. are employees of Microsoft Research Asia. The remaining authors declare no competing interests.

Author contributions

Author contributions are organized using the CRediT taxonomy. The contribution categories for this study are conceptualization, study design, data curation, annotation pipeline development, validation, statistical analysis, figure generation, writing of the original draft, manuscript review and editing, supervision and funding acquisition. Author initials and final role allocation: [to be added after the author list and order are confirmed]. All listed authors will approve the submitted version and will be accountable for their contributions.

Acknowledgements and funding

Acknowledgements and funding information, including grant numbers where applicable: [to be added]. Funders had no role in study design, data collection, analysis, interpretation, manuscript preparation or publication decisions unless stated otherwise.

LLM use statement

Large language models were objects of study and were also used as annotation and writing-assistance tools. Production labels were generated with LLM-assisted workflows and checked through parser rules, targeted update scripts and human validation. The authors also used LLM-based tools to assist with code drafting, debugging, workflow organization and language editing. No LLM or AI tool is listed as an author, and the authors take responsibility for the final manuscript, analyses and interpretation. Additional details are provided in Methods.

Reporting summary and editorial policy checklist

A Nature Portfolio Reporting Summary and the journal-required editorial policy checklist have been prepared as separate submission files. The manuscript and supplementary information report the observational design, sampling and filtering rules, annotation workflow, human validation evidence, statistical model specifications, bootstrap procedures, multiple-comparison handling and data/code availability statements needed to assess reproducibility.

Figure source data

Source data for the numeric main figures are provided in `source_data/figure2_source_data.csv`, `source_data/figure3_source_data.csv`, `source_data/figure4_source_data.csv`

and `source_data/figure5_source_data.csv`. Figure 1 is a conceptual framework figure and has no numeric source data. Supplementary table and regression source files are provided in `tables/` and `outputs/integrated_regression/`.

References

- [1] Chatterji, A. *et al.* How people use ChatGPT. Tech. Rep. Working Paper 34255, National Bureau of Economic Research (2025).
- [2] Noy, S. & Zhang, W. Experimental evidence on the productivity effects of generative artificial intelligence. *Science* **381**, 187–192 (2023).
- [3] Brynjolfsson, E., Li, D. & Raymond, L. Generative AI at work. *The Quarterly Journal of Economics* **140**, 889–942 (2025).
- [4] Anzai, Y. & Simon, H. A. The theory of learning by doing. *Psychological review* **86**, 124 (1979).
- [5] Kolb, D. A. *Experiential learning: Experience as the source of learning and development* (FT press, 2014).
- [6] Risko, E. F. & Gilbert, S. J. Cognitive offloading. *Trends in cognitive sciences* **20**, 676–688 (2016).
- [7] Gerlich, M. Ai tools in society: Impacts on cognitive offloading and the future of critical thinking. *Societies* **15**, 6 (2025).
- [8] Barke, S., James, M. B. & Polikarpova, N. Grounded copilot: How programmers interact with code-generating models. *Proceedings of the ACM on Programming Languages* **7**, 85–111 (2023).
- [9] Kobiella, C., Flores López, Y. S., Waltenberger, F., Draxler, F. & Schmidt, A. “If the machine is as good as me, then what use am I?”—how the use of ChatGPT changes young professionals’ perception of productivity and accomplishment. *Proceedings of the CHI Conference on Human Factors in Computing Systems* 1–16 (2024).
- [10] Eraut, M. Informal learning in the workplace. *Studies in Continuing Education* **26**, 247–273 (2004).
- [11] Marsick, V. J. & Watkins, K. *Informal and incidental learning in the workplace (Routledge Revivals)* (Routledge, 2015).
- [12] Schugurensky, D. The forms of informal learning: Towards a conceptualization of the field. Tech. Rep. NALL Working Paper 19, Centre for the Study of Education and Work, OISE/UT, Toronto (2000). URL <https://hdl.handle.net/1807/2733>.

- [13] Kasneci, E. *et al.* ChatGPT for good? On opportunities and challenges of large language models for education. *Learning and Individual Differences* **103**, 102274 (2023).
- [14] Meyer, J. *et al.* Using LLMs to bring evidence-based feedback into the classroom: AI-generated feedback increases secondary students’ text revision, motivation, and positive emotions. *Computers and Education: Artificial Intelligence* **6**, 100199 (2024).
- [15] Arif, T., Asthana, S. & Collins-Thompson, K. Generation and assessment of multiple-choice questions from video transcripts using large language models. *Proceedings of the Eleventh ACM Conference on Learning @ Scale* 530–534 (2024).
- [16] Jin, H., Lee, S., Shin, H. & Kim, J. Teach AI how to code: Using large language models as teachable agents for programming education. *Proceedings of the CHI Conference on Human Factors in Computing Systems* 1–28 (2024).
- [17] Lim, H., Choi, D. & Hong, H. Identify design problems through questioning: Exploring role-playing interactions with large language models to foster design questioning skills. *Companion Publication of the 2024 Conference on Computer-Supported Cooperative Work and Social Computing* 598–602 (2024).
- [18] Terzimehić, N., Bühler, B. & Kasneci, E. Conversational ai as a catalyst for informal learning: An empirical large-scale study on llm use in everyday learning. *Computers and Education: Artificial Intelligence* 100634 (2026).
- [19] VanLehn, K. The relative effectiveness of human tutoring, intelligent tutoring systems, and other tutoring systems. *Educational Psychologist* **46**, 197–221 (2011).
- [20] Fredricks, J. A., Blumenfeld, P. C. & Paris, A. H. School engagement: Potential of the concept, state of the evidence. *Review of Educational Research* **74**, 59–109 (2004).
- [21] Chi, M. T. & Wylie, R. The ICAP framework: Linking cognitive engagement to active learning outcomes. *Educational Psychologist* **49**, 219–243 (2014).
- [22] Koedinger, K. R. & Alevan, V. Exploring the assistance dilemma in experiments with cognitive tutors. *Educational Psychology Review* **19**, 239–264 (2007).
- [23] Zhao, W. *et al.* WildChat: 1M ChatGPT interaction logs in the wild (2024). URL <https://openreview.net/forum?id=Bl8u7ZRlbM>.
- [24] Zheng, L. *et al.* LMSYS-Chat-1M: A large-scale real-world LLM conversation dataset (2023). [arXiv:2309.11998](https://arxiv.org/abs/2309.11998).

- [25] Yan, Y., Nguyen, T., Su, B., Lieffers, M. & Le, T. ShareChat: A dataset of chatbot conversations in the wild (2025). [arXiv:2512.17843](#).
- [26] Handa, K. *et al.* Which economic tasks are performed with AI? Evidence from millions of Claude conversations (2025). [arXiv:2503.04761](#).
- [27] Wood, D., Bruner, J. S. & Ross, G. The role of tutoring in problem solving. *Journal of Child Psychology and Psychiatry* **17**, 89–100 (1976).
- [28] Van de Pol, J., Volman, M. & Beishuizen, J. Scaffolding in teacher–student interaction: A decade of research. *Educational Psychology Review* **22**, 271–296 (2010).
- [29] Winkler, R., Hobert, S., Salovaara, A., Söllner, M. & Leimeister, J. M. Sara, the lecturer: Improving learning in online education with a scaffolding-based conversational agent. *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems* 1–14 (2020).
- [30] Quintana, C. *et al.* A scaffolding design framework for software to support science inquiry. *The Journal of the Learning Sciences* **13**, 337–386 (2004).
- [31] Keuning, H., Jeuring, J. & Heeren, B. A systematic literature review of automated feedback generation for programming exercises. *ACM Transactions on Computing Education* **19**, 1–43 (2018).
- [32] Flower, L. & Hayes, J. R. A cognitive process theory of writing. *College Composition and Communication* **32**, 365–387 (1981).
- [33] Sommers, N. Revision strategies of student writers and experienced adult writers. *College Composition and Communication* **31**, 378–388 (1980).
- [34] Hattie, J. & Timperley, H. The power of feedback. *Review of Educational Research* **77**, 81–112 (2007).
- [35] Shute, V. J. Focus on formative feedback. *Review of Educational Research* **78**, 153–189 (2008).
- [36] Nicol, D. J. & Macfarlane-Dick, D. Formative assessment and self-regulated learning: A model and seven principles of good feedback practice. *Studies in Higher Education* **31**, 199–218 (2006).
- [37] Chi, M. T. H., Bassok, M., Lewis, M. W., Reimann, P. & Glaser, R. Self-explanations: How students study and use examples in learning to solve problems. *Cognitive Science* **13**, 145–182 (1989).
- [38] Renkl, A. Learning from worked-out examples: A study on individual differences. *Cognitive Science* **21**, 1–29 (1997).

- [39] Aleven, V., Stahl, E., Schworm, S., Fischer, F. & Wallace, R. Help seeking and help design in interactive learning environments. *Review of Educational Research* **73**, 277–320 (2003).
- [40] Reiser, B. J. Scaffolding complex learning: The mechanisms of structuring and problematizing student work. *The Journal of the Learning Sciences* **13**, 273–304 (2004).
- [41] Brown, J. S., Collins, A. & Duguid, P. Situated cognition and the culture of learning. *Educational researcher* **18**, 32–42 (1989).
- [42] Collins, A., Brown, J. S. & Newman, S. E. Cognitive apprenticeship: Teaching the crafts of reading, writing, and mathematics. *Knowing, learning, and instruction* 453–494 (2018).
- [43] Allen Institute for AI. WildChat-4.8M. <https://huggingface.co/datasets/allenai/WildChat-4.8M> (2025).
- [44] Lazer, D. *et al.* Computational social science. *Science* **323**, 721–723 (2009).
- [45] Salganik, M. J. *Bit by Bit: Social Research in the Digital Age* (Princeton University Press, 2017).
- [46] Gilardi, F., Alizadeh, M. & Kubli, M. ChatGPT outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences* **120**, e2305016120 (2023).
- [47] Ziems, C. *et al.* Can large language models transform computational social science? *Computational Linguistics* **50**, 237–291 (2024).
- [48] Matthews, B. W. Comparison of the predicted and observed secondary structure of T4 phage lysozyme. *Biochimica et Biophysica Acta* **405**, 442–451 (1975).
- [49] Gwet, K. L. Computing inter-rater reliability and its variance in the presence of high agreement. *British Journal of Mathematical and Statistical Psychology* **61**, 29–48 (2008).

Supplementary information

Supplementary information is organized to support reproducibility and interpretation. Appendix A records the corpus ledger, annotation ontology, production workflow, validation evidence, denominator rules and model specifications. Appendix B reports robustness checks and interpretive boundaries. Appendices C and D contain the supporting tables and figures referenced from the main text and Appendix A–B.

Appendix A Reproducibility protocol and construct ledger

A.1 Corpus provenance and filtering ledger

The main Methods summarize the three public corpora and six task settings. This section provides the auditable selection ledger behind those settings. The analyses use WildChat coding, WildChat writing, LMSYS Chat coding, LMSYS Chat writing, ShareChat coding and ShareChat writing. The analytic corpus contains 128,569 conversations and 981,470 total turns. WildChat contributes 31,878 coding conversations and 39,534 writing conversations; LMSYS Chat contributes 32,114 coding conversations and 21,023 writing conversations; ShareChat contributes 2,481 coding conversations and 1,539 writing conversations.

Corpus construction followed a common sequence: source records were converted to role-ordered user–assistant conversations; conversations were screened for English-oriented content and at least four message turns; coding and writing task ecologies were assigned with LLM-assisted semantic filters; and retained conversations were processed by the shared annotation workflow. ShareChat required an additional strict-English screen after task filtering. Public dataset-card turns, raw CSV message rows and role-normalized analysis turns are not identical counting units, so the ShareChat ledger reports platform-url conversations and role-labelled message rows before deriving user and assistant turns. Supplementary Table [A1](#) records the raw-to-analysis selection ledger and retained provenance artifacts for each corpus. Supplementary Table [A2](#) separates corpus-construction filters from analysis labels and reports validation or audit evidence for task ecology, user framing, language filtering and turn-level annotation.

Table A1 Corpus construction provenance artifacts. The table records the retained non-text provenance artifacts used to audit corpus construction. SHA-256 values are shown as 12-character prefixes for run scripts or manuscript-source scripts retained in the analysis archive. Raw conversation text remains governed by the original dataset releases and is not redistributed with the manuscript source.

Corpus	Source release and conversion unit	Selection ledger retained for this manuscript	Script/source-data provenance
WildChat	WildChat-4.8M public release; role-ordered user/assistant messages after task filtering.	Public release → English-oriented coding/writing semantic task filters with minimum four message turns → refreshed task-filter outputs (31,879 coding; 39,534 writing) → user-turn relabel/update pass → analytic sample (31,878 coding; 39,534 writing).	WildChat relabel/update run script SHA-256 prefix c8eecef55858 ; source-data exporter 206aaae44791 ; camera-ready figure script d9ecf7118f10 .
LMSYS Chat	LMSYS Chat-1M public release; one conversation record per public chat, converted to role-ordered messages.	1,000,000 source conversations → 777,453 English conversations → 236,372 English conversations with at least four turns → semantic task filters → analytic sample (32,114 coding; 21,023 writing).	Coding filter launcher f54c079fda08 ; writing filter launcher be970385e72e ; after-filter annotation/analysis launcher 326e11709c41 ; source-data exporter 206aaae44791 .
ShareChat	ShareChat Hugging Face release; public card reports 142,808 conversations and 660,293 public turns. Local conversion uses five final-language-filtered CSVs with one row per message and platform-url as the conversation key.	Five final-language-filtered CSVs (1,199,899 message rows; 129,584 platform-url conversations) → role-normalized user/bot rows (1,196,969 rows) → 69,915 conversations with at least four message turns → semantic task filters (3,494 coding; 1,663 writing) → strict-English screen → analytic sample (2,481 coding; 1,539 writing).	Coding filter launcher f2ff9ca543b8 ; writing filter launcher 2788818044e7 ; after-filter annotation/analysis launcher 501e26e5a728 ; source-data exporter 206aaae44791 .

Table A2 Filtering and classifier validation audit. The table separates corpus-construction filters from analysis labels and reports the strongest validation or audit evidence available in the current artifacts. Metrics from the writing-210 validation set are domain-level checks, whereas the user-framing audit deliberately sampled across corpus-by-task settings.

Variable or filter	How it was assigned	Validation or audit evidence	Interpretation boundary
Task ecology: coding versus writing	LLM-assisted semantic task filters selected conversations matching coding-oriented or writing-oriented use after minimum-turn screening. The filters were English-oriented and used a score threshold of 0.70 after keyword prefiltering where applicable.	The provenance ledger records retained counts and selection steps for each corpus (Supplementary Table A1). In the writing-210 validation artifacts, the conversation-topic field showed high human-LLM agreement ($F1=0.973$, $MCC=0.872$, $Gwet\ AC1=0.936$; $support=210$), but this is a domain-level check rather than a full cross-corpus task-filter audit.	Task ecology is treated as an analytic classification of conversation content, not native dataset metadata or randomized assignment.
User framing: intentional versus unintentional	Conversation-level LLM-assisted learning_intent annotation marked conversations explicitly oriented toward understanding, exploration or learning versus those without explicit learning-oriented framing.	A human-verified positive-oversampled audit reviewed 450 first user turns across all six corpus-by-task settings (75 per setting; 50 production-positive and 25 production-negative). Against the production labels, overall intentional-framing $F1$ was 0.857, MCC was 0.677 and $Gwet\ AC1$ was 0.671; task-stratified $F1$ was 0.924 for coding and 0.779 for writing.	User framing is a measured textual signal of explicit learning orientation, not a stable learner trait, motivation measure or causal exposure. Because the audit oversampled production-positive cases, it validates classification boundaries rather than estimating population prevalence.
Language filtering	WildChat and LMSYS Chat used English-oriented source or semantic filters. ShareChat received an additional strict-English screen after task filtering, retaining only records with final conversation-level language flag Eng11sh .	The ShareChat strict-English audit recorded zero malformed JSON files and explicit exclusions after task filtering: 1,013 coding and 124 writing conversations were removed by the final language screen. In the writing-210 validation artifacts, language showed near-perfect agreement ($F1=0.998$, $MCC=0.705$, $Gwet\ AC1=0.995$; $support=210$), with MCC lower because nearly all reviewed records were English.	The term strict-English is used only for ShareChat’s final post-task language screen. Mixed or non-English ShareChat records were excluded from the main analytic sample.
Turn-level labels: user engagement and scaffolded support	Production annotation used the final codebook, parser checks and post-processing rules. Constructive engagement and scaffolded support are the primary user-side outcome and assistant-side exposure.	Domain-stratified human–human codebook validation shows strong agreement for constructive engagement (coding $F1=0.85$; writing $F1=0.89$) and scaffolded support (coding $F1=0.87$; writing $F1=0.97$), with full per-label codebook metrics in Supplementary Table C4. Human-confirmed production audits across all six corpus-by-task settings showed acceptable or better reliability for user engagement and assistant labels: the 600-case user audit had constructive-vs-non-constructive $F1=0.922$, $MCC=0.840$ and $Gwet\ AC1=0.837$; the 180-case assistant audit had scaffolded-support $F1=0.912$, $MCC=0.590$ and $Gwet\ AC1=0.791$, with support-form $F1$ values ranging from 0.691 to 0.938 (Supplementary Table C5).	These labels support the main association and temporal analyses. The production audits are positive-oversampled and validate label boundaries rather than prevalence.

For each retained conversation, the pipeline stored the ordered turn sequence, role labels, available topic and language metadata, learning-intent indicators and model/source metadata where available. Model metadata are used only for robustness checks because model assignment in naturalistic public conversations is observational and may reflect time, user population, task mix and logging context.

A.2 Annotation ontology and operational examples

The codebook translates learning-science constructs into role-specific behavioural signatures observable in conversation logs. User turns and assistant turns are labelled with different constructs because they play different roles in the interaction. Supplementary Table A3 gives the operational definitions and boundary rules. Supplementary Table C1 gives paraphrased examples that illustrate the difference between engagement levels and support forms without reproducing raw user text. Supplementary Tables C2 and C3 provide complete de-identified illustrative transcripts, one coding-oriented and one writing-oriented, derived from retained conversations and lightly edited while preserving turn sequence, task structure and annotation rationales.

Table A3 Annotation codebook definitions. The table defines the labels used to translate learning-science constructs into observable turn-level signals. These labels are behavioural annotations of conversation text; they are not measures of durable learning, learner identity or instructional efficacy.

Family	Label	Definition	Boundary rule
User engagement	Passive (P)	Reception, acknowledgement or minimal continuation without visible transformation of the assistant response.	Includes short confirmations or requests to continue; excludes task-relevant follow-up or reasoning.
User engagement	Active (A)	Task-relevant action, follow-up, application or manipulation of supplied information.	The user does something with the response, but does not test, revise or extend an idea.
User engagement	Constructive (C)	Highest-depth observable user engagement: elaborating, testing, revising, transferring or extending an idea or explanation.	Requires visible user reasoning or transformation; not inferred from long text alone.
User expression	Emotional (E)	Explicit affective expression, such as frustration, gratitude, anxiety or excitement.	Coded separately from cognitive engagement and interpreted as visible expression in the log.
Assistant support	Scaffolding	Assistant support that structures, diagnoses, regulates or redirects the user's process while leaving cognitive or task responsibility with the user.	Distinguished from reference responses that simply provide an answer, draft or deliverable.
Support intent	I1 metacognitive	Support aimed at planning, monitoring, reflection or regulation of the user's work.	Captures the purpose of support, not its surface form.
Support intent	I2 cognitive	Support aimed at concepts, mechanisms, task knowledge or problem-solving content.	Most scaffolded support in this corpus is cognitive in purpose.
Support intent	I3 affective	Support aimed at encouragement, reassurance or affective regulation.	Used only when affective support is substantively present, not for polite tone alone.
Support form	M1 feedback	Evaluative information about the user's prior attempt, diagnosis or understanding.	Requires a target in the user's preceding contribution.
Support form	M2 hinting	A cue, prompt or partial path that helps the user proceed without fully resolving the task.	Leaves the next inferential or task step with the user.
Support form	M3 instructing	Explicit procedural guidance, ordered steps or strategy for doing the task.	More directive than hinting; may still count as scaffolding when it structures process.
Support form	M4 explaining	A rationale, mechanism or causal account that clarifies why something works or fails.	Requires explanatory content, not merely a restatement of an answer.
Support form	M5 modelling	A reusable pattern, worked example or demonstration intended to be adapted.	Distinguished from simply giving the final requested product.
Support form	M6 questioning	A question that advances diagnosis, reflection, planning or problem solving.	Excludes generic clarification questions that do not support the user's process.

User-turn labels describe observable engagement. Passive engagement (P) marks acknowledgement, reception or minimal continuation without visible transformation. Active engagement (A) marks task-relevant action or follow-up, such as applying supplied information or requesting a concrete modification. Constructive engagement (C) marks deeper work with ideas, such as testing an implication, revising a hypothesis, articulating a constraint, transferring a mechanism or extending an explanation. Emotional engagement is labelled separately and used descriptively because visible affective expression was rare.

Assistant-turn labels describe support. Scaffolded support denotes assistant behaviour that structures, regulates, diagnoses or redirects the user’s work while leaving some responsibility with the user. Non-scaffolded reference responses provide the requested information or deliverable without materially structuring the user’s process. Scaffolded turns can also receive support-intent labels and support-form labels. Support-intent labels distinguish metacognitive support (I1), cognitive support (I2) and affective support (I3). Support-form labels distinguish feedback (M1), hinting (M2), instructing (M3), explaining (M4), modelling (M5) and questioning (M6). Support-form labels are non-mutually exclusive because a single assistant turn can both explain a mechanism and provide an example or hint. Supplementary Tables C18 and C19 report support-intent and support-form profiles across the six task settings.

A.3 Production annotation workflow

The production workflow had three reproducible stages. First, the authors developed a theory-grounded codebook and prompt templates on validation examples and disagreement cases. This stage identified recurring failure modes, including over-labelling brief acknowledgements as active engagement, treating any long assistant answer as scaffolding, collapsing directive instructions into scaffolding without checking whether responsibility remained with the user and over-labelling sparse support forms. Second, parser checks enforced the expected schema, normalized label variants and flagged missing or malformed outputs. Third, production-scale annotation applied the frozen schema to the six analytic settings, followed by narrowly specified rule-based corrections for known boundary errors.

Production annotation used Azure OpenAI chat-completions deployments from the GPT-5.1 family for turn-level user and assistant labels, with GPT-5.1-chat used for the final writing-oriented calls. LLM-assisted semantic filtering used GPT-4o-family deployments. The released analysis code and archived annotation metadata retain the full prompt templates, functional prompt role, schema, deployment family, API version, retry settings and parser checks. In the manuscript and supplement we refer to prompt function rather than internal prompt-path names, because those path names are run-management identifiers rather than scientific constructs.

The production calls used structured JSON outputs where supported. The final writing calls used JSON-object or JSON-schema response constraints, 60-second request timeouts, maximum completion budgets of 300 tokens for conversation-level labels and 350 tokens for turn-level labels and did not pass an explicit temperature parameter. The final coding-from-filtered-conversation runner likewise did not pass an explicit temperature parameter. English-language checks used temperature 0.0 and

semantic task checks used temperature 0.1. Retry policies were specified in code: coding runs used up to five attempts with exponential backoff for rate limits and fixed-delay retry for other transient failures; writing runs used up to six attempts, disabled SDK-level retries, used exponential backoff with jitter for rate limits, timeouts and server errors and fell back from JSON-schema response format to JSON-object format if unsupported. Content-filter blocks returned empty analyses and were logged rather than silently imputed.

Post-processing was part of the predefined production annotation workflow, not an after-the-fact statistical adjustment. We use the term for two transparent operations. Schema-level processing normalized output variants and recovered fields that were present in the model explanation but missing from the structured JSON when the mapping was unambiguous. Boundary-rule processing corrected recurring, narrow false-positive patterns identified during validation, especially cases where brief user acknowledgements were labelled as Active or Constructive, generic emotional tone was labelled as Emotional or direct answer delivery was labelled as scaffolded support. Coding outputs did not use a broad post-processing chain in the final filtered-conversation runner, but WildChat user turns underwent the targeted user-turn review described in Section A.4. Writing outputs used the retained production chain for assistant scaffolding recovery, support-detail filling and strict user Active, Constructive and Emotional boundary filtering. API keys and service endpoints are intentionally not reproduced in the manuscript.

A.4 Human validation and production-label reliability

Human review was used during codebook development, prompt iteration and final validation. Two trained human reviewers with human–AI interaction and LLM-interaction research backgrounds reviewed validation cases. The purpose of human review changed over the development cycle: early review identified codebook boundary problems and prompt failure modes, whereas the final validation tables report agreement after the codebook, parser checks and production workflow had been stabilized. The final agreement metrics should therefore be read as evidence for the frozen workflow, not as a claim that high agreement was achieved at the start of development.

Validation evidence is reported in three layers. First, human–human agreement assesses whether trained reviewers could apply the codebook consistently; Supplementary Table C4 reports this codebook-operability layer. Second, production-label checks compare final or post-processed LLM-assisted production labels with human-confirmed labels; Supplementary Table C5 consolidates these human–LLM agreement checks. Third, targeted post-production audits test whether narrow rule-based reviews reduced known residual false positives after prompt and parser development. Coding validation combined three review rounds, with 658 task turns and 643 turns after deduplication. Writing validation used 210 reviewed conversations with final human-reviewed labels. Because many labels are sparse or prevalence-imbalanced, agreement is reported with per-label F1, Matthews correlation coefficient and Gwet’s AC1 rather than with a single global agreement statistic.

The final user-turn check contained 100 coding examples and 100 writing examples reviewed after the targeted user-turn review. These examples were used to identify residual false positives and calibrate narrow rule-based corrections, especially for Active and Passive user-turn labels. The final WildChat user-turn review scanned 31,878 coding conversations and 39,534 writing conversations, applying 1,815 coding corrections and 21 writing corrections. LMSYS Chat and ShareChat were then processed with the same refreshed annotation workflow, construct definitions and post-processing logic. The main manuscript reports the labels produced by these final production runs.

Because user framing carries an important contextual role, the consolidated production-label audit includes a human-verified check for the conversation-level learning-intent label. This audit sampled 450 first user turns, with 75 cases from each corpus-by-task setting and with production-positive cases deliberately oversampled. Overall intentional-framing agreement with the production label was $F1=0.857$, $MCC=0.677$ and $Gwet\ AC1=0.671$; task-stratified $F1$ was 0.924 for coding and 0.779 for writing. This audit validates the label boundary for user framing but is not used to estimate population prevalence because of the positive-oversampled design.

For user engagement production labels, the consolidated audit includes a 600-case check across the six corpus-by-task settings, with 100 user turns per setting and production constructive cases deliberately oversampled. The final labels were human-confirmed after item-level review. Against these human-confirmed labels, the production constructive-versus-non-constructive distinction had $F1=0.922$, $MCC=0.840$ and $Gwet\ AC1=0.837$. One-vs-rest metrics for the four user engagement levels were also high: $F1=0.922$ for constructive, 0.883 for active, 0.941 for passive and 0.930 for none (Supplementary Table C5). Because the sample is positive-oversampled, it validates label boundaries and production-label reliability rather than estimating population prevalence.

For assistant production labels, the consolidated audit includes a 180-case check across the six corpus-by-task settings, with 30 assistant turns per setting and scaffolded-support and support-form positives deliberately oversampled. The final labels were human-confirmed after item-level review. Broad scaffolded-support agreement was $F1=0.912$, $MCC=0.590$ and $Gwet\ AC1=0.791$. Support-form one-vs-rest $F1$ ranged from 0.691 to 0.938 across M1–M6 (Supplementary Table C5). Because the sample is positive-oversampled, it validates assistant-label boundaries rather than estimating support-form prevalence.

The validation evidence for turn-level engagement and support labels is therefore layered: human–human agreement supports codebook operability, targeted production audits directly check final user engagement and assistant scaffolding/support-form production labels across all six settings and post-update audits check known residual errors. Cross-corpus use of the labels is supported by applying the same schema, parser checks and post-processing rules to all three corpora and by reporting setting-level robustness checks in the Results and Supplementary Tables.

A.5 Derived variables and denominator rules

This section defines how turn labels become the quantities plotted or modelled in the Results. At the turn level, the focal user-side outcome is a binary indicator for constructive engagement. At the conversation level, we use both the count of constructive user turns and an indicator for whether at least one constructive user turn occurs. Constructive ratios divide constructive user turns by total user turns; group-level constructive-ratio contrasts are turn-weighted ratios computed from total constructive user turns over total user turns within each group. Cognitive engagement ratios include passive, active and constructive cognitive engagement. Unless otherwise noted, constructive ratios use all user turns in the corresponding conversation or group as the denominator.

At the assistant side, scaffolding is operationalized as scaffolded support. Conversation-level scaffolding presence indicates whether a conversation contains at least one scaffolded assistant turn. Support-form analyses condition on conversations or turns with scaffolded support and then compare the presence versus absence of each support-form label. These comparisons are descriptive because support-form labels can co-occur and because conversations are not randomly assigned to receive a support form.

Behavioural depth is computed from turn counts and turn ordering. Post-answer depth counts turns after the first assistant response. Persistence after failure is analysed only for coding-oriented conversations because debugging failures, error messages and failed-execution contexts are more directly observable in coding than in writing.

A.6 Statistical estimands and model specifications

This section maps each inferential quantity in the Results to its denominator, bootstrap unit, model formula, exposure specification and sensitivity checks. Descriptive prevalence analyses report proportions within the six task settings. Context analyses compare user framing, task ecology and depth groups. Conversation-level support associations use Poisson generalized linear models for constructive-turn counts and logistic models for binary constructive-engagement outcomes. Adjacent-turn analyses use assistant-to-user and user-to-assistant turn pairs to estimate conditional probabilities at the next-turn scale. Supplementary Tables C6, C8 and C9 provide the compact estimand map.

For the Section 2.3 adjusted conversation-level models, the Poisson outcome is the raw count of constructive user turns in a conversation, not a rate. The primary model does not include a user-turn offset. Conversation length is adjusted directly through total turns, so the exponentiated Poisson coefficient should be read as a multiplicative association with expected constructive-turn count conditional on observed conversation length and other covariates. The generic setting-specific mean model is $\mathbb{E}(C_i) = \exp(\alpha_s + \beta_s S_i + \gamma_s^\top X_i)$, where C_i is constructive-turn count, S_i is scaffolded-support presence and X_i contains the observed conversation-level context listed in Supplementary Table C8. The companion logistic model predicts whether a conversation contains any constructive user turn using scaffolded-support presence and the available setting-specific context. The user-framing stratified Poisson models use the

same Poisson specification after stratifying by framing and omitting user framing from the covariate set.

Covariates are included when they correspond to observed compositional differences at the analysis level. Depending on the analysis, these include user framing, task indicators in pooled models, turn counts, emotional ratios, error-related variables and turn index. Percentage-point contrasts are reported as effect sizes. Confidence intervals and p values for these contrasts use 1,000 conversation-level bootstrap resamples; adjacent-turn contrasts use 1,000 conversation-cluster bootstrap resamples rather than individual turn-pair resampling to account for repeated pairs within conversations. The adjusted Poisson and logistic estimates in Fig. 3 use model-based standard errors. Poisson overdispersion was checked with the Pearson χ^2 statistic divided by residual degrees of freedom. Dispersion ranged from 1.16 to 2.06 across the six task settings; scaling the Poisson standard errors with a quasi-Poisson dispersion adjustment preserved positive scaffolded-support associations and statistical distinguishability in every setting. To check exposure specification, we also refitted the same setting-specific Poisson mean model with user turns represented as the exposure offset; Supplementary Table C11 reports quasi-Poisson scaled rate ratios. Negative-binomial models were not used as the primary specification; the sensitivity analyses focus on variance misspecification and exposure specification under the same mean model rather than changing the main estimand. Support-form bracket tests use Benjamini–Hochberg false-discovery-rate adjustment within each six-form comparison family.

The Section 2.2 constructive-rate sensitivity uses a grouped-binomial logistic model at the conversation level. The numerator is the number of constructive user turns and the denominator is total user turns in the conversation. The predictors match the any-constructive context model: user framing, task ecology, length bucket and dataset fixed effects. Standard errors use a conversation-robust sandwich estimator. This sensitivity is included because any-constructive-turn outcomes partly reflect the number of turns available for observing at least one constructive turn.

The integrated adjacent-turn regression combines three feature groups in one logistic model: user state before the assistant turn, assistant scaffolding features and adjacent-turn context. The pooled specification includes dataset fixed effects. A model/source sensitivity uses exact WildChat model labels, LMSYS Chat model names recovered from conversation identifiers and ShareChat assistant/source families recovered from conversation identifiers. These adjusted estimates are interpreted as measured-covariate-adjusted associations. They do not remove unobserved user motivation, prior expertise, task difficulty or time-varying conversational state.

A.7 Temporal estimands and interpretive boundary

The temporal appendix clarifies the scale of the temporal claim. The logs identify within-conversation ordering, but they do not randomize assistant support or observe durable learning outcomes. We therefore distinguish adjacent-turn estimands from coarser within-conversation summaries. The adjacent-turn analyses ask whether the immediately following user turn is more likely to be constructive after scaffolded support than after a non-scaffolded reference response, whether that contrast depends on the prior user state and whether support-form signals differ across prior user states.

The reverse adjacent-turn analysis asks whether the assistant is more likely to provide scaffolded support after constructive, active or passive user turns.

Coarser before–after summaries ask a different question: whether constructive engagement is higher after the first scaffolded assistant turn than before it. These are retained only as exploratory boundary checks, not as treatment-effect or accumulation estimates. The first scaffolded turn may occur precisely because the user is struggling, because a difficult task has emerged or because the exchange is already developing. The main temporal claim is therefore scale-specific: the clearest evidence concerns adjacent-turn, state- and form-conditioned coupling rather than a demonstrated global trajectory or cumulative learning gain.

Appendix B Robustness checks and interpretive boundaries

B.1 Cross-dataset consistency

The primary robustness checks are cross-dataset checks, not WildChat-only checks. Whenever an estimand can be computed from the shared annotation schema, it is reported across all six WildChat, LMSYS Chat and ShareChat task settings. These six-setting checks include the main prevalence summaries, constructive-ratio contrasts, post-answer depth contrasts, adjacent-turn lifts, support-form supply profiles and setting-level adjacent-turn regressions (Supplementary Tables C10, C14 and C19; Supplementary Figures D1–D3).

The Results claims are checked at the level appropriate to each estimand. Section 2.1 is a descriptive prevalence claim: cognitive engagement, constructive engagement and scaffolded support were observable in all six settings, with constructive engagement consistently the highest-level and less frequent user-side signal. Section 2.2 is primarily descriptive: user framing, task ecology and interaction depth organize where constructive engagement appears, with the same directional depth gradient across settings. Formal uncertainty estimates are reported where the claim involves a contrast or model coefficient. Section 2.3 uses conversation-level bootstrap confidence intervals and p values for turn-weighted scaffolded versus non-scaffolded constructive-ratio contrasts, adjusted Poisson/logistic model estimates and an offset-rate Poisson sensitivity. Section 2.4 reports support-form confidence intervals and between-stratum FDR-adjusted q values. Section 2.5 reports conversation-cluster bootstrap confidence intervals for adjacent-turn lifts and cluster-robust regression estimates for the integrated adjacent-turn models.

The main percentage-point contrasts were directionally consistent across the six settings. Constructive-ratio contrasts comparing conversations with versus without scaffolded support were positive and statistically distinguishable from zero in all six settings, with effect sizes ranging from +0.99 to +7.34 percentage points. Adjacent-turn constructive lifts were also positive in all six settings, with five of six contrasts statistically distinguishable from zero; LMSYS writing was positive but weaker (+0.27 percentage points, 95% CI -0.01 to +0.56, $p = .061$). Supplementary Table C10 reports these effect sizes, confidence intervals and p values.

B.2 Model/source robustness

The only robustness displays restricted to WildChat are the model-snapshot figures in Supplementary Figures D5–D9. They are restricted to WildChat because WildChat exposes user-facing model metadata with sufficient coverage for repeated within-corpus snapshot plots. LMSYS Chat and ShareChat are still included in the cross-dataset and model/source-adjusted analyses, but they do not provide directly comparable model-snapshot structure for the same visual repetition. The WildChat model-snapshot analyses preserved the main directional patterns reported in the main text: coding generally showed higher cognitive engagement than writing; intentionality gaps were visible across most user-facing models; scaffolded conversations were generally more constructive; and adjacent-turn contrasts were more interpretable than coarser within-conversation contrasts.

Integrated adjacent-turn regressions were used as robustness checks rather than as the primary narrative claim. These models predicted whether the next user turn was constructive after combining user-state features, assistant-support features and adjacent-turn context. Because model metadata differ by corpus, we report both within-setting models and pooled specifications. The within-setting models show the dataset-by-task pattern directly, with recoverable model/source fixed effects added separately inside each setting (Supplementary Table C14). The pooled specification includes dataset fixed effects; a sensitivity replaces these with exact WildChat model labels, LMSYS model names recovered from conversation identifiers and ShareChat assistant/source families recovered from conversation identifiers. This is interpreted as model/source adjustment rather than as a causal comparison among model versions. Model/source heterogeneity was detectable, but broad scaffolded support still improved prediction after user-state, context and model/source controls (LR $\chi^2_1 = 509.7$, $p < .001$). Because M1–M6 are support forms nested within scaffolded support rather than independent covariates, we interpret the decomposed models as asking which forms carry additional signal among scaffolded turns, not as a test of scaffolded support after removing its own components. Adding M1–M6 descriptors within scaffolded support improved fit (LR $\chi^2_6 = 829.8$, $p < .001$), and the full scaffolding block remained strongly informative after controls (LR $\chi^2_7 = 1339.6$, $p < .001$; Supplementary Table C16). Adding prior-state $\times$ support-form interactions also improved fit (pooled model/source fixed effects: LR $\chi^2_{18} = 297.3$, $p < .001$; Supplementary Table C17). Fully decomposed coefficients clarify the source of this form-level signal: prior user state is the largest predictor of immediate constructive follow-up, and M4 explaining is the most stable positive support-form coefficient (Supplementary Table C15).

All model/source checks are used as robustness analyses rather than model-causal comparisons. Model labels are not experimentally assigned. They may be confounded with calendar time, user population, task mix, interface changes and logging context. For this reason, the manuscript does not claim that a specific model family caused the observed engagement differences.

B.3 Label uncertainty

The label scheme intentionally separates top-level constructs from finer support forms. Top-level labels such as constructive engagement and scaffolded support have clear behavioural anchors and the strongest validation evidence; they support the main prevalence, association and temporal claims. Finer support-form labels are interpreted as aggregate form-pattern signals rather than isolated item-level judgments. This is why the main claims rely most heavily on constructive engagement, scaffolded support presence, support-form patterns that are consistent across settings and temporal contrasts that can be expressed with simple conditional probabilities.

Labels with lower prevalence or more ambiguous boundaries are treated more cautiously. Visible emotional expression is reported descriptively because it was rare. Active engagement helps describe the engagement ecology, but the central outcome is constructive engagement. Support-form labels M1–M6 can co-occur and are not interpreted as mutually exclusive exposures. The human–human and human–LLM validation tables therefore support the use of the schema at the aggregate level, while the manuscript avoids claims that depend on adjudicating a single ambiguous turn in isolation.

B.4 Selection into support

The most important threat to causal interpretation is selection into support. More complex tasks, more motivated users, better-specified questions or already-developing exchanges may both elicit scaffolded support and produce constructive follow-up. This possibility is part of the interactional process that the paper studies. The empirical question is therefore whether scaffolded support remains associated with constructive participation after observed context and immediate prior user state are taken into account, not whether support was randomly assigned.

Conversation-level adjusted models reduce measured compositional differences but cannot remove unobserved selection. Adjacent-turn analyses narrow the temporal window and condition on the immediately preceding user state, which helps distinguish a broad conversation-level association from a state-conditioned coupling pattern. Reverse adjacent-turn summaries also show that user state predicts subsequent assistant scaffolding, consistent with a coupled process rather than a one-directional exposure. The safest interpretation is therefore state-dependent coupling between assistant support and user engagement, not a claim that scaffolded support alone caused learning.

B.5 Privacy-preserving reporting

The study analyses naturalistic conversations, so the manuscript avoids reproducing verbatim raw conversation logs. Operational examples are de-identified and paraphrased or lightly edited to illustrate decision rules without exposing user content. Aggregate statistics, model summaries and validation metrics are reported in the manuscript and supplementary tables. The analyses do not link conversations into user histories or estimate longitudinal user-level trajectories; temporal claims are restricted to within-conversation turn ordering.

Table C1 Operational examples for turn-level labels. Examples are de-identified and paraphrased from validation review patterns; they illustrate the decision rule rather than reproducing raw conversation text.

Construct	Paraphrased turn	Annotation rationale
P passive user engagement	"Thanks, that makes sense."	Minimal uptake or acknowledgement without new reasoning, evaluation or task transformation.
A active user engagement	"Can you convert that answer into a command I can run in my setup?"	Task-relevant follow-up or manipulation of supplied information, but without elaborating or testing an idea.
C constructive user engagement	"If the cache expires before the asynchronous callback returns, would the same race condition still happen?"	The user tests an implication, boundary condition or transfer of the prior explanation.
Scaffolding	"First isolate the failing step; that will show whether the issue is parsing or state management."	The assistant structures the user's process and leaves diagnostic or revision work with the user.
M1 feedback	"Your diagnosis is close; the bug is in the index offset, not the loop boundary."	Evaluative feedback on the user's prior attempt or understanding.
M2 hinting	"Check the parser log around the first malformed token before changing the model."	A cue or partial path that helps the user proceed without fully resolving the task.
M3 instructing	"First split the condition into two branches, then add a regression test."	Explicit steps or strategy that organizes the user's work.
M4 explaining	"The exception occurs because the event loop is already running, so a nested call cannot acquire it."	A rationale or mechanism that clarifies why the issue occurs.
M5 modeling	"A small pattern you can adapt is: validate input, normalize it, then pass it to the writer."	A sample approach intended to be emulated or adapted rather than merely copied.
M6 questioning	"Which constraint matters more here: preserving tone or shortening the paragraph?"	A question that advances problem solving or reflection.

Appendix C Supporting tables

Table C2 Complete de-identified coding case with turn-level labels. The illustrative transcript is derived from a retained conversation and lightly edited while preserving turn sequence, task structure and annotation decisions. User labels are P passive, A active and C constructive.

Case overview. WildChat coding; intentional framing; SQL transaction handling. The exchange illustrates constructive user follow-up around a concrete programming mechanism and scaffolded assistant support through explanation, procedure and modelling.		
Turn	De-identified transcript	Label and annotation cue
U1 User	“How does <code>XACT_ABORT</code> work in a SQL Server script? If my delete statements are inside a transaction so failed deletes should roll back, should <code>XACT_ABORT</code> be on or off? Does it apply only to the transaction or to the whole script?”	C constructive. The user probes a mechanism, asks about scope and links the answer to rollback behaviour.
A1 Assistant	The assistant explains that <code>XACT_ABORT ON</code> makes many run-time errors abort and roll back the current transaction, distinguishes transaction-level rollback from statements outside the transaction and explains why the setting is commonly used for update/delete scripts.	Scaffolded support. I2; M4. The reply explains the mechanism connecting errors, transactions and rollback.
U2 User	“Do I need to set <code>XACT_ABORT OFF</code> at the end because of scope? What about using <code>TRY/CATCH</code> outside a transaction instead?”	C constructive. The user extends the prior answer by comparing scope and an alternative error-handling strategy.
A2 Assistant	The assistant explains the session-scoped setting, gives a cautious cleanup pattern, contrasts <code>XACT_ABORT</code> with <code>TRY/CATCH</code> and sketches a safer transaction pattern with <code>BEGIN TRY, COMMIT, BEGIN CATCH and ROLLBACK</code> .	Scaffolded support. I2; M3, M4, M5. The reply combines procedure, explanation and an adaptable model.

Table C3 Complete de-identified writing case with turn-level labels. The illustrative transcript is derived from a retained conversation and lightly edited while preserving turn sequence, task structure and annotation decisions.

Case overview. WildChat writing; intentional framing; formal report justification. The exchange illustrates a shift from direct explanation use to constructive reframing, followed by direct drafting support under a new style constraint.		
Turn	De-identified transcript	Label and annotation cue
U1 User	“Help me explain why a machine-learning dataset should include difficult samples and not just clean samples, so that the trained model will be more robust. Give me a more detailed explanation.”	A active. The user requests a detailed explanation but does not yet evaluate, revise or extend an idea.
A1 Assistant	The assistant gives a direct explanation: diverse and difficult samples expose the model to real-world variation, reduce dependence on overly clean patterns and improve generalization to noisier cases.	Non-scaffolded reference. The reply supplies the requested explanatory content without structuring the user’s next reasoning or writing process.
U2 User	“Suppose I am explaining this to someone who does not believe that a model trained only on clean data generalizes worse than one trained with diverse, difficult data. Help me explain why this is the case in more detail.”	C constructive. The user reframes the task around persuasion and justification, extending the prior explanation into a more demanding reasoning context.
A2 Assistant	The assistant explains the mechanism with an analogy and rationale: clean-only training narrows the problem space, whereas difficult examples expose edge cases, noise and variation that make the learned decision boundary less brittle.	Scaffolded support. I2; M4. The reply teaches the conceptual mechanism behind robustness rather than merely restating the claim.
U3 User	“I need a different type of explanation for a formal report. I should justify the need for a diverse dataset without relying on analogies. Can you help me do that?”	A active. The user specifies a new writing constraint and asks for a format shift, but the turn is mainly a directive refinement.
A3 Assistant	The assistant provides a formal report-style explanation linking sample diversity to generalization, deployment variation, robustness and evaluation risk.	Non-scaffolded reference. The reply directly supplies the requested formal wording rather than adding process guidance or feedback on the user’s own draft.

Table C4 Domain-stratified human–human codebook agreement for labels used in the main analyses. Coding uses the latest three-round validation rerun over 658 task turns, 643 after deduplication; writing uses the final writing-210 validation labels. These validation sets assess whether human reviewers could apply the coding and writing codebook consistently; they are not, by themselves, direct estimates of production LLM-label accuracy and are not separately powered within each public corpus. F1, Matthews correlation coefficient (MCC) and Gwet’s AC1 are reported per binary label. Coding MCC and AC1 are recomputed from the exported support and positive-count fields; writing metrics are from the final writing agreement report. Rare degenerate labels with no positive cases in a validation set are omitted. Production-label checks are reported in Supplementary Table C5, and post-update audits are described in Supplementary Section A.4.

Domain	Role	Label	F1	MCC	Gwet AC1	Support
Coding	User	Active (A)	0.75	0.72	0.87	315
Coding	User	Constructive (C)	0.85	0.82	0.90	315
Coding	User	Passive (P)	0.88	0.87	0.96	315
Coding	User	Emotional (E)	0.83	0.82	0.97	315
Coding	Assistant	Scaffolding	0.87	0.69	0.70	328
Coding	Assistant	Feedback (M1)	0.90	0.89	0.95	328
Coding	Assistant	Hinting (M2)	0.79	0.75	0.89	328
Coding	Assistant	Instructing (M3)	0.81	0.77	0.92	328
Coding	Assistant	Explaining (M4)	0.86	0.79	0.84	328
Coding	Assistant	Modeling (M5)	0.84	0.79	0.87	328
Coding	Assistant	Questioning (M6)	0.82	0.81	0.95	328
Writing	User	Active (A)	0.73	0.68	0.81	90
Writing	User	Constructive (C)	0.89	0.83	0.86	90
Writing	User	Passive (P)	0.80	0.81	0.99	90
Writing	User	Emotional (E)	0.80	0.79	0.98	90
Writing	Assistant	Scaffolding	0.97	0.80	0.93	120
Writing	Assistant	Feedback (M1)	0.89	0.80	0.80	120
Writing	Assistant	Hinting (M2)	0.82	0.70	0.71	120
Writing	Assistant	Instructing (M3)	0.71	0.67	0.86	120
Writing	Assistant	Explaining (M4)	0.84	0.70	0.70	120
Writing	Assistant	Modeling (M5)	0.94	0.91	0.95	120
Writing	Assistant	Questioning (M6)	0.96	0.96	0.99	120

Table C5 Consolidated human–LLM production-label agreement. User-framing audit sampled 450 first user turns across the six corpus-by-task settings, with production-positive intentional-framing cases deliberately oversampled. User engagement audit sampled 600 user turns, with 100 cases per setting and production constructive cases deliberately oversampled. Assistant audit sampled 180 assistant turns, with 30 cases per setting and scaffolded-support and support-form positives deliberately oversampled. Final human labels were confirmed after item-level review. Metrics compare final LLM-assisted production labels with human-confirmed labels using one-vs-rest binary coding for each label. Because the samples are positive-oversampled, these metrics validate label boundaries rather than estimating population prevalence.

Role	Label	Audit N	Prod. +	Human +	Precision	Recall	F1	MCC	Gwet AC1	Accuracy
<i>Conversation-level framing</i>										
Conversation	Intentional framing	450	300	239	0.770	0.967	0.857	0.677	0.671	0.829
<i>User engagement</i>										
User	Constructive (C)	600	300	327	0.963	0.884	0.922	0.840	0.837	0.912
User	Active (A)	600	120	111	0.850	0.919	0.883	0.856	0.935	0.958
User	Passive (P)	600	90	80	0.889	1.000	0.941	0.934	0.978	0.983
User	None	600	90	82	0.889	0.976	0.930	0.920	0.974	0.980
<i>Assistant scaffolding and support forms</i>										
Assistant	Scaffolding	180	144	139	0.896	0.928	0.912	0.590	0.791	0.861
Assistant	Feedback (M1)	180	45	51	1.000	0.882	0.938	0.918	0.945	0.967
Assistant	Hinting (M2)	180	43	45	0.930	0.889	0.909	0.880	0.930	0.956
Assistant	Instructing (M3)	180	51	55	0.745	0.691	0.717	0.600	0.715	0.833
Assistant	Explaining (M4)	180	87	110	1.000	0.791	0.883	0.772	0.747	0.872
Assistant	Modelling (M5)	180	48	79	0.958	0.582	0.724	0.631	0.642	0.806
Assistant	Questioning (M6)	180	35	20	0.543	0.950	0.691	0.675	0.873	0.906

Table C6 Analytic model families used in the main analyses. The table summarizes descriptive and conversation-level model families. Adjusted models are interpreted as covariate-adjusted associations in naturalistic data, not as causal treatment-effect estimates.

Analysis family	Unit and outcome	Main contrast or model	Scope
Engagement prevalence	User turns and conversations; cognitive, constructive and emotional engagement indicators.	Descriptive proportions by task setting, task ecology and user framing.	Establishes where learning-oriented engagement is visible; no causal contrast is implied.
Engagement composition	Cognitively engaged user turns; passive, active and constructive depth levels.	Composition of cognitive engagement within each task setting.	Motivates constructive engagement as the focal high-level outcome while preserving the broader cognitive-engagement hierarchy.
Interaction depth	Conversations; at least one constructive user turn.	Probability of constructive engagement across behavioural-depth buckets.	Depth is interpreted as interactional opportunity and task development, not as a manipulated exposure.
Scaffolding presence	Conversations; constructive-turn count and any constructive turn.	Setting-specific Poisson GLM for constructive-turn count with scaffolded-support presence and observed conversation-level context; no user-turn offset in the primary count model. Setting-specific logistic model for whether any constructive user turn occurs. Estimable covariates are listed in Supplementary Table C8.	Adjusted estimates reduce measured compositional differences but remain associational because support is selected into by task difficulty and user state. Pearson overdispersion was checked and quasi-Poisson scaling preserved the direction of the Poisson estimates.

Table C7 Additional analytic model families and robustness checks. These analyses qualify the main associations by support form, temporal scale and available model or source metadata.

Analysis family	Unit and outcome	Main contrast or model	Scope
Support-form associations	Scaffolded conversations; constructive engagement ratios.	Mean differences for conversations containing a given support-means label versus scaffolded conversations without that label, summarized by user framing and task ecology.	Descriptive because means labels are non-exclusive and can co-occur in the same assistant turn.
Local temporal coupling	Adjacent assistant–user or user–assistant turn pairs; next constructive user turn or next scaffolded assistant turn.	Conditional probability contrasts around adjacent turns; detailed estimands are listed in Supplementary Table C9.	Interpreted as local ordering and state-dependent coupling, not as randomized support effects.
Integrated adjacent-turn regression	Adjacent assistant–user turn pairs; next constructive user turn.	Logistic regression combining prior user state, scaffolded-support presence, support means, task/framing context, prior-state $\times$ support-form interactions and dataset or model/source fixed effects.	Tests whether the local coupling pattern depends jointly on user state and support form; standard errors are clustered by conversation.
Coarse temporal boundary checks	Conversations containing scaffolded support; constructive engagement before and after the first scaffolded turn.	Mean after-minus-before contrast reported as an exploratory appendix check.	Not used as evidence of accumulation because the first scaffolded turn may be selected by difficulty, uncertainty or an already developing exchange.
WildChat model robustness	WildChat conversations or turns; main prevalence, association and temporal outcomes.	Repetition across user-facing model snapshots or coarse model families.	Robustness summaries only; model labels are observational and may be confounded with time, user mix and task mix.

Table C8 Setting-specific covariate sets for the Section 2.3 adjusted models. Each model was fitted separately within the six task settings. Task ecology is therefore fixed within a setting and is not interpreted as a within-setting coefficient. The table lists the estimable covariates used to adjust the scaffolded-support association; pooled context models use task ecology separately as shown in Supplementary Tables C12 and C13.

Settings	Primary Poisson count model	Logistic any-constructive model	Sensitivities and notes
WC coding; LMSYS coding; SC coding; WC writing; LMSYS writing; SC writing	Scaffolded-support presence; user framing; total turns; emotional-expression ratio; observable error marker; persistence-after-failure marker; high-persistence marker.	Scaffolded-support presence; user framing; total turns.	The offset-rate sensitivity used the same Poisson mean model but represented user turns with an exposure offset. User-framing stratified Poisson models omitted user framing because it was fixed within stratum. Error and persistence markers are most interpretable in coding-oriented contexts but are retained as observed conversation-level context when present.

Table C9 Temporal estimands for adjacent-turn support–engagement coupling. Main temporal analyses use adjacent turn pairs. Coarser before–after summaries are exploratory boundary checks because scaffolded turns are selected by task difficulty, user motivation and conversation trajectory.

Estimand	Definition	Reported in main text	Scope
Overall adjacent-turn lift	Next-turn constructive probability after scaffolded support versus a non-scaffolded reference response, estimated within each task setting.	Figure 5a	Local association between support type and immediately following constructive engagement; not a randomized support effect.
Prior-state adjacent contrast	Same next-turn contrast, stratified by prior constructive, active and passive user states.	Figure 5b	Tests whether local coupling depends on the user's immediately preceding state.
Reverse state-conditioned support	Probability that the next assistant turn is scaffolded after constructive, active and passive user turns.	Figure 5c	Describes how scaffolding supply varies with the user's current engagement state.
Prior-state support-form regression	Adjacent-turn logistic regression for next-turn constructive engagement, adding prior-state $\times$ M1–M6 interactions.	Figure 5d and Supplementary Table C17	Tests whether local coupling depends jointly on prior user state and support form.
Around-first-scaffold contrast	Mean constructive engagement after versus before the first scaffolded assistant turn, among conversations containing scaffolded support.	Supplementary Figure D1	Exploratory boundary check only; the first scaffolded turn may be selected by difficulty or an already developing exchange.

Table C10 Uncertainty estimates for key support–engagement contrasts.

Constructive-ratio contrasts are turn-weighted conversation-bootstrap estimates, with ratios computed as total constructive user turns divided by total user turns within scaffolded versus non-scaffolded conversation groups. Post-answer depth contrasts use conversation-level bootstrap resampling. Adjacent-turn contrasts bootstrap conversations over assistant-to-user pairs. Percentage-point and turn differences are the effect sizes reported in the main figures.

Setting	Contrast	Effect	95% CI	p value
WC coding	Scaffolded – reference constructive ratio	+4.56 pp	[+4.15, +4.99]	< .001
WC coding	Scaffolded – reference post-answer depth	+2.23 turns	[+2.11, +2.35]	< .001
LMSYS coding	Scaffolded – reference constructive ratio	+2.50 pp	[+2.19, +2.83]	< .001
LMSYS coding	Scaffolded – reference post-answer depth	+1.42 turns	[+1.34, +1.52]	< .001
SC coding	Scaffolded – reference constructive ratio	+7.34 pp	[+5.29, +9.27]	< .001
SC coding	Scaffolded – reference post-answer depth	+3.36 turns	[+2.69, +4.01]	< .001
WC writing	Scaffolded – reference constructive ratio	+0.99 pp	[+0.82, +1.16]	< .001
WC writing	Scaffolded – reference post-answer depth	+3.13 turns	[+2.96, +3.30]	< .001
LMSYS writing	Scaffolded – reference constructive ratio	+1.09 pp	[+0.87, +1.31]	< .001
LMSYS writing	Scaffolded – reference post-answer depth	+2.10 turns	[+1.91, +2.28]	< .001
SC writing	Scaffolded – reference constructive ratio	+3.20 pp	[+1.55, +4.89]	< .001
SC writing	Scaffolded – reference post-answer depth	+3.38 turns	[+1.92, +4.64]	< .001
WC coding	Adjacent scaffolded – reference lift	+2.68 pp	[+2.22, +3.09]	< .001
LMSYS coding	Adjacent scaffolded – reference lift	+2.01 pp	[+1.54, +2.50]	< .001
SC coding	Adjacent scaffolded – reference lift	+6.21 pp	[+4.35, +8.10]	< .001
WC writing	Adjacent scaffolded – reference lift	+1.13 pp	[+0.88, +1.41]	< .001
LMSYS writing	Adjacent scaffolded – reference lift	+0.27 pp	[−0.01, +0.56]	.061
SC writing	Adjacent scaffolded – reference lift	+3.35 pp	[+0.80, +6.62]	.025

Table C11 Offset-rate sensitivity for scaffolded-support Poisson models. The primary Fig. 3b Poisson model uses raw constructive-turn counts with conversation length entered as a covariate. As a rate sensitivity, this table refits the same setting-specific mean model with `offset(log(user_turns))`. Cells report the scaffolded-support rate ratio with 95% confidence intervals and two-sided p values using quasi-Poisson scaled standard errors.

Setting	Offset rate ratio [95% CI], p value	Pearson dispersion	Conversations
WC coding	1.57 [1.47, 1.68], < .001	1.29	31,878
LMSYS coding	1.47 [1.38, 1.56], < .001	1.13	32,114
SC coding	1.58 [1.31, 1.91], < .001	1.29	2,481
WC writing	1.36 [1.26, 1.46], < .001	1.15	39,534
LMSYS writing	1.70 [1.50, 1.94], < .001	1.12	21,023
SC writing	1.65 [1.27, 2.15], < .001	1.32	1,539

Table C12 Conversation-level context model for constructive engagement. The outcome is whether a conversation contains at least one constructive user turn. The logistic regression includes user framing, task ecology, conversation-length bucket and dataset fixed effects. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-robust standard errors. The reference category is unintentional framing, writing-oriented task, 2–3 user turns and WildChat.

Predictor	OR [95% CI], p value
Intentional framing	3.03 [2.93, 3.13], < .001
Coding task ecology	1.89 [1.71, 2.08], < .001
4–6 user turns	2.07 [1.99, 2.15], < .001
7+ user turns	3.58 [3.42, 3.74], < .001
LMSYS Chat	0.70 [0.68, 0.73], < .001
ShareChat	2.29 [2.12, 2.46], < .001
Conversations	128,569

Table C13 Conversation-level constructive-rate sensitivity model. The outcome is the number of constructive user turns out of total user turns in each conversation, modelled with a grouped-binomial logistic regression. The model includes user framing, task ecology, conversation-length bucket and dataset fixed effects, using user-turn counts as denominators to reduce the mechanical opportunity effect in any-constructive-turn outcomes. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-robust standard errors. The reference category is unintentional framing, writing-oriented task, 2–3 user turns and WildChat.

Predictor	OR [95% CI], p value
Intentional framing	2.78 [2.68, 2.88], < .001
Coding task ecology	1.66 [1.47, 1.88], < .001
4–6 user turns	1.13 [1.09, 1.18], < .001
7+ user turns	1.06 [1.01, 1.11], .010
LMSYS Chat	0.72 [0.69, 0.75], < .001
ShareChat	2.30 [2.13, 2.50], < .001
Conversations	128,569
User turns	491,685
Constructive user turns	24,227

Table C14 Setting-level integrated adjacent-turn models. Each column reports a separate within-setting logistic regression for whether the next user turn is constructive. The scaffolded-support row comes from a broad scaffolded-support model. Support-form rows come from a decomposed model that includes broad scaffolded-support presence and co-occurring M1–M6 forms, so form coefficients compare variation within scaffolded support rather than replacing the broad scaffolded-support contrast. All models include prior user state, intentional framing, assistant-turn index and recoverable model/source fixed effects where available. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors.

Predictor	WC coding	LMSYS coding	SC coding	WC writing	LMSYS writing	SC writing
A2U pairs	92,042	71,910	9,001	124,214	55,580	5,825
Conversations	31,809	31,606	2,474	39,455	20,936	1,537
Model/source FE count	13	23	4	13	23	3
Scaffolded support	1.13 [1.08, 1.19] < .001	1.25 [1.16, 1.34] < .001	1.24 [1.09, 1.42] < .001	1.57 [1.38, 1.78] < .001	1.04 [0.86, 1.26] .679	1.34 [0.98, 1.85] .068
Prior user constructive	5.20 [4.84, 5.59] < .001	6.01 [5.43, 6.66] < .001	4.13 [3.46, 4.93] < .001	8.28 [6.73, 10.19] < .001	11.89 [9.10, 15.54] < .001	6.23 [4.43, 8.76] < .001
Prior user active	1.67 [1.57, 1.76] < .001	1.58 [1.46, 1.71] < .001	1.56 [1.32, 1.83] < .001	1.91 [1.68, 2.18] < .001	2.18 [1.80, 2.64] < .001	1.47 [1.07, 2.01] .017
Prior user passive	2.19 [1.65, 2.92] < .001	2.01 [1.38, 2.92] < .001	1.82 [1.25, 2.66] .002	2.49 [1.69, 3.65] < .001	2.32 [1.29, 4.17] .005	0.86 [0.40, 1.84] .703
Intentional framing	1.57 [1.49, 1.66] < .001	1.46 [1.35, 1.58] < .001	1.57 [1.34, 1.83] < .001	1.02 [0.89, 1.16] .798	0.86 [0.70, 1.06] .150	1.37 [0.93, 2.02] .107
M1 feedback	1.19 [1.07, 1.31] .001	1.21 [0.98, 1.50] .079	1.28 [1.03, 1.60] .025	0.69 [0.53, 0.91] .008	0.81 [0.51, 1.31] .397	1.74 [1.12, 2.71] .013
M2 hinting	1.04 [0.95, 1.14] .429	0.90 [0.72, 1.11] .332	0.85 [0.65, 1.10] .209	1.52 [1.20, 1.93] < .001	1.63 [1.09, 2.44] .018	0.76 [0.45, 1.29] .314
M3 instructing	0.89 [0.83, 0.96] .003	0.84 [0.69, 1.04] .103	0.93 [0.77, 1.13] .491	1.36 [1.04, 1.78] .027	1.06 [0.71, 1.58] .779	0.94 [0.45, 1.94] .860
M4 explaining	1.21 [1.09, 1.35] < .001	1.36 [1.08, 1.72] .008	1.77 [1.35, 2.33] < .001	1.81 [1.40, 2.34] < .001	1.12 [0.75, 1.68] .570	1.09 [0.58, 2.06] .789
M5 modelling	0.86 [0.80, 0.93] < .001	0.73 [0.63, 0.86] < .001	0.76 [0.60, 0.96] .021	0.69 [0.47, 1.01] .055	1.04 [0.64, 1.68] .872	1.57 [0.94, 2.61] .084
M6 questioning	0.73 [0.62, 0.86] < .001	0.30 [0.21, 0.44] < .001	0.70 [0.54, 0.90] .006	1.13 [0.74, 1.72] .567	0.66 [0.36, 1.21] .179	1.66 [0.90, 3.08] .106

Table C15 Integrated adjacent-turn regression combining user state, assistant scaffolding and model controls. The outcome is whether the next user turn is constructive. The scaffolded-support row is estimated in a broad scaffolded-support model; M1, M4 and M6 rows are estimated in a decomposed support-form model that includes broad scaffolded-support presence and co-occurring support forms. Estimates are odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors. The primary pooled model uses dataset fixed effects; the model/source sensitivity adds recoverable assistant model or source-family fixed effects.

Predictor	Pooled dataset FE	Pooled model/source FE	WildChat model FE
Scaffolded support	1.50 [1.44, 1.55], < .001	1.47 [1.41, 1.52], < .001	1.56 [1.49, 1.63], < .001
Prior user constructive	8.90 [8.45, 9.38], < .001	8.55 [8.11, 9.01], < .001	9.11 [8.53, 9.74], < .001
Prior user active	2.62 [2.52, 2.72], < .001	2.57 [2.47, 2.67], < .001	2.85 [2.71, 3.00], < .001
Prior user passive	2.08 [1.76, 2.44], < .001	2.07 [1.76, 2.44], < .001	2.32 [1.85, 2.90], < .001
Intentional framing	1.40 [1.35, 1.46], < .001	1.42 [1.37, 1.48], < .001	1.40 [1.33, 1.47], < .001
Coding task	1.37 [1.22, 1.54], < .001	1.37 [1.22, 1.54], < .001	1.21 [1.05, 1.40], .009
M1 feedback	0.79 [0.73, 0.86], < .001	0.80 [0.74, 0.86], < .001	0.74 [0.67, 0.82], < .001
M4 explaining	2.15 [1.99, 2.32], < .001	2.12 [1.96, 2.30], < .001	2.11 [1.92, 2.31], < .001
M6 questioning	0.82 [0.72, 0.93], .002	0.84 [0.74, 0.96], .008	1.08 [0.93, 1.25], .340
Model/source block	Reference	LR $\chi^2_{41} = 714.4$, $p < .001$	LR $\chi^2_{13} = 553.2$, $p < .001$

Table C16 Incremental contribution of scaffolding features in integrated adjacent-turn models. Likelihood-ratio block tests compare nested logistic regressions for whether the next user turn is constructive. All models include prior user state, user framing, task ecology, assistant-turn index and the indicated dataset or model/source fixed effects.

Scope	Added feature block	LR χ^2 (df)	p value
six task settings, dataset FE	Broad scaffolded support after user/context controls	584.6 (1)	< .001
six task settings, dataset FE	M1–M6 descriptors within scaffolded support	896.7 (6)	< .001
six task settings, dataset FE	Full scaffolding block after user/context controls	1481.3 (7)	< .001
six task settings, model/source FE	Broad scaffolded support after user/context controls	509.7 (1)	< .001
six task settings, model/source FE	M1–M6 descriptors within scaffolded support	829.8 (6)	< .001
six task settings, model/source FE	Full scaffolding block after user/context controls	1339.6 (7)	< .001
WildChat only, model FE	Broad scaffolded support after user/context controls	418.6 (1)	< .001
WildChat only, model FE	M1–M6 descriptors within scaffolded support	519.4 (6)	< .001
WildChat only, model FE	Full scaffolding block after user/context controls	937.9 (7)	< .001

Table C17 Prior user state and all six support forms jointly structure adjacent-turn constructive engagement. The first panel reports likelihood-ratio tests adding prior-state $\times$ M1–M6 interactions to the integrated adjacent-turn logistic model. The second panel reports pooled model/source fixed-effect estimates for each support form within each prior user state. Estimates are odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors.

Panel	Scope / prior state	Term or test	Estimate
Interaction block	six task settings, dataset FE	Prior state $\times$ M1–M6	LR $\chi^2_{18} = 305.4$, < .001
Interaction block	six task settings, model/source FE	Prior state $\times$ M1–M6	LR $\chi^2_{18} = 297.3$, < .001
Interaction block	WildChat only, model FE	Prior state $\times$ M1–M6	LR $\chi^2_{18} = 186.0$, < .001
State-stratified model/source FE	prior constructive	M1 feedback	0.72 [0.64, 0.81], < .001
	prior constructive	M2 hinting	0.91 [0.79, 1.04], .155
	prior constructive	M3 instructing	0.92 [0.81, 1.05], .236
	prior constructive	M4 explaining	1.72 [1.45, 2.04], < .001
	prior constructive	M5 modelling	0.81 [0.70, 0.93], .003
	prior constructive	M6 questioning	0.95 [0.72, 1.25], .694
	prior active	M1 feedback	0.83 [0.72, 0.95], .008
	prior active	M2 hinting	1.14 [1.03, 1.27], .016
	prior active	M3 instructing	0.98 [0.89, 1.08], .722
	prior active	M4 explaining	2.13 [1.88, 2.41], < .001
	prior active	M5 modelling	0.83 [0.76, 0.91], < .001
	prior active	M6 questioning	0.75 [0.61, 0.93], .009
	prior passive	M1 feedback	2.16 [1.16, 4.04], .015
	prior passive	M2 hinting	1.79 [0.95, 3.37], .071
	prior passive	M3 instructing	1.52 [0.67, 3.43], .312
	prior passive	M4 explaining	2.88 [1.55, 5.35], < .001
	prior passive	M5 modelling	0.46 [0.20, 1.04], .061
	prior passive	M6 questioning	2.19 [1.12, 4.27], .021

Table C18 Support-intent labels are dominated by cognitive support. Scaffolded assistant turns can receive metacognitive (I1), cognitive (I2) or affective (I3) intent labels. Scaffolded-turn prevalence reports the share of scaffolded assistant turns containing each intent label. Constructive ratios compare scaffolded conversations containing at least one assistant turn with the intent label with scaffolded conversations without that intent label, using turn-weighted constructive user-turn ratios. Intent labels can co-occur, and affective intent is sparse, so these contrasts are interpreted descriptively.

Scope	Intent label	Scaffolded turns (%)	With intent (%)	Without intent (%)	Diff. (pp)	Conv. with / without
Pooled	I1 metacognitive	8.7	5.0	6.9	-2.0	10,924 / 55,892
Pooled	I2 cognitive	91.6	6.8	2.8	+4.0	61,942 / 4,874
Pooled	I3 affective	0.8	12.4	6.4	+6.0	1,052 / 65,764
Coding	I1 metacognitive	8.2	7.0	9.6	-2.6	6,782 / 36,179
Coding	I2 cognitive	92.3	9.4	3.9	+5.6	40,279 / 2,682
Coding	I3 affective	1.0	16.1	8.9	+7.2	778 / 42,183
Writing	I1 metacognitive	9.6	2.5	2.8	-0.3	4,142 / 19,713
Writing	I2 cognitive	90.4	2.8	1.8	+0.9	21,663 / 2,192
Writing	I3 affective	0.6	3.9	2.7	+1.2	274 / 23,581

Table C19 Support-supply profile within scaffolded assistant turns. Values show the percentage of scaffolded assistant turns containing each scaffolding-means label. Support-form labels are not mutually exclusive, so rows do not sum to 100%. Coding-oriented scaffolded turns are dominated by explanation-heavy support, whereas writing-oriented scaffolded turns show a more mixed support profile.

Setting	M1 Feedback	M2 Hinting	M3 Instructing	M4 Explaining	M5 Modelling	M6 Questioning
WildChat coding	8.8	14.7	29.7	84.0	23.9	6.9
LMSYS coding	4.4	7.6	9.7	76.7	16.5	13.9
ShareChat coding	13.3	12.7	25.7	81.2	18.4	18.6
WildChat writing	16.6	9.6	59.2	29.5	15.0	6.9
LMSYS writing	11.3	10.9	52.3	26.9	10.3	13.4
ShareChat writing	24.2	23.2	22.9	45.2	23.5	38.1

Appendix D Supporting figures

The supporting figures separate cross-dataset extensions from model-snapshot checks. Supplementary Figures D1 and D2 extend the main analyses across the six WildChat, LMSYS Chat and ShareChat task settings. Supplementary Figure D3 reports the support-form supply profile across the same six settings, and Supplementary Figure D4 provides paraphrased support-form cards for the six scaffolded-support forms. Supplementary Figures D5–D9 are a separate WildChat-only model-snapshot robustness layer. They are not substitutes for the six-setting checks; they ask whether the WildChat portion of the findings is concentrated in a particular user-facing model snapshot. LMSYS Chat and ShareChat remain represented in the six-setting figures and in the model/source-adjusted regressions, but they do not provide directly comparable user-facing model-snapshot structure for the same visual repetition. All model-snapshot figures are descriptive because model labels in naturalistic corpora are observational rather than experimentally assigned.




Fig. D1 Coarse around-first-scaffold contrasts are exploratory boundary checks across the six task settings. Bars show the difference in constructive engagement after versus before the first scaffolded assistant turn within conversations containing scaffolded support. Error bars show 95% conversation-bootstrap confidence intervals. These before–after contrasts are not used as treatment-effect or accumulation estimates because the first scaffolded turn may be selected by difficulty, uncertainty or an already developing exchange.




Fig. D2 Scaffolded-support associations by user framing across the six task settings. Stratified Poisson count-ratio estimates compare the association between scaffolded-support presence and constructive participation in intentional versus unintentional conversations. Error bars show 95% model-based confidence intervals. The figure is interpreted as descriptive moderation rather than evidence that framing causally changes support effectiveness, because user framing, task complexity and persistence are not randomly assigned.

Support-form supply within scaffolded assistant turns (%)						
WC coding	8.8	14.7	29.7	84.0	23.9	6.9
LMSYS coding	4.4	7.6	9.7	76.7	16.5	13.9
SC coding	13.3	12.7	25.7	81.2	18.4	18.6
WC writing	16.6	9.6	59.2	29.5	15.0	6.9
LMSYS writing	11.3	10.9	52.3	26.9	10.3	13.4
SC writing	24.2	23.2	22.9	45.2	23.5	38.1
	M1	M2	M3	M4	M5	M6

WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat. Support-form labels are non-exclusive.

Fig. D3 Support-form supply within scaffolded assistant turns across the six task settings. Cells report the percentage of scaffolded assistant turns containing each non-exclusive support-form label. Rows do not sum to 100% because a single scaffolded turn can contain multiple support forms. The figure provides task-context background for the support-form association contrasts in Fig. 4.

M1 feedback <i>Evaluates a user’s prior attempt</i> “Your diagnosis is close: the race is not in the callback itself, but in when the cache key is reused. Check that key before changing the loop.” Role: Returns the user’s attempt as something to repair, revise or extend.	M2 hinting <i>Offers a partial cue</i> “Before rewriting the parser, inspect the first malformed token and compare it with the token immediately before it.” Role: Leaves solution work open while directing attention to a useful diagnostic path.
M3 instructing <i>Provides procedural steps</i> “First isolate the failing branch, then add a regression test, and only then refactor the duplicated condition.” Role: Organizes the user’s next actions through explicit sequencing or strategy.	M4 explaining <i>Makes a mechanism visible</i> “The exception occurs because the event loop is already running; a nested call asks for a resource that cannot be acquired twice.” Role: Provides a rationale that can be tested, adapted or transferred.
M5 modelling <i>Shows an adaptable pattern</i> “A reusable structure is: validate input, normalize fields, then pass the cleaned object to the writer. You can adapt this skeleton to your case.” Role: Supplies an example or pattern intended for adaptation rather than only completion.	M6 questioning <i>Elicits a constraint or reflection</i> “Which constraint matters more here: preserving the author’s tone or shortening the paragraph? That choice should determine the revision strategy.” Role: Prompts the user to specify priorities, constraints or next reasoning steps.

Fig. D4 Representative support-form cards for the six scaffolded-support forms. Cards show de-identified, paraphrased assistant replies that illustrate the annotation decision rule for each support-form label rather than reproducing raw conversation text. Support-form labels are non-exclusive, so a single assistant turn can receive more than one form label.




Fig. D5 WildChat model-snapshot prevalence profiles preserve the main domain contrast. Across user-facing WildChat model snapshots, coding-oriented conversations generally show higher cognitive engagement and higher assistant scaffolding rates than writing-oriented conversations. This is a WildChat-only model-snapshot check; cross-dataset prevalence and association checks are reported separately across all six settings. The model labels are observational and may reflect time, user mix and traffic composition, so the figure is not a model-causal comparison.




Fig. D6 Intentionality gaps are visible across WildChat model snapshots. The figure reports model-specific differences between intentional and unintentional conversations in cognitive and constructive engagement, separately for WildChat coding and writing. It complements the six-setting user-framing analyses by checking whether the WildChat pattern is concentrated in a particular model snapshot.




Fig. D7 Conversation-level support associations are broadly positive across WildChat model snapshots. The figure repeats major conversation-level analyses by user-facing WildChat model snapshot, including constructive participation, adjusted support associations and exploratory coarse before–after contrasts around the first scaffolded assistant turn. It is a within-WildChat robustness check; cross-dataset scaffolded-support contrasts are reported across all six settings in the supporting tables and main figures.




Fig. D8 Local temporal signals vary across WildChat model snapshots. The figure reports adjacent-turn constructive lifts and emotion-to-support responsiveness by WildChat model snapshot. Adjacent-turn constructive lifts are often positive, especially in coding-oriented conversations, but the magnitude varies across snapshots; emotion-to-support responsiveness is more heterogeneous and is treated as exploratory.




Fig. D9 Coarse WildChat model-family aggregation preserves the main directional patterns. WildChat model snapshots are grouped into coarse model families to summarize cognitive engagement prevalence, adjusted support associations and adjacent-turn constructive lift. Coding remains higher than writing across families, and the support-association and adjacent-turn lift patterns remain directionally consistent on average.