

Informal learning in everyday human–LLM interaction

First Author^{1,2*}, Second Author^{2,3†} and Third Author^{1,2†}

^{1*}Department, Organization, Street, City, 100190, State, Country.

²Department, Organization, Street, City, 10587, State, Country.

³Department, Organization, Street, City, 610101, State, Country.

*Corresponding author(s). E-mail(s): iauthor@gmail.com;
Contributing authors: iauthor@gmail.com; iiiauthor@gmail.com;

[†]These authors contributed equally to this work.

Abstract

As LLM-based assistants become embedded in work and study, human–LLM interaction is becoming a consequential site of informal learning. Learning science links deeper overt engagement to stronger learning outcomes, making engagement a useful observable process indicator in informal, user-generated interactions whose goals are not pre-specified by curricula or study protocols. We analyse 128,569 naturalistic conversations from WildChat, LMSYS Chat and ShareChat, translating learning-science constructs into turn-level labels for 981,470 turns using a human-reviewed LLM-assisted workflow. Across 491,685 user turns, cognitive engagement appeared in 31.9% and constructive engagement—the deepest observable form of user sense-making—in 4.9%; scaffolded support appeared in 31.7% of assistant turns. Scaffolded support marked richer constructive participation across settings, but not as a uniform exposure: its association depended on user framing, task ecology, support form, timing and prior user state. Feedback-like and explanatory forms aligned most clearly with constructive engagement, and adjacent-turn analyses showed local, state- and form-dependent coupling rather than conversation-wide accumulation. These findings make behavioural signatures of informal learning observable at scale and recast AI assistance around cognitive responsibility: whether systems merely accelerate answer delivery or also preserve opportunities for people to reason, test ideas and build understanding while solving real problems.

Keywords: human–AI interaction, informal learning, large language models, scaffolding, learner engagement

1 Introduction

Large language models (LLMs) are becoming embedded in the cognitive work of everyday life. People use them to write, code, search, troubleshoot and make sense of unfamiliar problems, often outside classrooms, training programmes or purpose-built tutoring systems [1–3]. In such moments, the interaction can become more than task completion: assistant responses may create occasions for users to clarify uncertainties, test assumptions, revise ideas and decide how to proceed [4, 5]. Everyday human–LLM interaction therefore constitutes a rapidly expanding setting for informal learning, in which understanding can emerge through self-directed inquiry and problem solving rather than formal curricula [6–8].

Yet existing research on LLMs and learning has focused primarily on settings where learning support is deliberately designed, such as automated feedback, question generation, retrieval practice and tutoring systems [9–13]. These studies show that LLMs can enact instructional strategies when learning is the explicit design target [14, 15]; they tell us less about general-purpose LLM use, where users usually seek task progress rather than instruction. This gap matters because the educational value of everyday AI use depends not only on what models provide, but on what users do with that support: whether an answer becomes material for questioning, testing, revision and extension, or whether it simply accelerates completion while leaving little visible user sense-making [16–20]. Learning science provides a way to make this distinction measurable: deeper forms of overt engagement are associated with stronger learning outcomes, so engagement can serve as an observable process indicator in informal, user-generated interactions whose goals and endpoints are not pre-specified by curricula or study protocols [21, 22]. We therefore ask whether learning-oriented engagement is visible in everyday human–LLM interaction, where its deepest constructive form appears, and how assistant scaffolding is coupled with further constructive participation.

We analyse 128,569 naturalistic conversations from three public conversational corpora: WildChat, LMSYS Chat and ShareChat [23–25]. Across these corpora, we focus on coding and writing as two common task ecologies in which LLMs are used to produce, revise and troubleshoot knowledge work [1, 4, 26], yielding 491,685 user turns and 489,785 assistant turns. Our analyses move from engagement prevalence, to the contextual organization of constructive engagement, to how assistant scaffolding is associated with user participation at conversation-level, support-form and adjacent-turn scales. To make these questions measurable in naturalistic logs, we translate theories of engagement and scaffolding into role-specific, turn-level behavioural signatures: user turns are coded for passive, active or constructive uptake [21, 22], while assistant turns are coded for scaffolded support and further characterized by support intent and support form [27–30]. We treat constructive engagement as the deepest observable indicator of learning-oriented participation available in these logs, using it to examine how deeper user sense-making is organized by task context, assistant support and local turn ordering.

The results trace a progression from the presence of learning-oriented engagement, to the conditions under which its constructive form appears, to the ways assistant support is coupled with further user participation. First, learning-oriented engagement was recurrent in everyday human–LLM interaction, but its deepest constructive




Fig. 1 Analytic framework for studying informal learning in everyday human-LLM interaction. **a**, Public conversations are situated by data source, task ecology and user framing. **b**, Engagement and scaffolding constructs are translated into role-specific turn-level labels for user engagement and assistant support. **c**, These labels support analyses of engagement presence, contextual organization and support-engagement coupling at conversation-level, support-form and adjacent-turn scales. Labels indicate observable learning-oriented participation, not durable learning outcomes or causal instructional effects.

form was selective: cognitive engagement appeared in nearly one third of user turns, whereas constructive engagement appeared in about five percent. Second, constructive engagement was ecologically organized: it was more common under explicit learning-oriented user framing, varied across task ecologies and became most visible in sustained exchanges where users returned to, interrogated and reshaped the assistant’s response. Third, assistant scaffolding marked richer constructive participation, but not as a uniform exposure. Its association depended on how support was positioned within the exchange: feedback-like and explanatory forms aligned most clearly with constructive engagement at the conversation level, while adjacent-turn analyses showed local, state- and form-dependent coupling rather than evidence that scaffolding accumulated uniformly across conversations. Together, these findings locate everyday human-LLM interaction in a middle space between ordinary tool use and designed instruction. They suggest that the educational value of future AI assistants lies not in making every interaction tutor-like, but in calibrating when to answer, explain, give feedback or step back, preserving the cognitive responsibility through which people question, test, revise and build judgment while solving real problems.

2 Results

The results follow the analytic sequence in Fig. 1: engagement prevalence, contextual organization and support-engagement coupling. We first quantify cognitive and constructive engagement in everyday human-LLM conversations, then locate constructive engagement across user framing, task ecology and interaction depth. We then turn to assistant support, asking whether scaffolded conversations contain richer constructive participation, which support forms carry this association and how support is ordered with prior and subsequent user engagement at the adjacent-turn scale.

Table 1 Dataset and behavioural overview. Summary statistics for the WildChat, LMSYS Chat and ShareChat task settings analysed in the study. Constructive engagement is reported as the deepest observable level within cognitive engagement and serves as the focal user-side measure in subsequent analyses.

Setting	Sample size			User framing	Assistant support	Cognitive engagement		Emotional engagement
	Conversations	User turns	Assistant turns	Intentional conv. (%)	Scaffolded turns (%)	Overall (%)	Constructive level (%)	Emotional (%)
WildChat coding	31,878	124,073	124,623	37.13	51.03	50.21	9.23	0.95
WildChat writing	39,534	163,705	164,824	25.80	23.53	16.34	2.22	0.57
LMSYS coding	32,114	107,065	105,174	41.49	29.07	45.26	5.44	0.91
LMSYS writing	21,023	77,906	76,279	20.64	19.43	15.30	1.53	0.21
ShareChat coding	2,481	11,541	11,522	45.63	50.67	45.00	15.17	3.38
ShareChat writing	1,539	7,395	7,363	29.69	22.42	33.05	5.18	0.34
Total / pooled	128,569	491,685	489,785	32.12	31.70	31.94	4.93	0.74

Note. User framing is reported as the percentage of conversations with explicit learning-oriented framing. Cognitive, constructive and emotional engagement are percentages of user turns; scaffolded support is computed over assistant turns. The constructive column uses all user turns as the denominator, whereas Fig. 2-a decomposes cognitively engaged user turns into passive, active and constructive levels.

2.1 Everyday human–LLM conversations contain learning-oriented engagement

Learning-oriented engagement was visible across all six task settings. Across 491,685 user turns, 31.9% showed cognitive engagement and 4.9% reached constructive engagement, the deepest observable user-side signal in the annotation scheme (Table 1). Visible emotional expression was uncommon, appearing in 0.74% of user turns. Thus, everyday human–LLM conversations often contained observable user uptake of assistant responses, but deeper constructive sense-making was selective rather than routine.

Figure 2a shows how cognitively engaged turns were composed. Active uptake was the dominant form in every setting, accounting for 62.0–86.4% of cognitively engaged turns, whereas constructive engagement accounted for 10.0–33.7% and passive receipt for 1.2–14.6%. This composition separates broad engagement from deeper sense-making. Cognitive engagement shows that everyday LLM use often involves users working with assistant output; active uptake captures use, application or follow-up; and constructive engagement identifies the cases in which users elaborate, test, revise or extend ideas. The remaining analyses therefore focus on where constructive engagement appears and how it is associated with user framing, task ecology, interaction depth and assistant support.

2.2 Constructive engagement is patterned by user framing, task ecology and sustained exchange

Having established that everyday human–LLM conversations contain learning-oriented engagement, we next examined where its constructive form became visible. Explicit learning-oriented framing provided the first contrast (Fig. 2b). In every setting, conversations framed around learning, understanding or exploration had higher cognitive and constructive user-turn rates than conversations without such framing. The cognitive-engagement lift ranged from +17.7 to +51.4 percentage points across settings, and the constructive-engagement lift ranged from +2.5 to +12.8 points. Yet constructive

participation was also present in conversations without explicit learning-oriented framing. User framing therefore amplified constructive engagement, but did not define it: learning-relevant participation also emerged in task-oriented exchanges where learning was incidental to solving the problem at hand.

Task ecology provided a second source of variation. Across all three corpora, coding-oriented conversations showed higher cognitive and constructive engagement than writing-oriented conversations. Cognitive engagement ranged from 45.00–50.21% in coding conversations, compared with 15.30–33.05% in writing conversations; constructive engagement ranged from 5.44–15.17% in coding and 1.53–5.18% in writing. A conversation-level context model showed the same direction after including user framing, conversation-length bucket and dataset fixed effects, with coding task ecology positively associated with the presence of at least one constructive user turn (OR=1.89, 95% CI 1.71–2.08, $p < .001$; Supplementary Table C9).

The coding–writing contrast points to task affordances that organize baseline opportunities for visible constructive participation. Coding tasks often externalize uncertainty through errors, constraints and executable tests, creating occasions for users to diagnose failures, refine assumptions, compare alternatives and test candidate solutions. Writing tasks also involve substantial revision and evaluation, but the interaction can often proceed through accepting, adapting or requesting a draft revision without requiring the user to articulate the rationale for the change [31–33]. The observed task difference therefore consistent with the ecological interpretation: constructive engagement is more likely to appear when the task makes uncertainty, evaluation and revision explicit in the exchange.

The distinction between broad cognitive engagement and constructive engagement reinforces this task-ecology account. LMSYS Chat coding illustrates this distinction: it had high cognitive engagement, yet its cognitively engaged turns were dominated by active uptake rather than constructive elaboration (Fig. 2a). This pattern illustrates that users can work actively with assistant output—applying code, requesting repairs or asking for the next step—without making their reasoning, tests or revisions explicit enough to surface as constructive engagement. Taken together with the coding–writing contrast, this shows that task ecology patterned not only how often users worked with assistant output, but also whether that work became visible as articulated questioning, testing, refinement or extension.

Sustained exchange provided the third context in which constructive participation became visible (Fig. 2c). We operationalized sustained exchange with conversation-length buckets, asking whether a conversation contained at least one constructive user turn. Across all six settings, conversations with more user turns were more likely to contain at least one constructive user turn. The gradient was steeper in coding-oriented settings, but the same direction appeared in writing-oriented settings. This pattern indicates that constructive engagement is most visible when everyday LLM use develops beyond a one-off request for an answer. In such exchanges, users have room to return to the assistant’s response, introduce new constraints, test implications, request revisions and reshape their understanding. Because this panel uses an any-event outcome, we treat the length gradient as a visibility pattern; a grouped-binomial




Fig. 2 Learning-oriented engagement varies by engagement form, user framing and conversation length. **a**, Composition of cognitively engaged user turns by passive, active and constructive levels across the six task settings. **b**, Cognitive and constructive user-turn rates in conversations with versus without explicit learning-oriented framing; annotations show the framed-minus-unframed difference in percentage points. **c**, Share of conversations containing at least one constructive user turn by conversation-length bucket. Panel a uses cognitively engaged user turns as the denominator; panel b reports percentages of user turns; panel c reports conversation-level shares. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.

constructive-rate sensitivity attenuated but preserved the positive length association (4–6 user turns: OR=1.13; 7+ user turns: OR=1.06; Supplementary Table C10).

Together, these results show that constructive engagement was ecologically organized rather than evenly distributed across everyday use. It was amplified by explicit learning-oriented framing, patterned by task affordances that create occasions for testing and revision, and most visible in sustained exchanges. A conversation-level context model with user framing, task ecology, conversation-length bucket and dataset fixed effects supported the same organization, and the rate sensitivity showed that the length association was attenuated but preserved after accounting for user-turn denominators (Supplementary Tables C9 and C10). Having located these user-side and contextual patterns, we next ask whether assistant scaffolding marks conversations with richer constructive participation.

2.3 Scaffolded support marks richer constructive participation

Having located the user-side and contextual conditions under which constructive engagement became visible, we next turned to assistant-side support. At the conversation level, conversations containing at least one scaffolded assistant turn had higher turn-weighted constructive user-turn ratios than reference conversations without scaffolded support in all six task settings (Fig. 3a). The constructive-ratio differences

Table 2 Summary of key support–engagement associations across the three corpora. Constructive-ratio contrasts compare conversations with versus without scaffolded support using turn-weighted constructive user-turn ratios. Adjusted models report the association of scaffolded-support presence with constructive-turn count (Poisson count ratio) and with having at least one constructive turn (logit odds ratio). Adjacent-turn lift is the difference in the probability that the next user turn is constructive after scaffolded versus non-scaffolded reference assistant turns.

Setting	Constructive ratio Has Scaf. / No Scaf.	Difference (pp)	Depth diff. (turns)	Poisson count ratio	Logit OR	Adjacent lift (pp)
WildChat coding	0.101 / 0.056	+4.6	+2.23	1.852	1.761	+2.68
LMSYS coding	0.065 / 0.040	+2.5	+1.42	1.622	1.544	+2.01
ShareChat coding	0.163 / 0.090	+7.3	+3.36	1.893	2.140	+6.21
WildChat writing	0.027 / 0.017	+1.0	+3.13	1.569	1.437	+1.13
LMSYS writing	0.022 / 0.011	+1.1	+2.10	1.985	1.764	+0.27
ShareChat writing	0.066 / 0.034	+3.2	+3.38	2.491	1.757	+3.35

ranged from +0.99 to +7.34 percentage points across WildChat, LMSYS Chat and ShareChat, and all six contrasts were distinguished from zero in conversation-level bootstrap tests ($p < .001$; Supplementary Table C7). Scaffolded conversations also showed more turns after the first assistant response in every setting (+1.42 to +3.38 turns; Fig. 3d), indicating that scaffolded support marked exchanges with more sustained participation. Together, these descriptive contrasts show that scaffolded support co-occurred with richer constructive participation across corpora and task settings.

The association remained positive after accounting for observed user framing and interactional context in covariate-adjusted models. In setting-specific models, scaffolded-support presence was associated with higher constructive-turn counts, with Poisson count ratios ranging from 1.569 to 2.491, and with higher odds that a conversation contained at least one constructive user turn, with logistic odds ratios ranging from 1.437 to 2.140 (Fig. 3b). The same positive direction appeared within both explicitly learning-oriented and non-explicitly framed conversations (Fig. 3c), indicating that the support–engagement association was not limited to interactions users announced as learning episodes. A rate sensitivity treating user turns as the exposure preserved positive scaffolded-support associations in all six settings (rate ratios 1.356–1.703; all $p < .001$; Supplementary Table C8). Across these checks, scaffolded support remained a consistent conversation-level marker of richer constructive participation. This conversation-level signal motivates a finer analysis of the support forms through which scaffolded support was expressed.

2.4 Support form differentiates the scaffolding–engagement association

The conversation-level analyses established scaffolded support as a consistent marker of richer constructive participation. We next asked which features of scaffolded support carried this signal. For each scaffolded assistant turn, the annotation scheme recorded two parallel descriptors: a support-intent label indicating whether the turn primarily targeted metacognitive, cognitive or affective support (I1–I3), and one or more support-form labels indicating how the support was delivered—feedback, hinting, instructing, explaining, modelling or questioning (M1–M6). This two-axis scheme follows scaffolding accounts that distinguish the purpose of assistance from the means




Fig. 3 Scaffolded support marks richer constructive participation across task settings. **a**, Turn-weighted constructive user-turn ratios in conversations containing scaffolded support versus reference conversations without scaffolded support. **b**, Covariate-adjusted within-setting associations between scaffolded-support presence and constructive participation, reported as Poisson count ratios and logistic odds ratios. **c**, Poisson count ratios for scaffolded-support presence stratified by explicit versus non-explicit learning-oriented framing. **d**, Difference in the number of turns after the first assistant response between scaffolded and reference conversations. Numbers indicate scaffolded-minus-reference differences, except panels b and c, which report model estimates. Error bars in panels b and c show model-based 95% confidence intervals; bootstrap confidence intervals for the descriptive contrasts in panels a and d are reported in Supplementary Table C7. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.

through which assistance is provided [27–30]. Support intent offered limited contrast because scaffolded turns were predominantly cognitive in purpose (Supplementary Table C15). We therefore focused on the form axis, comparing scaffolded conversations containing each support form with scaffolded conversations without that form and reporting turn-weighted constructive user-turn ratio differences (Fig. 4a,b).

Support form sharply differentiated the scaffolded-support association. Feedback-like support (M1) showed the largest constructive-engagement contrast in both explicitly framed and non-explicitly framed conversations, with the association strongest under explicit learning-oriented framing (+14.7 versus +6.7 percentage points; framing difference, +8.0 points; Fig. 4a). Explaining (M4) showed the same direction (+6.7 versus +3.0 points; framing difference, +3.6 points), and hinting (M2) was positive but smaller. Instructing (M3) and questioning (M6), by contrast, were lower or negative, especially in explicitly framed conversations. Task-stratified estimates gave a complementary view (Fig. 4b): feedback-like support (M1) and explaining (M4) were positive in both coding- and writing-oriented task ecologies, whereas instructing (M3), modelling (M5) and questioning (M6) varied more strongly by task.

This ordering is consistent with learning-science accounts of how assistance positions the learner. Formative feedback is most useful when it gives information about a learner’s prior attempt and supports revision, evaluation or self-regulation [34–36]. Feedback-like support (M1) therefore has a natural link to constructive engagement because it responds to something the user has already tried and can make revision or extension the next available move. Explaining (M4) makes mechanisms or rationales available for inspection, giving users material to test, adapt or connect to their own understanding [37, 38]. Hinting (M2) offers partial guidance while leaving some solution work open. Instructing (M3), modelling (M5) and questioning (M6) can also be useful forms of support, especially for task progress, examples and clarification. However, assistance-dilemma and scaffolding accounts suggest that support which makes the next step directly executable, supplies a model to imitate or manages missing information may not always require the user to generate, evaluate or revise an idea in the next move [20, 39, 40]. The conversation-level contrasts are consistent with this distinction: the forms most consistently aligned with constructive engagement were those that left an observable role for user evaluation, revision or adaptation.

The support-form supply profile places these associations in task context (Fig. 4c). Panel c is not an additional effect estimate; it shows which forms constituted the ordinary scaffolded-support environment in each task setting. Explaining (M4) dominated scaffolded coding turns across the three corpora, whereas writing-oriented support was more mixed and more often included instructing (M3), explaining (M4) or questioning (M6). Thus, coding and writing differed not only in how support forms aligned with constructive engagement, but also in which forms users typically encountered when assistant turns were scaffolded.

Together, these analyses refine the conversation-level support result: scaffolded support marked richer constructive participation, but the signal was concentrated in particular forms and embedded in task-specific support supply. This form-level result makes the temporal question sharper. If support matters partly through how it responds to the user’s ongoing work, then its local association with constructive engagement should depend on the user’s immediately preceding state. We therefore next examine support–engagement coupling at the adjacent-turn scale.

2.5 Support–engagement coupling is local, state- and form-dependent

The preceding analyses established two things: scaffolded conversations contained richer constructive participation, and this association was carried unevenly by support form. These results identify where assistant support is associated with constructive engagement, but not the interactional scale at which the association appears. We therefore turned to adjacent-turn ordering, analysing how the user’s current state, the assistant’s support form and the immediately following user turn were linked.

At this adjacent-turn scale, scaffolded assistant turns were more likely than reference assistant turns to be followed by constructive engagement in all six task settings (Fig. 5a). The next-turn constructive lift was positive in every setting, ranging from +0.27 to +6.21 percentage points. Five of the six adjacent-turn lifts were statistically distinguishable from zero in conversation-cluster bootstrap tests; LMSYS Chat




Fig. 4 Support form and support supply differentiate scaffolded interaction. **a**, Constructive-engagement differences associated with each support form among scaffolded conversations, stratified by intentional versus unintentional user framing. **b**, The same support-form contrasts stratified by coding- versus writing-oriented task ecology. **c**, Percentage of scaffolded assistant turns containing each support form by task setting. In panels a and b, bars are turn-weighted constructive user-turn ratio differences, comparing scaffolded conversations containing a support form with scaffolded conversations without that form; error bars show 95% confidence intervals, and bracket annotations show between-stratum differences and Benjamini-Hochberg FDR-adjusted q values within each panel family. Support-form labels correspond to M1–M6 and are not mutually exclusive, so rows do not sum to 100%.

writing was positive but weaker (+0.27 percentage points, 95% CI -0.01 to +0.56, $p = .061$; Supplementary Table C7). The temporal signal was therefore consistently positive across settings, but heterogeneous in magnitude.

This local coupling was organized by the user’s immediately preceding state. State-conditioned contrasts were generally positive but varied across prior constructive, active and passive turns (Fig. 5b). The user-to-assistant direction showed the same contingency: constructive and active user turns were more often followed by scaffolded support than passive turns, especially in coding-oriented conversations (Fig. 5c). These two directions suggest a coupled sequence rather than a one-way support pattern: engaged user states more often preceded scaffolded support, and scaffolded support was then more often followed by further constructive participation.

To test whether support form further structured this local coupling, we fitted integrated adjacent-turn regressions with next-turn constructive engagement as




Fig. 5 Support-engagement coupling is local, state- and form-dependent across task settings. **a**, Next-turn constructive lift after scaffolded versus reference assistant turns. **b**, Probability of a constructive next user turn by prior user state and assistant support condition; annotations show scaffolded-minus-reference differences. **c**, Probability that the next assistant turn contains scaffolded support by prior user state. **d**, Pooled model/source fixed-effect odds ratios for selected support forms within prior user states. M1 feedback and M4 explaining are shown because Section 2.4 identified them as focal positive conversation-level forms; full M1–M6 interaction estimates are reported in Supplementary Table C14. Error bars show 95% confidence intervals where displayed. Panel d reports adjacent-turn logistic models with conversation-cluster robust standard errors. WC, WildChat; LMSYS, LMSYS Chat; SC, ShareChat.

the outcome. These models combined prior user state, assistant-scaffolding descriptors, task/framing variables, turn position and model/source fixed effects. Scaffolding descriptors improved fit beyond user state and context (full scaffolding block LR $\chi^2_7 = 1337.5$, $p < .001$; Supplementary Table C13), and adding prior-state $\times$ support-form interactions improved fit further (pooled model/source fixed effects: LR $\chi^2_{18} = 300.2$, $p < .001$; Supplementary Table C14). This indicates that support form added local information beyond the user’s immediately preceding engagement state, and that its role depended on that state. Focusing on feedback-like support and explaining—the two positive forms identified in the conversation-level analysis—Fig. 5d shows that explaining was the more stable local marker: it predicted next-turn constructive engagement after prior constructive, active and passive turns (ORs 1.72, 2.12 and 2.88, respectively; all $p < .001$). Feedback was more state-contingent, with lower next-turn odds after prior constructive and active turns but higher odds after prior passive turns. Thus, the integrated models sharpen the conversation-level finding: feedback

was a strong conversation-level marker, whereas explaining was the more consistent adjacent-turn support form.

Coarser within-conversation summaries bounded the temporal claim. Before–after contrasts around the first scaffolded assistant turn did not show a consistent conversation-wide increase: coding settings were small positive, whereas writing settings were near-zero or negative (Table 2; Supplementary Fig. D1). We therefore treat these summaries as exploratory boundary checks. The strongest temporal evidence remained adjacent-turn coupling, linking prior user state, assistant support form and the immediately following user turn.

Taken together, the temporal analyses characterize scaffolding as a local interactional process rather than a uniform exposure. Constructive engagement was most clearly coupled to support when the user’s current state and the assistant’s support form aligned: engaged user states more often preceded scaffolded support, explanatory support was most consistently followed by further constructive work and feedback shifted with the state it followed. Learning-oriented engagement in everyday human–LLM interaction is therefore assembled turn by turn, through sequences that preserve reasons, uncertainty and room for the user to continue working with ideas.

3 Discussion

Everyday human–LLM interaction as an informal learning ecology

This study identifies everyday human–LLM interaction as a middle space between ordinary tool use and designed instruction. Across naturalistic WildChat, LMSYS Chat and ShareChat coding and writing conversations, learning-oriented engagement was recurrent, but its constructive form was selective. Engagement labels are useful here because they make learning-oriented participation observable in logs. Within that broader behavioural trace, constructive engagement provides the deepest user-side signal: moments in which users make their reasoning visible by testing a boundary, revising an idea, articulating a constraint or extending an explanation. Our analysis shows that this deeper participation was not simply a property of a user, a task or an assistant response. It appeared more often when explicit learning-oriented user framing, task affordances, sustained exchange and assistant support preserved reasons and opportunities for users to keep working with ideas. Feedback-like support and explanation aligned more clearly with constructive engagement than more directive or question-based moves, and the temporal evidence was strongest in adjacent-turn coupling. Together, these findings suggest that learning-oriented participation in everyday LLM use emerges through situated exchanges in which users and assistants maintain a shared space for testing, revising and extending understanding while pursuing practical goals.

This interactional view also clarifies how everyday informal learning differs from designed LLM tutoring. Recent studies show that AI-enabled tutoring can improve learning when pedagogical structure is deliberately engineered into the system, including carefully designed scaffolding, feedback and real-time guidance [41–43]. Yet

learning-science and intelligent-tutoring research also shows that help is not automatically productive: students vary in how they seek, use or bypass hints and feedback, and too much or too direct assistance creates the assistance dilemma [20, 44–46]. Designed tutoring approaches remain valuable when learning is explicit, but everyday human–LLM interaction is usually organized around task progress. The relevant question is therefore not whether an exchange looks pedagogical, but whether it keeps the user involved in the work of understanding while the task is being solved. In our results, constructive participation appeared beyond explicitly learning-oriented conversations and became more visible when exchanges gave users room to surface uncertainty, test alternatives, revise assumptions and extend ideas. The broader implication is that informal learning in everyday AI use depends on how interactional conditions distribute cognitive responsibility: whether assistant support resolves the task for the user, or preserves enough reasons, uncertainty and agency for the user to keep working with ideas.

Scaffolding as situated support

Assistant support becomes part of this informal learning ecology by shaping what remains open for the user after an answer. The conversation-level results establish scaffolded support as a marker of exchanges in which constructive participation is more visible, but this binary distinction is only a starting point. In task-oriented LLM use, assistance can function as an endpoint that supplies the next step, or as a working object that leaves a reason, error, mechanism, constraint or partial solution available for inspection. This distinction is central to scaffolding theory, which treats support not simply as help-giving, but as contingent assistance that structures activity while gradually returning responsibility to the learner [20, 27, 28, 30, 40]. In everyday human–LLM interaction, the relevant question is therefore not whether a response looks pedagogical, but whether it keeps some part of the problem available for the user to examine, adapt or test.

The support-form results give empirical content to this account. Feedback-like support and explanation were most closely aligned with constructive engagement because they tend to turn assistance into material for further user work. Feedback makes a prior attempt evaluable: it can expose a mismatch, partial success or error that the user can repair or extend [34, 35]. Explanation makes reasons and mechanisms inspectable, giving users material for comparison, self-explanation and transfer [22, 37]. More directive instruction, modelling and questioning can also be useful, especially when users need a next step, an example or clarification. In these conversations, however, they were less consistently aligned with visible constructive elaboration than feedback and explanation. The broader insight is that informal-learning support depends on the kind of cognitive space an answer leaves behind. Support is most learning-relevant when it helps users make progress while preserving a productive gap between receiving an answer and understanding it—a space in which they can compare, diagnose, justify, revise or transfer. This reframes scaffolding for general-purpose AI: not as adding more pedagogical moves to a tool, but as designing task assistance that keeps users’ epistemic agency active while the problem is being solved.

From local coupling to learning continuity

The temporal results clarify the scale at which everyday AI use can contribute to informal learning. The strongest evidence was local: scaffolded support and constructive engagement were coupled in adjacent user–assistant–user sequences, whereas coarser within-conversation summaries did not support a uniform conversation-wide trajectory. This places everyday human–LLM interaction at a different scale from formal instruction. Formal learning environments organize accumulation through curricula, sequencing, shared standards, assessment and repeated opportunities to abstract beyond the immediate task. Everyday AI use is more episodic and task-embedded. Its contribution lies in point-of-need moments when users encounter uncertainty, compare alternatives, repair failed attempts or adapt ideas to concrete problems.

Seen this way, local coupling is not merely a short-run association. It identifies a unit of learning continuity: a moment in which a problem state becomes visible, support is provided and the next user turn either carries reasoning forward or lets it close. These moments are not durable learning outcomes by themselves, but they are the kinds of episodes that future systems could help users connect across time. A learning-oriented assistant would not need to turn every task into a formal lesson. It could preserve the rationale, constraint or uncertainty that a local exchange surfaced; help users recognize recurring breakdowns; and support consolidation after a task has been solved.

This complementarity reframes the role of personal AI assistants in education. Formal instruction provides progression, accountability and shared criteria. Everyday AI assistance can carry learning into the intervals between formal learning episodes, where people solve real problems and test ideas under practical constraints. The human-centred stake is whether AI-mediated work makes reasoning disappear into fluent answers, or whether it leaves recoverable traces of why a solution worked, what uncertainty remained and how a related problem might be approached next. In that sense, the promise of lifelong AI assistance is not to replace formal education, but to give continuity, memory and personalization to the informal learning that happens while people work.

Towards a science of everyday AI learning

The present study establishes a large-scale empirical baseline for understanding informal learning in everyday human–LLM interaction. A necessary first step in evaluating whether AI support improves retention, transfer or learner autonomy is to identify where learning-oriented participation already appears in ordinary use and how it is organized by task ecology, user framing, support form and interactional timing. Naturalistic conversational logs are particularly valuable for this first step. They reveal behavioural regularities across hundreds of thousands of interactions, making visible patterns that are difficult to observe within the narrower scope of classroom studies, laboratory experiments or purpose-built tutoring systems. Across three public conversational corpora, the framework introduced here identifies a set of measurable behavioural signatures through which learning-relevant activity becomes observable in everyday problem solving.

More broadly, the findings suggest that informal learning in the age of AI is becoming increasingly intertwined with everyday work. As AI systems become embedded in writing, coding, searching and problem solving, opportunities for learning may arise not only in classrooms or explicitly educational settings, but also through repeated episodes of explanation, evaluation, revision and testing that occur while people pursue practical goals. Understanding these processes requires moving beyond questions of answer quality or task efficiency alone. It requires understanding how AI systems shape the opportunities people have to participate in the work of understanding itself. In this sense, the contribution of the present study is not only methodological but conceptual: it makes everyday human–LLM interaction observable as a site of informal learning and opens a broader question about the role of AI in human development.

Limitations

Several boundaries define the scope of these claims. The analyses are observational and should be interpreted as structured associations in naturalistic digital traces, not as causal estimates of learning effects [47, 48]. Scaffolded support, task context, user framing and model/source conditions were observed rather than assigned, and unobserved task difficulty or user motivation may partly explain where constructive engagement appears. The engagement labels also capture behavioural signatures of learning-oriented participation rather than verified learning outcomes; they do not measure retention, transfer, prior knowledge or durable skill change. Finally, the analytic corpora emphasize public English-language coding and writing conversations and rely on LLM-assisted annotation with human validation, so future work should test the framework in other domains, languages, platforms and longitudinal learner-level settings.

Conclusion

As AI becomes part of the infrastructure of everyday life, the central question shifts from how individuals learn from intelligent systems to how learning itself is organized in AI-mediated societies. The issue is not only what AI can do for people, but whether AI-mediated environments continue to create opportunities for the forms of participation through which judgment, understanding and intellectual growth emerge. The long-term significance of AI will therefore depend as much on whether it sustains opportunities for people to construct, question and extend understanding through active engagement in the world as on the knowledge it delivers.

4 Methods

4.1 Study design

This study is an observational analysis of naturalistic public human–LLM conversations from WildChat, LMSYS Chat-1M (reported below as LMSYS Chat) and the ShareChat strict-English subset (reported below as ShareChat) across two everyday task ecologies, coding and writing. It asks whether learning-oriented engagement is visible in ordinary use, where constructive engagement is concentrated and how assistant

scaffolding is ordered with user engagement across conversation-level and adjacent-turn timescales. The design is descriptive and associational, following the logic of large-scale digital trace and computational social-science analyses [47, 49]. It does not estimate randomized exposure effects, causal instructional efficacy or durable learning outcomes [48].

The empirical sequence mirrors the Results section. We first estimate the prevalence and composition of user engagement, then examine how constructive engagement varies by user framing, task ecology and interaction depth. We next test whether scaffolded assistant support marks conversations with richer constructive participation, separate scaffolded support by support form and analyse local temporal coupling between assistant support and user engagement. Dataset and task ecology are treated as substantive dimensions of the interaction environment rather than nuisance variables to be averaged away.

4.2 Data sources and analytic samples

The analytic corpus contains 128,569 conversations and 981,470 total turns, including 491,685 user turns and 489,785 assistant turns. The six task settings are WildChat coding, WildChat writing, LMSYS Chat coding, LMSYS Chat writing, ShareChat coding and ShareChat writing. WildChat and LMSYS Chat provide large public human-LLM interaction logs from deployed chat systems, whereas ShareChat provides a smaller publicly shared conversation corpus after English-language filtering; together they offer complementary naturalistic settings rather than experimentally matched platform conditions [23–25]. WildChat coding contributes 31,878 conversations with 124,073 user turns and 124,623 assistant turns. WildChat writing contributes 39,534 conversations with 163,705 user turns and 164,824 assistant turns. LMSYS Chat coding contributes 32,114 conversations with 107,065 user turns and 105,174 assistant turns. LMSYS Chat writing contributes 21,023 conversations with 77,906 user turns and 76,279 assistant turns. ShareChat contributes 2,481 coding conversations with 11,541 user turns and 11,522 assistant turns and 1,539 writing conversations with 7,395 user turns and 7,363 assistant turns.

We also processed two newer corpora with the same pipeline as boundary checks. SWE-chat provides a recent agentic software-engineering setting and retained 1,536 coding conversations after filtering, so we report it as a coding-specific external check rather than as part of the balanced six-setting analysis [50]. ThoughtTrace retained only 56 coding and 61 writing conversations after the final task-and-language filter, with near-saturated scaffolded support, and is therefore kept as a feasibility check in the Appendix [51].

The six task settings were analysed with a common Level 1–4 annotation and analysis framework so that prevalence, contextual variation, support–engagement association and temporal ordering were computed from the same construct definitions. The manuscript reports the refreshed production outputs after the latest user-turn update pass and the subsequent LMSYS Chat and ShareChat annotation runs rather than earlier pre-update batches. Supplementary Section A.1 describes the corpus construction and refresh provenance in more detail.

Corpus filtering and classification are reported in Supplementary Table A1. In brief, coding and writing task ecologies were assigned by English-oriented LLM-assisted semantic filters rather than by native topic metadata alone. ShareChat additionally underwent a strict-English screen after task filtering, retaining only conversations whose final conversation-level language flag was English. Supplementary Table A2 summarizes the available validation and audit evidence for the task, language and user-framing classifiers.

4.3 Conversation metadata and task ecologies

Each conversation is represented as an ordered sequence of user and assistant turns together with available metadata. The main contextual variables are data source, task domain, user framing, topic and language. Task domain and user framing are analytic variables generated by the filtering and annotation pipeline; topic and language are retained as corpus or pipeline metadata where available. Coding and writing were selected because they are common everyday LLM use cases but differ in their interactional affordances: coding often involves debugging, error diagnosis and iterative testing, whereas writing often involves drafting, revision and directive editing [31–33]. User framing is a conversation-level LLM-assisted label that distinguishes conversations explicitly oriented toward understanding, exploration or learning from conversations without such explicit learning-oriented framing.

Behavioural depth is measured from the number and ordering of turns in the conversation. Depth is used descriptively, as a marker of interactional opportunity and sustained exchange, not as an independent manipulation. Post-answer depth is measured as the number of turns after the first assistant response.

4.4 Turn-level constructs

The unit of annotation is the conversational turn. User and assistant turns are labelled with different constructs because they play different roles in the interaction. User turns are labelled for behavioural, emotional and cognitive engagement, drawing on engagement theory and the ICAP distinction between passive, active and constructive participation [21, 22]. The main analyses emphasize cognitive engagement and its depth levels. Passive engagement (P) captures reception or acknowledgement without substantial transformation. Active engagement (A) captures task-relevant action, follow-up or manipulation of supplied information. Constructive engagement (C) captures elaboration, revision, testing, extension or transformation of ideas. Constructive engagement is treated as the highest-level behavioural proxy for learning-oriented participation available in these logs, not as a direct measure of knowledge gain.

Assistant turns are first distinguished as S1 or S2. S1 denotes a non-scaffolded reference response, including straightforward task fulfilment or stand-alone information provision. S2 denotes scaffolded support: assistant behaviour that structures, regulates or redirects the user’s process while leaving some cognitive work with the user. This definition follows the scaffolding tradition, especially contingency, transfer of responsibility and potential fading [27, 28, 30]. Scaffolded turns can also receive intention labels (I1 metacognitive, I2 cognitive and I3 affective support) and non-mutually exclusive

means labels (M1 feedback, M2 hinting, M3 instructing, M4 explaining, M5 modelling and M6 questioning). Operational examples are provided in Supplementary Table C1.

4.5 Annotation workflow and verification

The annotation workflow combined theory-grounded codebook development, prompt iteration, parser checks, human review and production-scale LLM-assisted labelling. We used LLM-assisted annotation as a scalable text-analysis workflow while retaining parser checks and human verification because prior work shows both the promise and the need for validation when language models are used as annotators [52, 53]. Prompt and post-processing rules were first developed on validation batches to reduce common failure modes, including over-labelling brief acknowledgements as active engagement and misclassifying directive assistant responses as scaffolding. After the pipeline stabilized, the same annotation logic was applied to the WildChat, LMSYS Chat and ShareChat analytic corpora. All examples reproduced in the manuscript are de-identified and paraphrased rather than copied from raw logs.

Human verification was used as a quality-control input rather than as a separate source of outcome measurement. The final coding validation rerun combines three review rounds, with 658 task turns and 643 turns after deduplication. The final writing validation used 210 reviewed conversations. These validation sets were organized by task ecology rather than independently powered within each public corpus. WildChat also underwent a final targeted user-turn update check; LMSYS Chat and ShareChat were then processed with the same codebook, parser checks and post-processing logic, but the current artifacts do not include separate human-agreement estimates for each added corpus. Agreement was summarized with per-label F1, Matthews correlation coefficient and Gwet’s AC1 because the labels are sparse, prevalence-imbalanced and not mutually exclusive in all cases [54, 55]. Agreement was strongest for constructive, passive, emotional, scaffolded-support and support-form labels. Boundary-heavy labels such as writing A and S1 were less stable; these labels support descriptive composition and the reference category, but the primary claims rely on constructive engagement, scaffolded support and support-form labels rather than on writing active-uptake or S1 boundary distinctions. Supplementary Table C2 reports the per-label agreement metrics, and Supplementary Section A.4 describes the final user-turn update check and corpus-transfer scope.

4.6 Outcome construction

The primary user-side outcome is constructive engagement. At the turn level, this is a binary indicator for whether the user turn is labelled C. At the conversation level, we use both the count of constructive user turns and an indicator for whether a conversation contains at least one constructive user turn. Constructive ratios are defined as constructive user turns divided by user turns.

The primary assistant-side exposure is scaffolded support. At the turn level, this is the distinction between scaffolded support (S2) and non-scaffolded reference response (S1). At the conversation level, we use an indicator for whether a conversation contains at least one scaffolded assistant turn. Support-form analyses condition

on scaffolded turns and compare conversations containing each means label with scaffolded conversations that do not contain that label. Because means labels can co-occur, these are descriptive support-form contrasts rather than mutually exclusive treatment comparisons.

4.7 Statistical analyses

Descriptive analyses estimate engagement prevalence, scaffolding prevalence, passive/active/constructive composition, emotional-engagement rates and behavioural-depth summaries within each task setting. Context analyses compare intentional versus unintentional conversations and coding versus writing settings. Depth analyses group conversations into behavioural-depth buckets and estimate the probability that a conversation contains at least one constructive turn. Because any-event outcomes mechanically become easier to observe in longer conversations, we also fit a grouped-binomial constructive-rate sensitivity model in which the outcome is the constructive user-turn count out of total user turns, with user framing, task ecology, length bucket and dataset fixed effects as predictors.

Conversation-level associations are analysed by comparing conversations with and without at least one scaffolded assistant turn. Constructive-ratio contrasts use turn-weighted ratios within each comparison group, defined as total constructive user turns divided by total user turns, with uncertainty from 1,000 conversation-level bootstrap resamples. The adjusted Section 2.3 count model is a setting-specific Poisson generalized linear model for the raw constructive-turn count per conversation, not for a rate. For each setting s , the generic mean model was $\mathbb{E}(C_i) = \exp(\alpha_s + \beta_s S_i + \gamma_s^\top X_i)$, where C_i is the constructive-turn count, S_i indicates at least one scaffolded assistant turn and X_i contains observed conversation-level context such as user framing, conversation length, emotional-expression ratio and error or persistence markers where available. We did not use a user-turn offset in the primary count model; conversation length entered directly as a covariate. The exponentiated coefficient for scaffolded-support presence is reported as a multiplicative association with expected constructive-turn count conditional on observed covariates. As a rate sensitivity, we refitted the same setting-specific mean model with user turns represented as an exposure offset and quasi-Poisson scaled standard errors.

The companion binary model for Section 2.3 is a setting-specific logistic regression for whether a conversation contains at least one constructive user turn, using scaffolded-support presence and the same core conversation-level context available within settings. Task ecology is fixed within each setting and is therefore used in pooled/context models, not interpreted as a within-setting coefficient. Supplementary Table C5 lists the estimable covariate sets used for the setting-specific count, binary and offset-rate models. Model-based standard errors and 95% confidence intervals are shown for the adjusted Poisson and logistic estimates. We checked Poisson overdispersion using the Pearson χ^2 statistic divided by residual degrees of freedom; the dispersion estimates ranged from 1.16 to 2.06 across the six settings. A quasi-Poisson sensitivity that scaled the Poisson standard errors by this dispersion preserved the direction and statistical distinguishability of the scaffolded-support association

in every setting, and the offset-rate sensitivity likewise preserved positive scaffolded-support associations in every setting. Negative-binomial models were not used as the primary specification because the estimand was the common pre-specified multiplicative association under the same mean model across settings; the overdispersion and offset checks were used to assess variance and exposure-specification sensitivity rather than to change the main estimand. Binary outcomes in other analyses, including whether the next user turn is constructive and whether the next assistant turn is scaffolded, are modelled with logistic regression or reported as conditional probability contrasts, depending on the estimand [56, 57]. Percentage-point contrasts are reported as effect sizes; adjacent-turn contrasts use 1,000 conversation-cluster bootstrap resamples because multiple turn pairs can come from the same conversation. Support-form bracket tests in Fig. 4 use Benjamini–Hochberg false-discovery-rate adjustment within each displayed family of six support-form comparisons. Supplementary Tables C3–C6 summarize the main model families and temporal estimands.

To assess whether the adjacent-turn pattern was reducible to a single feature family, we also fitted integrated logistic regressions that combined user-state features, assistant-scaffolding features and adjacent-turn context in one model. The outcome was whether the next user turn was constructive. The pooled model included dataset fixed effects; a sensitivity replaced these with recoverable assistant model or source-family fixed effects, using exact `chat_model` labels where available, LMSYS Chat model names recovered from conversation identifiers and ShareChat assistant/source families recovered from conversation identifiers. Standard errors for these adjacent-turn regressions were clustered by conversation. We report broad scaffolded-support models separately from support-form decompositions because M1–M6 are nested descriptors of scaffolded support rather than independent exposures. Nested likelihood-ratio block tests first add broad scaffolded-support presence and then test whether support-form descriptors add signal within scaffolded support. These model/source checks are interpreted as robustness analyses rather than model-causal comparisons.

Temporal analyses focus on adjacent-turn ordering. The primary analyses compare the probability that the next user turn is constructive after scaffolded versus non-scaffolded assistant turns, estimate the same contrast within prior user states and estimate the probability that the next assistant turn is scaffolded after different prior user states. We also fit adjacent-turn regressions with prior-state $\times$ support-form interactions to test whether local coupling depends jointly on the user’s immediately preceding engagement state and the form of scaffolding provided. Coarser within-conversation before–after summaries are reported only as exploratory boundary checks, not as evidence of cumulative learning trajectories, because the first scaffolded turn may be selected by task difficulty, user motivation or an already developing exchange.

4.8 Robustness, privacy and LLM use

WildChat robustness analyses repeat the main patterns across user-facing model snapshots and coarse model families. These checks assess whether the headline patterns

are reducible to a single WildChat model label. They are not interpreted as model-causal comparisons, because model labels are not experimentally assigned and may be confounded with time, user population, task mix and logging context.

The analyses are bounded by the observational design and by the nature of the labels. They support claims that everyday human–LLM conversations contain measurable learning-oriented engagement; that constructive engagement varies by intent, task ecology and depth; that assistant scaffolding is associated with richer constructive participation; and that support and engagement show local reciprocal ordering. They do not support claims that scaffolded support causes durable learning gains, that all constructive engagement reflects actual knowledge change or that the labelled conversational states perfectly measure learning itself.

LLMs were used as objects of study and as annotation tools. Production labels were generated with LLM-assisted workflows and then checked through parser rules, targeted update scripts and human validation. The authors also used LLM-based tools to assist with code drafting, debugging, analysis workflow organization and language editing. LLM-generated assistance was not used as a substitute for human verification of the validation samples or for interpretation of the statistical results.

All reported results are aggregate turn-level or conversation-level analyses. The study does not link conversations into user histories, estimate user-level longitudinal trajectories or publish raw conversation text; temporal analyses are confined to ordering within a conversation.

Data availability

This study analyses public human–LLM conversation corpora released by their original data providers: WildChat, LMSYS Chat-1M, ShareChat, SWE-chat and Thought-Trace. Readers should obtain the raw public conversation data from the original releases and comply with the corresponding dataset licences, terms of use and redistribution restrictions. The manuscript does not redistribute raw message text, user identifiers or linked user histories.

Derived data sufficient to verify the reported aggregate figures, tables and statistical summaries are provided with the analysis repository as figure source-data files in `source_data/`, table files in `tables/` and model/statistical output files in `outputs/integrated_regression/`. These files include the numeric values underlying the main figures, table summaries, confidence intervals, p values, false-discovery-rate q values and model-output summaries used in the manuscript. Frozen archival data record: [Zenodo/OSF DOI or accession to be added]. Any request for additional derived label files that are not included in the public deposit will be assessed against the original dataset licences and privacy constraints.

Code availability

Custom code used to filter corpora, run the LLM-assisted annotation workflow, post-process labels, compute statistical analyses and generate figures and tables is maintained in the project repository at https://github.com/CinderD/Jun_10_NHB_Informal_Learning. Frozen archival code release: [Zenodo/OSF DOI or accession to be

added]. The repository excludes API keys, raw message text and any files that cannot be redistributed under the original corpus terms. Reproducing the full annotation pipeline requires user-provided API credentials and access to the raw public corpora from their original sources.

Ethics statement

The study is a secondary analysis of public, naturalistic human–LLM conversation datasets. The authors did not recruit participants, intervene in conversations or attempt to identify users. Analyses are reported at aggregate turn or conversation level, and examples in the manuscript are paraphrased rather than copied from raw logs. Human validation was used to assess annotation quality on reviewed examples and did not involve interaction with the original dataset users. Institutional ethics review/IRB determination has been obtained for this secondary analysis; institutional name and approval or exemption identifier: [to be added].

Competing interests

The authors declare no competing interests.

Author contributions

Author contributions are organized using the CRediT taxonomy. The contribution categories for this study are conceptualization, study design, data curation, annotation pipeline development, validation, statistical analysis, figure generation, writing of the original draft, manuscript review and editing, supervision and funding acquisition. Author initials and final role allocation: [to be added after the author list and order are confirmed]. All listed authors will approve the submitted version and will be accountable for their contributions.

Acknowledgements and funding

Acknowledgements and funding information, including grant numbers where applicable: [to be added]. Funders had no role in study design, data collection, analysis, interpretation, manuscript preparation or publication decisions unless stated otherwise.

LLM use statement

Large language models were objects of study and were also used as annotation and writing-assistance tools. Production labels were generated with LLM-assisted workflows and checked through parser rules, targeted update scripts and human validation. The authors also used LLM-based tools to assist with code drafting, debugging, workflow organization and language editing. No LLM or AI tool is listed as an author, and the authors take responsibility for the final manuscript, analyses and interpretation. Additional details are provided in Methods.

Reporting summary and editorial policy checklist

A completed Nature Portfolio Reporting Summary and any journal-required editorial policy checklist will be provided at submission. The manuscript and supplementary information report the observational design, sampling and filtering rules, annotation workflow, human validation evidence, statistical model specifications, bootstrap procedures, multiple-comparison handling and data/code availability statements needed to assess reproducibility.

Figure source data

Source data for the numeric main figures are provided in `source_data/figure2_source_data.csv`, `source_data/figure3_source_data.csv`, `source_data/figure4_source_data.csv` and `source_data/figure5_source_data.csv`. Figure 1 is a conceptual framework figure and has no numeric source data. Supplementary table and regression source files are provided in `tables/` and `outputs/integrated_regression/`.

References

- [1] Chatterji, A. *et al.* How people use ChatGPT. Tech. Rep. Working Paper 34255, National Bureau of Economic Research (2025).
- [2] Noy, S. & Zhang, W. Experimental evidence on the productivity effects of generative artificial intelligence. *Science* **381**, 187–192 (2023).
- [3] Brynjolfsson, E., Li, D. & Raymond, L. Generative AI at work. *The Quarterly Journal of Economics* **140**, 889–942 (2025).
- [4] Barke, S., James, M. B. & Polikarpova, N. Grounded copilot: How programmers interact with code-generating models. *Proceedings of the ACM on Programming Languages* **7**, 85–111 (2023).
- [5] Kobiella, C., Flores López, Y. S., Waltenberger, F., Draxler, F. & Schmidt, A. “If the machine is as good as me, then what use am I?”—how the use of ChatGPT changes young professionals’ perception of productivity and accomplishment. *Proceedings of the CHI Conference on Human Factors in Computing Systems* 1–16 (2024).
- [6] Eraut, M. Informal learning in the workplace. *Studies in Continuing Education* **26**, 247–273 (2004).
- [7] Marsick, V. J. & Watkins, K. *Informal and incidental learning in the workplace (Routledge Revivals)* (Routledge, 2015).
- [8] Schugurensky, D. The forms of informal learning: Towards a conceptualization of the field. Tech. Rep. NALL Working Paper 19, Centre for the Study of Education

and Work, OISE/UT, Toronto (2000). URL <https://hdl.handle.net/1807/2733>.

- [9] Kasneci, E. *et al.* ChatGPT for good? On opportunities and challenges of large language models for education. *Learning and Individual Differences* **103**, 102274 (2023).
- [10] Meyer, J. *et al.* Using LLMs to bring evidence-based feedback into the classroom: AI-generated feedback increases secondary students’ text revision, motivation, and positive emotions. *Computers and Education: Artificial Intelligence* **6**, 100199 (2024).
- [11] Arif, T., Asthana, S. & Collins-Thompson, K. Generation and assessment of multiple-choice questions from video transcripts using large language models. *Proceedings of the Eleventh ACM Conference on Learning @ Scale* 530–534 (2024).
- [12] Jin, H., Lee, S., Shin, H. & Kim, J. Teach AI how to code: Using large language models as teachable agents for programming education. *Proceedings of the CHI Conference on Human Factors in Computing Systems* 1–28 (2024).
- [13] Lim, H., Choi, D. & Hong, H. Identify design problems through questioning: Exploring role-playing interactions with large language models to foster design questioning skills. *Companion Publication of the 2024 Conference on Computer-Supported Cooperative Work and Social Computing* 598–602 (2024).
- [14] Terzimehić, N., Bühler, B. & Kasneci, E. Conversational AI as a catalyst for informal learning: An empirical large-scale study on LLM use in everyday learning. *Computers and Education: Artificial Intelligence* 100634 (2026).
- [15] VanLehn, K. The relative effectiveness of human tutoring, intelligent tutoring systems, and other tutoring systems. *Educational Psychologist* **46**, 197–221 (2011).
- [16] Bastani, H. *et al.* Generative AI without guardrails can harm learning: Evidence from high school mathematics. *Proceedings of the National Academy of Sciences* **122**, e2422633122 (2025).
- [17] Lee, H.-P. *et al.* The impact of generative AI on critical thinking: Self-reported reductions in cognitive effort and confidence effects from a survey of knowledge workers. *Proceedings of the CHI Conference on Human Factors in Computing Systems* 1–22 (2025).
- [18] Fan, Y. *et al.* Beware of metacognitive laziness: Effects of generative artificial intelligence on learning motivation, processes, and performance. *British Journal of Educational Technology* **56**, 489–530 (2025).

- [19] Salomon, G., Perkins, D. N. & Globerson, T. Partners in cognition: Extending human intelligence with intelligent technologies. *Educational Researcher* **20**, 2–9 (1991).
- [20] Koedinger, K. R. & Alevan, V. Exploring the assistance dilemma in experiments with cognitive tutors. *Educational Psychology Review* **19**, 239–264 (2007).
- [21] Fredricks, J. A., Blumenfeld, P. C. & Paris, A. H. School engagement: Potential of the concept, state of the evidence. *Review of Educational Research* **74**, 59–109 (2004).
- [22] Chi, M. T. & Wylie, R. The ICAP framework: Linking cognitive engagement to active learning outcomes. *Educational Psychologist* **49**, 219–243 (2014).
- [23] Zhao, W. *et al.* WildChat: 1M ChatGPT interaction logs in the wild (2024). [arXiv:2405.01470](#).
- [24] Zheng, L. *et al.* LMSYS-Chat-1M: A large-scale real-world LLM conversation dataset (2023). [arXiv:2309.11998](#).
- [25] Yan, Y., Nguyen, T., Su, B., Lieffers, M. & Le, T. ShareChat: A dataset of chatbot conversations in the wild (2025). [arXiv:2512.17843](#).
- [26] Handa, K. *et al.* Which economic tasks are performed with AI? Evidence from millions of Claude conversations (2025). [arXiv:2503.04761](#).
- [27] Wood, D., Bruner, J. S. & Ross, G. The role of tutoring in problem solving. *Journal of Child Psychology and Psychiatry* **17**, 89–100 (1976).
- [28] Van de Pol, J., Volman, M. & Beishuizen, J. Scaffolding in teacher–student interaction: A decade of research. *Educational Psychology Review* **22**, 271–296 (2010).
- [29] Winkler, R., Hobert, S., Salovaara, A., Söllner, M. & Leimeister, J. M. Sara, the lecturer: Improving learning in online education with a scaffolding-based conversational agent. *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems* 1–14 (2020).
- [30] Quintana, C. *et al.* A scaffolding design framework for software to support science inquiry. *The Journal of the Learning Sciences* **13**, 337–386 (2004).
- [31] Keuning, H., Jeuring, J. & Heeren, B. A systematic literature review of automated feedback generation for programming exercises. *ACM Transactions on Computing Education* **19**, 1–43 (2018).
- [32] Flower, L. & Hayes, J. R. A cognitive process theory of writing. *College Composition and Communication* **32**, 365–387 (1981).

- [33] Sommers, N. Revision strategies of student writers and experienced adult writers. *College Composition and Communication* **31**, 378–388 (1980).
- [34] Hattie, J. & Timperley, H. The power of feedback. *Review of Educational Research* **77**, 81–112 (2007).
- [35] Shute, V. J. Focus on formative feedback. *Review of Educational Research* **78**, 153–189 (2008).
- [36] Nicol, D. J. & Macfarlane-Dick, D. Formative assessment and self-regulated learning: A model and seven principles of good feedback practice. *Studies in Higher Education* **31**, 199–218 (2006).
- [37] Chi, M. T. H., Bassok, M., Lewis, M. W., Reimann, P. & Glaser, R. Self-explanations: How students study and use examples in learning to solve problems. *Cognitive Science* **13**, 145–182 (1989).
- [38] Renkl, A. Learning from worked-out examples: A study on individual differences. *Cognitive Science* **21**, 1–29 (1997).
- [39] Aleven, V., Stahl, E., Schworm, S., Fischer, F. & Wallace, R. Help seeking and help design in interactive learning environments. *Review of Educational Research* **73**, 277–320 (2003).
- [40] Reiser, B. J. Scaffolding complex learning: The mechanisms of structuring and problematizing student work. *The Journal of the Learning Sciences* **13**, 273–304 (2004).
- [41] Kestin, G., Miller, K., Klales, A., Milbourne, T. & Ponti, G. AI tutoring outperforms in-class active learning: an RCT introducing a novel research-based design in an authentic educational setting. *Scientific Reports* **15**, 17458 (2025).
- [42] Wang, R. E., Ribeiro, A. T., Robinson, C. D., Loeb, S. & Demszky, D. Tutor CoPilot: A human-AI approach for scaling real-time expertise (2024). [arXiv:2410.03017](https://arxiv.org/abs/2410.03017).
- [43] Jin, H. *et al.* RelianceScope: An analytical framework for examining students’ reliance on generative AI chatbots in problem solving. *Proceedings of the Thirteenth ACM Conference on Learning @ Scale* 136–147 (2026).
- [44] Aleven, V., Roll, I., McLaren, B. M. & Koedinger, K. R. Help helps, but only so much: Research on help seeking with intelligent tutoring systems. *International Journal of Artificial Intelligence in Education* **26**, 205–223 (2016).
- [45] Roll, I., Aleven, V., McLaren, B. M. & Koedinger, K. R. Improving students’ help-seeking skills using metacognitive feedback in an intelligent tutoring system. *Learning and Instruction* **21**, 267–280 (2011).

- [46] Baker, R. S. J. d. *et al.* Why students engage in “gaming the system” behavior in interactive learning environments. *Journal of Interactive Learning Research* **19**, 185–224 (2008). URL <https://www.learntechlib.org/primary/p/24328/>.
- [47] Salganik, M. J. *Bit by Bit: Social Research in the Digital Age* (Princeton University Press, 2017).
- [48] Hernán, M. A. & Robins, J. M. *Causal Inference: What If* (Chapman & Hall/CRC, Boca Raton, 2020).
- [49] Lazer, D. *et al.* Computational social science. *Science* **323**, 721–723 (2009).
- [50] Baumann, J. *et al.* SWE-chat: Coding agent interactions from real users in the wild (2026). [arXiv:2604.20779](https://arxiv.org/abs/2604.20779).
- [51] Jin, C. *et al.* ThoughtTrace: Understanding user thoughts in real-world LLM interactions (2026). [arXiv:2605.20087](https://arxiv.org/abs/2605.20087).
- [52] Gilardi, F., Alizadeh, M. & Kubli, M. ChatGPT outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences* **120**, e2305016120 (2023).
- [53] Ziems, C. *et al.* Can large language models transform computational social science? *Computational Linguistics* **50**, 237–291 (2024).
- [54] Matthews, B. W. Comparison of the predicted and observed secondary structure of T4 phage lysozyme. *Biochimica et Biophysica Acta* **405**, 442–451 (1975).
- [55] Gwet, K. L. Computing inter-rater reliability and its variance in the presence of high agreement. *British Journal of Mathematical and Statistical Psychology* **61**, 29–48 (2008).
- [56] McCullagh, P. & Nelder, J. A. *Generalized Linear Models* 2 edn (Chapman & Hall, London, 1989).
- [57] Agresti, A. *Categorical Data Analysis* 3 edn (Wiley, Hoboken, 2013).

Supplementary information

Supplementary information is provided below.

Appendix A Supporting methods

A.1 Corpus construction and analytic settings

The analyses use six task settings formed by crossing three public conversational corpora with two task ecologies: WildChat coding, WildChat writing, LMSYS coding, LMSYS writing, ShareChat coding and ShareChat writing. Coding and writing were selected because both are common everyday LLM use cases, yet they differ in the opportunities they create for visible learning-oriented engagement. Coding conversations often contain error messages, debugging attempts, execution failures and iterative repair. Writing conversations more often involve drafting, rewriting, shortening, tone adjustment and direct editing.

The analytic corpus contains 128,569 conversations and 981,470 total turns. WildChat contributes 31,878 coding conversations and 39,534 writing conversations; LMSYS contributes 32,114 coding conversations and 21,023 writing conversations; ShareChat contributes 2,481 coding conversations and 1,539 writing conversations. All reported analyses use the refreshed Level 1–4 outputs after the latest user-turn update pass and the subsequent LMSYS and ShareChat annotation runs.

Supplementary Table A1 summarizes how raw corpus records were converted into these analytic samples. It distinguishes source or pipeline metadata from analysis variables: coding/writing task ecology comes from LLM-assisted semantic filters, user framing comes from the conversation-level `learning_intent` annotation, and ShareChat strict-English status comes from a final conversation-level language flag applied after task filtering.

SWE-chat and ThoughtTrace were also processed with the same pipeline as feasibility and boundary checks (Supplementary Table C17) [50, 51]. SWE-chat retained a usable coding sample (1,536 conversations) but only three writing conversations after filtering, so it can serve only as a coding-specific stress test rather than as part of the main coding–writing comparison. ThoughtTrace retained 56 coding and 61 writing conversations after the final task-and-language filter, but the sample remains small and scaffolded support is near-saturated, especially in coding where all conversations contain at least one scaffolded assistant turn. These properties make ThoughtTrace useful as a feasibility check but not as stable evidence for conversation-level support-association estimation.

For each conversation, the pipeline retained the ordered turn sequence, role labels, available topic and language metadata, learning-intent indicators and model metadata where available. Model metadata are used only for robustness checks where the corpus exposes usable model labels, because model assignment in naturalistic public conversations is observational and may reflect time, user population, task mix and logging context.

Table A1 Corpus filtering and classification pipeline. The table records the auditable steps used to construct the six main analytic settings. Task domain and user framing are analysis variables generated by the filtering and annotation pipeline, not native corpus metadata.

Corpus	Source records and first screen	Language and minimum-turn filter	Task classification source	Final analytic sample	User framing and audit trail
WildChat	Public WildChat release. The retained manuscript artifact begins from refreshed task-filter outputs rather than the full raw-to-task selection ledger.	Retained task-filter outputs have at least four message turns. The task filters were English-oriented semantic filters; source language metadata were retained for analysis but were not used as an additional strict post-hoc exclusion in the main WildChat sample.	LLM-assisted semantic task filters labelled conversations as coding-english or writing-english with score threshold $\geq .70$ after keyword prefiltering. The retained filter outputs contain 31,879 coding conversations and 39,534 writing conversations; all retained records had task-filter scores above threshold.	31,878 coding; 39,534 writing.	User framing came from the conversation-level learning_intent annotation (INTENTIONAL versus UNINTENTIONAL), not from native metadata. The final WildChat user-turn update pass scanned 31,879 coding and 39,534 writing conversations; one coding conversation with missing learning_intent was omitted from Level 2 and downstream analyses.
LMSYS Chat	LMSYS Chat-1M source release: 1,000,000 conversations.	The conversion retained 777,453 English conversations and excluded 222,547 non-target-language conversations. Among retained English conversations, 236,372 had at least four turns.	The same LLM-assisted semantic task filters were run on the English minimum-turn pool. After completing the full English minimum-turn pool, the coding filter selected 32,114 coding conversations from 777,453 scanned records; the writing filter scanned the same 777,453 records and selected 21,023 writing conversations.	32,114 coding; 21,023 writing.	User framing came from the downstream learning_intent annotation. The conversion log also records redaction, moderation and model/source metadata where available; these fields were retained for robustness checks rather than used as task labels.
ShareChat	ShareChat Hugging Face release. The local Copilot-style conversion produced 129,584 conversations and 1,196,969 message rows after role mapping validation (user=600,222; bot=596,747).	The first screen retained 69,915 conversations with at least four message turns. After task filtering, a strict-English screen retained only conversations with final conversation-level language flag English; conversations flagged non-English or mixed/non-English by this final flag were excluded.	The same LLM-assisted semantic filters selected 3,494 coding and 1,663 writing conversations before the strict-English screen. The strict-English screen retained 2,481 of 3,494 coding conversations and 1,539 of 1,663 writing conversations.	2,481 coding; 1,539 writing.	User framing came from the downstream learning_intent annotation. The strict-English summary recorded zero malformed JSON files; excluded counts were 1,013 coding and 124 writing conversations.

Note. A “conversation” is a role-ordered user–assistant record after corpus-specific conversion. The minimum-turn screen uses message count, not user–assistant round count. The term “strict-English” is reserved for the additional ShareChat final-language screen; WildChat and LMSYS Chat used English-oriented source or semantic filters as described above. Human validation was organized by task ecology rather than independently powered within each corpus; Supplementary Table C2 reports the domain-stratified agreement metrics.

Table A2 Filtering and classifier validation audit. The table separates corpus-construction filters from analysis labels and reports the strongest validation or audit evidence available in the current artifacts. Metrics from the writing-210 validation set are domain-level checks, not corpus-stratified validation for each public corpus.

Variable or filter	How it was assigned	Validation or audit evidence	Interpretation boundary
Task ecology: coding versus writing	LLM-assisted semantic task filters selected conversations matching coding-oriented or writing-oriented use after minimum-turn screening. The filters were English-oriented and used a score threshold of 0.70 after keyword prefiltering where applicable.	The filtering ledger records retained counts and thresholds for each corpus (Supplementary Table A1). In the writing-210 validation artifacts, the conversation-topic field showed high human-LLM agreement ($F1=0.973$, $MCC=0.872$, $Gwet\ AC1=0.936$; support=210), but this is a domain-level check rather than a full cross-corpus task-filter audit.	Task ecology is treated as an analytic classification of conversation content, not native dataset metadata or randomized assignment.
User framing: intentional versus unintentional	Conversation-level LLM-assisted learning.intent annotation marked conversations explicitly oriented toward understanding, exploration or learning versus those without explicit learning-oriented framing.	The writing-210 validation artifacts contain a learning-intent check, but the reviewed writing sample had no positive human cases for this binary label, so the high apparent agreement is not prevalence-informative. The manuscript therefore uses user framing as an observed contextual label and reports robustness with dataset fixed effects rather than claiming corpus-specific classifier accuracy.	User framing is a measured textual signal of explicit learning orientation, not a stable learner trait, motivation measure or causal exposure.
Language filtering	WildChat and LMSYS Chat used English-oriented source or semantic filters. ShareChat received an additional strict-English screen after task filtering, retaining only records with final conversation-level language flag <code>Eng1ish</code> .	The ShareChat strict-English audit recorded zero malformed JSON files and explicit exclusions after task filtering: 1,013 coding and 124 writing conversations were removed by the final language screen. In the writing-210 validation artifacts, language showed near-perfect agreement ($F1=0.998$, $MCC=0.705$, $Gwet\ AC1=0.995$; support=210), with MCC lower because nearly all reviewed records were English.	The term strict-English is used only for ShareChat’s final post-task language screen. Mixed or non-English ShareChat records were excluded from the main analytic sample.
Turn-level labels: constructive engagement and scaffolded support	Production annotation used the final codebook, parser checks and post-processing rules. Constructive engagement and scaffolded support are the primary user-side outcome and assistant-side exposure.	Domain-stratified human validation shows strong agreement for constructive engagement (coding $F1=0.85$; writing $F1=0.89$) and scaffolded support (coding $F1=0.87$; writing $F1=0.97$), with full per-label metrics in Supplementary Table C2.	These labels support the main association and temporal analyses. Validation is task-domain stratified, not independently powered within every public corpus.

A.2 Annotation ontology

The codebook translates learning-science constructs into role-specific behavioural signatures. User turns are labelled for engagement. Passive engagement (P) marks acknowledgement, reception or minimal continuation without visible transformation. Active engagement (A) marks task-relevant action or follow-up, such as asking for a reformatted answer or applying supplied information. Constructive engagement (C) marks deeper work with ideas, such as testing an implication, revising a hypothesis, articulating a constraint, transferring a mechanism or extending an explanation. Emotional engagement is labelled separately because affective expressions were rare and not the main analytic outcome.

Assistant turns are labelled for support. S1 denotes a non-scaffolded reference response: the assistant provides the requested information or deliverable without materially structuring the user’s process. S2 denotes scaffolded support: the assistant structures, regulates, diagnoses or redirects the user’s work while leaving some responsibility with the user. Scaffolded turns can also receive support-intent labels and support-means labels. Intention labels distinguish metacognitive support (I1), cognitive support (I2) and affective support (I3). Means labels distinguish feedback (M1), hinting (M2), instructing (M3), explaining (M4), modelling (M5) and questioning (M6). Means labels are non-mutually exclusive because a single assistant turn can both explain a mechanism and provide an example or hint. Supplementary Tables C15 and C16 report the support-intent and support-means profiles across the six task settings.

A.3 LLM-assisted annotation pipeline

The annotation pipeline combined prompt-based labelling with deterministic parser checks and targeted post-processing. During development, prompts were iterated against validation examples and disagreement cases. Common failure modes included treating brief acknowledgements as active engagement, treating any long assistant answer as scaffolding, collapsing directive instructions into scaffolding without checking whether responsibility remained with the user, and over-labelling sparse support means.

Production labels were generated at turn level. Parser checks enforced the expected label schema, normalized label variants and flagged missing or malformed outputs. Post-processing rules were used only when a recurring error pattern could be specified transparently, such as a narrow class of user-turn false positives. Prompt versions, run metadata, disagreement files and update outputs were retained with the analysis artifacts so that validation cases can be traced without reproducing raw conversation text in the manuscript.

A.4 Human verification and update provenance

Human verification was used to quantify label reliability and identify boundary cases. Coding validation combined three review rounds, with 658 task turns and 643 turns after deduplication. Writing validation used 210 reviewed conversations. Because many labels are sparse or prevalence-imbalanced, agreement is reported with per-label F1,

Matthews correlation coefficient and Gwet’s AC1 rather than with a single global agreement statistic.

The final user-turn check contained 100 coding examples and 100 writing examples reviewed after the update pass. These examples were used to identify residual false positives and calibrate patch-style update filters, especially for Active and Passive user-turn labels. The last WildChat user-turn update pass scanned 31,878 coding conversations and 39,534 writing conversations, applying 1,815 coding updates and 21 writing updates. LMSYS Chat and ShareChat were then processed with the same refreshed annotation workflow, construct definitions and post-processing logic. The main manuscript reports the refreshed labels produced by these final production runs.

The validation evidence is therefore domain-stratified rather than fully corpus-stratified. It supports the reliability of the construct definitions and the final coding/writing labelling logic, but it should not be read as separate human-agreement evidence for WildChat, LMSYS Chat and ShareChat individually. Cross-corpus use of the labels is supported by applying the same schema, parser checks and post-processing rules to all three corpora and by reporting setting-level robustness checks in the Results and Supplementary Tables, rather than by a separate manual audit for each added corpus.

A.5 Outcome construction

At the turn level, the focal user-side outcome is a binary indicator for constructive engagement. At the conversation level, we use both the count of constructive user turns and an indicator for whether at least one constructive user turn occurs. Constructive ratios divide constructive user turns by total user turns; group-level constructive-ratio contrasts are turn-weighted ratios computed from total constructive user turns over total user turns within each group. Cognitive engagement ratios include passive, active and constructive cognitive engagement.

At the assistant side, scaffolded support is the analysis term for turns assigned the S2 code. Conversation-level scaffolding presence indicates whether a conversation contains at least one scaffolded assistant turn. Support-form analyses condition on conversations or turns with scaffolded support and then compare the presence versus absence of each means label. These comparisons are descriptive because means labels can co-occur and because conversations are not randomly assigned to receive a support form.

Behavioural depth is computed from turn counts and turn ordering. Post-answer depth counts turns after the first assistant response. Persistence after failure is analysed only for coding-oriented conversations because debugging failures, error messages and failed-execution contexts are more directly observable in coding than in writing.

A.6 Statistical model specification

The main analyses use simple estimands matched to the claim being made. Descriptive prevalence analyses report proportions within the six task settings. Context analyses compare user framing, task ecology and depth groups. Conversation-level support associations use Poisson generalized linear models for constructive-turn counts and

logistic models for binary constructive-engagement outcomes. Adjacent-turn analyses use assistant-to-user and user-to-assistant turn pairs to estimate local conditional probabilities.

For the Section 2.3 adjusted conversation-level models, the Poisson outcome is the raw count of constructive user turns in a conversation, not a rate. The primary model does not include a user-turn offset. Conversation length is adjusted directly through total turns, so the exponentiated Poisson coefficient should be read as a multiplicative association with expected constructive-turn count conditional on observed conversation length and other covariates. The generic setting-specific mean model is $\mathbb{E}(C_i) = \exp(\alpha_s + \beta_s S_i + \gamma_s^\top X_i)$, where C_i is constructive-turn count, S_i is scaffolded-support presence and X_i contains the observed conversation-level context listed in Supplementary Table C5. The companion logistic model predicts whether a conversation contains any constructive user turn using scaffolded-support presence and the available setting-specific context. The user-framing stratified Poisson models use the same Poisson specification after stratifying by framing and omitting user framing from the covariate set.

Covariates are included when they correspond to observed compositional differences at the analysis level. Depending on the analysis, these include user framing, task indicators in pooled models, turn counts, emotional ratios, error-related variables and turn index. Percentage-point contrasts are reported as effect sizes. Confidence intervals and p values for these contrasts use 1,000 conversation-level bootstrap resamples; adjacent-turn contrasts use 1,000 conversation-cluster bootstrap resamples rather than individual turn-pair resampling to account for repeated pairs within conversations. The adjusted Poisson and logistic estimates in Fig. 3 use model-based standard errors. Poisson overdispersion was checked with the Pearson χ^2 statistic divided by residual degrees of freedom. Dispersion ranged from 1.16 to 2.06 across the six task settings; scaling the Poisson standard errors with a quasi-Poisson dispersion adjustment preserved positive scaffolded-support associations and statistical distinguishability in every setting. To check exposure specification, we also refitted the same setting-specific Poisson mean model with user turns represented as the exposure offset; Supplementary Table C8 reports quasi-Poisson scaled rate ratios. Negative-binomial models were not used as the primary specification; the sensitivity analyses focus on variance misspecification and exposure specification under the same mean model rather than changing the main estimand. Support-form bracket tests use Benjamini–Hochberg false-discovery-rate adjustment within each six-form comparison family.

The Section 2.2 constructive-rate sensitivity uses a grouped-binomial logistic model at the conversation level. The numerator is the number of constructive user turns and the denominator is total user turns in the conversation. The predictors match the any-constructive context model: user framing, task ecology, length bucket and dataset fixed effects. Standard errors use a conversation-robust sandwich estimator. This sensitivity is included because any-constructive-turn outcomes partly reflect the number of turns available for observing at least one constructive turn.

The integrated adjacent-turn regression combines three feature groups in one logistic model: user state before the assistant turn, assistant scaffolding features, and adjacent-turn context. The pooled specification includes dataset fixed effects.

A model/source sensitivity uses exact WildChat model labels, LMSYS model names recovered from conversation identifiers and ShareChat assistant/source families recovered from conversation identifiers. These adjusted estimates are interpreted as measured-covariate-adjusted associations. They do not remove unobserved user motivation, prior expertise, task difficulty or time-varying conversational state.

A.7 Temporal scale

The temporal analyses focus on adjacent-turn ordering because it is the scale at which the logs most directly identify sequence. The adjacent-turn analyses ask whether the immediately following user turn is more likely to be constructive after scaffolded support than after a non-scaffolded reference response, whether that contrast depends on the prior user state and whether support-form signals differ across prior user states. The reverse adjacent-turn analysis asks whether the assistant is more likely to provide scaffolded support after constructive, active or passive user turns.

Coarser before-after summaries ask a different question: whether constructive engagement is higher after the first scaffolded assistant turn than before it. These are retained only as exploratory boundary checks, not as treatment-effect or accumulation estimates. The first scaffolded turn may occur precisely because the user is struggling, because a difficult task has emerged or because the exchange is already developing. The main temporal claim is therefore scale-specific: the clearest evidence concerns local, state- and form-conditioned coupling rather than a demonstrated global trajectory.

Appendix B Robustness and interpretive boundaries

B.1 WildChat model-snapshot robustness

WildChat includes user-facing model metadata that make it possible to repeat major descriptive patterns across model snapshots and coarse model families. These analyses preserved the main directional patterns reported in the main text: coding generally showed higher cognitive engagement than writing; intentionality gaps were visible across most user-facing models; scaffolded conversations were generally more constructive; and adjacent-turn contrasts were more interpretable than coarser within-conversation contrasts. Supplementary Figures D1 and D2 report six-setting temporal and framing diagnostics. Supplementary Figures D3–D7 are WildChat-only model-snapshot checks because comparable user-facing model-snapshot structure is not available across all three corpora.

These checks are used as robustness analyses rather than model-causal comparisons. WildChat model labels are not experimentally assigned. They may be confounded with calendar time, user population, task mix, interface changes and logging context. For this reason, the manuscript does not claim that a specific model family caused the observed engagement differences.

B.2 Cross-dataset and model/source regression checks

The Results claims were checked for cross-dataset consistency at the level appropriate to each estimand. Section 2.1 is a descriptive prevalence claim: cognitive engagement, constructive engagement and scaffolded support were observable in all six settings, with constructive engagement consistently the highest-level and less frequent user-side signal. Section 2.2 is primarily descriptive: user framing, task ecology and interaction depth organize where constructive engagement appears, with the same directional depth gradient across settings. As a simple model-based check, we fitted a conversation-level logistic regression for whether a conversation contained at least one constructive user turn, using user framing, task ecology, length bucket and dataset fixed effects (Supplementary Table C9). Because this any-constructive outcome is partly affected by the number of user turns available for observing at least one constructive turn, we also fitted a grouped-binomial constructive-rate sensitivity model, using constructive user-turn counts out of total user turns as the outcome and the same contextual predictors (Supplementary Table C10). Formal uncertainty estimates are otherwise reported where the claim involves a contrast or model coefficient. Section 2.3 uses conversation-level bootstrap confidence intervals and p values for turn-weighted scaffolded versus non-scaffolded constructive-ratio contrasts, adjusted Poisson/logistic model estimates and an offset-rate Poisson sensitivity using user turns as the exposure (Supplementary Table C8). Section 2.4 reports support-form confidence intervals and between-stratum FDR-adjusted q values in the main figure. Section 2.5 reports conversation-cluster bootstrap confidence intervals for adjacent-turn lifts and cluster-robust regression estimates for the integrated local models.

The main percentage-point contrasts were directionally consistent across the six settings. Constructive-ratio contrasts comparing conversations with versus without scaffolded support were positive and statistically distinguishable from zero in all six settings, with effect sizes ranging from +0.99 to +7.34 percentage points. Adjacent-turn constructive lifts were also positive in all six settings, with five of six contrasts statistically distinguishable from zero; LMSYS writing was positive but weaker (+0.27 percentage points, 95% CI -0.01 to +0.56, $p = .061$). Supplementary Table C7 reports these effect sizes, confidence intervals and p values.

Integrated adjacent-turn regressions were used as robustness checks rather than as the primary narrative claim. These models predicted whether the next user turn was constructive after combining user-state features, assistant-scaffolding features and adjacent-turn context. Because model metadata differ by corpus, we report both within-setting models and pooled specifications. The within-setting models show the dataset-by-task pattern directly, with recoverable model/source fixed effects added separately inside each setting (Supplementary Table C11). The pooled specification includes dataset fixed effects; a sensitivity replaces these with exact WildChat model labels, LMSYS model names recovered from conversation identifiers and ShareChat assistant/source families recovered from conversation identifiers. This is interpreted as model/source adjustment rather than as a causal comparison among model versions. Model/source heterogeneity was detectable, but it did not absorb the scaffolding signal. Adding broad scaffolded support to user/context/model controls improved fit (LR

$\chi^2_1 = 509.5, p < .001$). Because M1–M6 are support forms nested within scaffolded support rather than independent covariates, we interpret the decomposed models as asking which forms carry additional signal among scaffolded turns, not as a test of scaffolded support after removing its own components. Adding M1–M6 descriptors within scaffolded support improved fit (LR $\chi^2_6 = 827.9, p < .001$), and the full scaffolding block remained strongly informative after controls (LR $\chi^2_7 = 1337.5, p < .001$; Supplementary Table C13). Adding prior-state $\times$ support-form interactions also improved fit (pooled model/source fixed effects: LR $\chi^2_{18} = 300.2, p < .001$; Supplementary Table C14). Fully decomposed coefficients clarify the source of this form-level signal: prior user state is the largest predictor of immediate constructive follow-up, and M4 explaining is the most stable positive support-form coefficient (Supplementary Table C12).

B.3 Label uncertainty

The label scheme intentionally separates top-level constructs from finer support forms. Top-level labels such as constructive engagement and scaffolded support have clearer behavioural anchors. Fine-grained labels, especially sparse or boundary-heavy labels, are less stable. This is why the main claims rely most heavily on constructive engagement, scaffolded support presence, support-form patterns that are consistent across settings and temporal contrasts that can be expressed with simple conditional probabilities.

B.4 Selection into support

The most important threat to causal interpretation is selection into support. More complex tasks, more motivated users or already-developing exchanges may both elicit scaffolded support and produce constructive follow-up. Conversation-level adjusted models reduce measured compositional differences but cannot remove unobserved selection. Adjacent-turn analyses reduce the temporal window and condition on local state, but they still cannot identify randomized exposure effects. The safest interpretation is therefore local, state-dependent coupling between assistant support and user engagement.

B.5 Privacy-preserving reporting

The study analyses naturalistic conversations, so the manuscript avoids reproducing raw conversation text. Operational examples are de-identified and paraphrased to illustrate decision rules without exposing user content. Aggregate statistics, model summaries and validation metrics are reported in the manuscript and supplementary tables. The analyses do not link conversations into user histories or estimate longitudinal user-level trajectories; temporal claims are restricted to within-conversation turn ordering.

Appendix C Supporting tables

Table C1 Operational examples for turn-level labels. Examples are de-identified and paraphrased from validation review patterns; they illustrate the decision rule rather than reproducing raw conversation text.

Construct	Paraphrased turn	Annotation rationale
P passive user engagement	"Thanks, that makes sense."	Minimal uptake or acknowledgement without new reasoning, evaluation or task transformation.
A active user engagement	"Can you convert that answer into a command I can run in my setup?"	Task-relevant follow-up or manipulation of supplied information, but without elaborating or testing an idea.
C constructive user engagement	"If the cache expires before the asynchronous callback returns, would the same race condition still happen?"	The user tests an implication, boundary condition or transfer of the prior explanation.
S1 non-scaffolded reference	"Here is the requested function/email draft."	The assistant primarily completes the requested deliverable without regulating the user's process.
S2 scaffolded support	"First isolate the failing step; that will show whether the issue is parsing or state management."	The assistant structures the user's process and leaves diagnostic or revision work with the user.
M1 feedback	"Your diagnosis is close; the bug is in the index offset, not the loop boundary."	Evaluative feedback on the user's prior attempt or understanding.
M2 hinting	"Check the parser log around the first malformed token before changing the model."	A cue or partial path that helps the user proceed without fully resolving the task.
M3 instructing	"First split the condition into two branches, then add a regression test."	Explicit steps or strategy that organizes the user's work.
M4 explaining	"The exception occurs because the event loop is already running, so a nested call cannot acquire it."	A rationale or mechanism that clarifies why the issue occurs.
M5 modeling	"A small pattern you can adapt is: validate input, normalize it, then pass it to the writer."	A sample approach intended to be emulated or adapted rather than merely copied.
M6 questioning	"Which constraint matters more here: preserving tone or shortening the paragraph?"	A question that advances problem solving or reflection.

Table C2 Domain-stratified human–human agreement for labels used in the main analyses. Coding uses the latest three-round validation rerun over 658 task turns, 643 after deduplication; writing uses the latest writing-210 best-pipeline validation. These validation sets assess the coding and writing labelling logic and are not separately powered within each public corpus. F1, Matthews correlation coefficient (MCC) and Gwet’s AC1 are reported per binary label. Coding MCC and AC1 are recomputed from the exported support and positive-count fields; writing metrics are from the exported agreement report. Rare degenerate labels with no positive cases in a validation set are omitted. Writing Active and S1 labels are boundary-heavy and less stable; the primary Results rely on constructive engagement, scaffolded support and support-form labels, for which agreement is substantially stronger.

Domain	Role	Label	F1	MCC	Gwet AC1	Support
Coding	User	Active (A)	0.75	0.72	0.87	315
Coding	User	Constructive (C)	0.85	0.82	0.90	315
Coding	User	Passive (P)	0.88	0.87	0.96	315
Coding	User	Emotional (E)	0.83	0.82	0.97	315
Coding	Assistant	Non-scaffolded reference (S1)	0.78	0.73	0.88	328
Coding	Assistant	Scaffolded support (S2)	0.87	0.69	0.70	328
Coding	Assistant	Feedback (M1)	0.90	0.89	0.95	328
Coding	Assistant	Hinting (M2)	0.79	0.75	0.89	328
Coding	Assistant	Instructing (M3)	0.81	0.77	0.92	328
Coding	Assistant	Explaining (M4)	0.86	0.79	0.84	328
Coding	Assistant	Modeling (M5)	0.84	0.79	0.87	328
Coding	Assistant	Questioning (M6)	0.82	0.81	0.95	328
Writing	User	Active (A)	0.38	0.38	0.69	90
Writing	User	Constructive (C)	0.89	0.83	0.86	90
Writing	User	Passive (P)	0.80	0.81	0.99	90
Writing	User	Emotional (E)	0.67	0.65	0.96	90
Writing	Assistant	Non-scaffolded reference (S1)	0.46	0.43	0.86	120
Writing	Assistant	Scaffolded support (S2)	0.97	0.80	0.93	120
Writing	Assistant	Feedback (M1)	0.89	0.80	0.80	120
Writing	Assistant	Hinting (M2)	0.82	0.70	0.71	120
Writing	Assistant	Instructing (M3)	0.71	0.67	0.86	120
Writing	Assistant	Explaining (M4)	0.84	0.70	0.70	120
Writing	Assistant	Modeling (M5)	0.94	0.91	0.95	120
Writing	Assistant	Questioning (M6)	0.96	0.96	0.99	120

Table C3 Analytic model families used in the main analyses. The table summarizes descriptive and conversation-level model families. Adjusted models are interpreted as covariate-adjusted associations in naturalistic data, not as causal treatment-effect estimates.

Analysis family	Unit and outcome	Main contrast or model	Scope
Engagement prevalence	User turns and conversations; cognitive, constructive and emotional engagement indicators.	Descriptive proportions by task setting, task ecology and user framing.	Establishes where learning-oriented engagement is visible; no causal contrast is implied.
Engagement composition	Cognitively engaged user turns; passive, active and constructive depth levels.	Composition of cognitive engagement within each task setting.	Motivates constructive engagement as the focal high-level outcome while preserving the broader cognitive-engagement hierarchy.
Interaction depth	Conversations; at least one constructive user turn.	Probability of constructive engagement across behavioural-depth buckets.	Depth is interpreted as interactional opportunity and task development, not as a manipulated exposure.
Scaffolding presence	Conversations; constructive-turn count and any constructive turn.	Setting-specific Poisson GLM for constructive-turn count with scaffolded-support presence and observed conversation-level context; no user-turn offset in the primary count model. Setting-specific logistic model for whether any constructive user turn occurs. Estimable covariates are listed in Supplementary Table C5.	Adjusted estimates reduce measured compositional differences but remain associational because support is selected into by task difficulty and user state. Pearson overdispersion was checked and quasi-Poisson scaling preserved the direction of the Poisson estimates.

Table C4 Additional analytic model families and robustness checks. These analyses qualify the main associations by support form, temporal scale and available model or source metadata.

Analysis family	Unit and outcome	Main contrast or model	Scope
Support-form associations	Scaffolded conversations; constructive engagement ratios.	Mean differences for conversations containing a given support-means label versus scaffolded conversations without that label, summarized by user framing and task ecology.	Descriptive because means labels are non-exclusive and can co-occur in the same assistant turn.
Local temporal coupling	Adjacent assistant–user or user–assistant turn pairs; next constructive user turn or next scaffolded assistant turn.	Conditional probability contrasts around adjacent turns; detailed estimands are listed in Supplementary Table C6.	Interpreted as local ordering and state-dependent coupling, not as randomized support effects.
Integrated adjacent-turn regression	Adjacent assistant–user turn pairs; next constructive user turn.	Logistic regression combining prior user state, scaffolded-support presence, support means, task/framing context, prior-state $\times$ support-form interactions and dataset or model/source fixed effects.	Tests whether the local coupling pattern depends jointly on user state and support form; standard errors are clustered by conversation.
Coarse temporal boundary checks	Conversations containing scaffolded support; constructive engagement before and after the first scaffolded turn.	Mean after-minus-before contrast reported as an exploratory appendix check.	Not used as evidence of accumulation because the first scaffolded turn may be selected by difficulty, uncertainty or an already developing exchange.
WildChat model robustness	WildChat conversations or turns; main prevalence, association and temporal outcomes.	Repetition across user-facing model snapshots or coarse model families.	Robustness summaries only; model labels are observational and may be confounded with time, user mix and task mix.

Table C5 Setting-specific covariate sets for the Section 2.3 adjusted models. Each model was fitted separately within the six task settings. Task ecology is therefore fixed within a setting and is not interpreted as a within-setting coefficient. The table lists the estimable covariates used to adjust the scaffolded-support association; pooled context models use task ecology separately as shown in Supplementary Tables [C9](#) and [C10](#).

Settings	Primary Poisson count model	Logistic any-constructive model	Sensitivities and notes
WC coding; LMSYS coding; SC coding; WC writing; LMSYS writing; SC writing	Scaffolded-support presence; user framing; total turns; emotional-expression ratio; observable error marker; persistence-after-failure marker; high-persistence marker.	Scaffolded-support presence; user framing; total turns.	The offset-rate sensitivity used the same Poisson mean model but represented user turns with an exposure offset. User-framing stratified Poisson models omitted user framing because it was fixed within stratum. Error and persistence markers are most interpretable in coding-oriented contexts but are retained as observed conversation-level context when present.

Table C6 Temporal estimands used to characterize local support–engagement coupling. The main temporal analyses use adjacent assistant–user or user–assistant turn pairs. Coarser before–after summaries are retained only as exploratory boundary checks because scaffolded turns are selected by task difficulty, user motivation and conversation trajectory.

Estimand	Definition	Reported in main text	Scope
Overall adjacent-turn lift	Difference in the probability that the next user turn is constructive after scaffolded support versus after a non-scaffolded reference response, estimated within each task setting.	Figure 5a	Local association between assistant support type and immediately following constructive engagement; not a randomized support effect.
Prior-state adjacent contrast	Difference in next-turn constructive probability after scaffolded support versus reference response, separately for prior constructive, active and passive user states.	Figure 5b	Tests whether local support–engagement coupling depends on the user’s immediately preceding state.
Reverse state-conditioned support	Probability that the next assistant turn is scaffolded, estimated after constructive, active and passive user turns.	Figure 5c	Describes how assistant scaffolding supply varies with the user’s current engagement state.
Prior-state support-form regression	Adjacent-turn logistic regression for whether the next user turn is constructive, adding prior-state $\times$ M1–M6 support-form interactions to the integrated model.	Figure 5d and Supplementary Table C14	Tests whether local coupling depends jointly on the user’s immediately preceding engagement state and the form of scaffolded support.
Around-first-scaffold contrast	Mean constructive engagement after the first scaffolded assistant turn minus mean constructive engagement before that turn, among conversations containing scaffolded support.	Supplementary Figure D1	Exploratory boundary check only; the first scaffolded turn may be selected by difficulty, uncertainty or an already developing exchange.

Table C7 Uncertainty estimates for key support–engagement contrasts.

Constructive-ratio contrasts are turn-weighted conversation-bootstrap estimates, with ratios computed as total constructive user turns divided by total user turns within scaffolded versus non-scaffolded conversation groups. Post-answer depth contrasts use conversation-level bootstrap resampling. Adjacent-turn contrasts bootstrap conversations over assistant-to-user pairs. Percentage-point and turn differences are the effect sizes reported in the main figures.

Setting	Contrast	Effect	95% CI	p value
WC coding	Scaffolded – reference constructive ratio	+4.56 pp	[+4.15, +4.99]	< .001
WC coding	Scaffolded – reference post-answer depth	+2.23 turns	[+2.11, +2.35]	< .001
LMSYS coding	Scaffolded – reference constructive ratio	+2.50 pp	[+2.19, +2.83]	< .001
LMSYS coding	Scaffolded – reference post-answer depth	+1.42 turns	[+1.34, +1.52]	< .001
SC coding	Scaffolded – reference constructive ratio	+7.34 pp	[+5.29, +9.27]	< .001
SC coding	Scaffolded – reference post-answer depth	+3.36 turns	[+2.69, +4.01]	< .001
WC writing	Scaffolded – reference constructive ratio	+0.99 pp	[+0.82, +1.16]	< .001
WC writing	Scaffolded – reference post-answer depth	+3.13 turns	[+2.96, +3.30]	< .001
LMSYS writing	Scaffolded – reference constructive ratio	+1.09 pp	[+0.87, +1.31]	< .001
LMSYS writing	Scaffolded – reference post-answer depth	+2.10 turns	[+1.91, +2.28]	< .001
SC writing	Scaffolded – reference constructive ratio	+3.20 pp	[+1.55, +4.89]	< .001
SC writing	Scaffolded – reference post-answer depth	+3.38 turns	[+1.92, +4.64]	< .001
WC coding	Adjacent scaffolded – reference lift	+2.68 pp	[+2.22, +3.09]	< .001
LMSYS coding	Adjacent scaffolded – reference lift	+2.01 pp	[+1.54, +2.50]	< .001
SC coding	Adjacent scaffolded – reference lift	+6.21 pp	[+4.35, +8.10]	< .001
WC writing	Adjacent scaffolded – reference lift	+1.13 pp	[+0.88, +1.41]	< .001
LMSYS writing	Adjacent scaffolded – reference lift	+0.27 pp	[−0.01, +0.56]	.061
SC writing	Adjacent scaffolded – reference lift	+3.35 pp	[+0.80, +6.62]	.025

Table C8 Offset-rate sensitivity for scaffolded-support Poisson models. The primary Fig. 3b Poisson model uses raw constructive-turn counts with conversation length entered as a covariate. As a rate sensitivity, this table refits the same setting-specific mean model with `offset(log(user_turns))`. Cells report the scaffolded-support rate ratio with 95% confidence intervals and two-sided p values using quasi-Poisson scaled standard errors.

Setting	Offset rate ratio [95% CI], p value	Pearson dispersion	Conversations
WC coding	1.57 [1.47, 1.68], < .001	1.29	31,878
LMSYS coding	1.47 [1.38, 1.56], < .001	1.13	32,114
SC coding	1.58 [1.31, 1.91], < .001	1.29	2,481
WC writing	1.36 [1.26, 1.46], < .001	1.15	39,534
LMSYS writing	1.70 [1.50, 1.94], < .001	1.12	21,023
SC writing	1.65 [1.27, 2.15], < .001	1.32	1,539

Table C9 Conversation-level context model for constructive engagement. The outcome is whether a conversation contains at least one constructive user turn. The logistic regression includes user framing, task ecology, conversation-length bucket and dataset fixed effects. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-robust standard errors. The reference category is unintentional framing, writing-oriented task, 2–3 user turns and WildChat.

Predictor	OR [95% CI], p value
Intentional framing	3.03 [2.93, 3.13], < .001
Coding task ecology	1.89 [1.71, 2.08], < .001
4–6 user turns	2.07 [1.99, 2.15], < .001
7+ user turns	3.58 [3.42, 3.74], < .001
LMSYS Chat	0.70 [0.68, 0.73], < .001
ShareChat	2.29 [2.12, 2.46], < .001
Conversations	128,569

Table C10 Conversation-level constructive-rate sensitivity model. The outcome is the number of constructive user turns out of total user turns in each conversation, modelled with a grouped-binomial logistic regression. The model includes user framing, task ecology, conversation-length bucket and dataset fixed effects, using user-turn counts as denominators to reduce the mechanical opportunity effect in any-constructive-turn outcomes. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-robust standard errors. The reference category is unintentional framing, writing-oriented task, 2–3 user turns and WildChat.

Predictor	OR [95% CI], p value
Intentional framing	2.78 [2.68, 2.88], < .001
Coding task ecology	1.66 [1.47, 1.88], < .001
4–6 user turns	1.13 [1.09, 1.18], < .001
7+ user turns	1.06 [1.01, 1.11], .010
LMSYS Chat	0.72 [0.69, 0.75], < .001
ShareChat	2.30 [2.13, 2.50], < .001
Conversations	128,569
User turns	491,685
Constructive user turns	24,227

Table C11 Setting-level integrated adjacent-turn models. Each column reports a separate within-setting logistic regression for whether the next user turn is constructive. The scaffolded-support row comes from a broad scaffolded-support model. Support-form rows come from a decomposed model that includes broad scaffolded-support presence and co-occurring M1–M6 forms, so form coefficients compare variation within scaffolded support rather than replacing the broad scaffolded-support contrast. All models include prior user state, intentional framing, assistant-turn index and recoverable model/source fixed effects where available. Cells report odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors.

Predictor	WC coding	LMSYS coding	SC coding	WC writing	LMSYS writing	SC writing
A2U pairs	92,042	71,910	9,001	124,214	55,580	5,825
Conversations	31,809	31,606	2,474	39,455	20,936	1,537
Model/source FE count	13	23	4	13	23	3
Scaffolded support	1.13 [1.08, 1.19] < .001	1.25 [1.16, 1.34] < .001	1.24 [1.09, 1.42] .001	1.57 [1.38, 1.78] < .001	1.04 [0.86, 1.26] .679	1.34 [0.98, 1.85] .068
Prior user constructive	5.20 [4.84, 5.59] < .001	6.01 [5.43, 6.66] < .001	4.13 [3.46, 4.93] < .001	8.28 [6.73, 10.19] < .001	11.89 [9.10, 15.54] < .001	6.23 [4.43, 8.76] < .001
Prior user active	1.67 [1.57, 1.76] < .001	1.58 [1.46, 1.71] < .001	1.56 [1.32, 1.83] < .001	1.91 [1.68, 2.18] < .001	2.18 [1.80, 2.64] < .001	1.47 [1.07, 2.01] .017
Prior user passive	2.19 [1.65, 2.92] < .001	2.01 [1.38, 2.92] < .001	1.82 [1.25, 2.66] .002	2.49 [1.69, 3.65] < .001	2.32 [1.29, 4.17] .005	0.86 [0.40, 1.84] .703
Intentional framing	1.57 [1.49, 1.66] < .001	1.46 [1.35, 1.58] < .001	1.57 [1.34, 1.83] < .001	1.02 [0.89, 1.16] .798	0.86 [0.70, 1.06] .150	1.37 [0.93, 2.02] .107
M1 feedback	1.19 [1.07, 1.31] .001	1.21 [0.98, 1.50] .079	1.28 [1.03, 1.60] .025	0.69 [0.53, 0.91] .008	0.81 [0.51, 1.31] .397	1.74 [1.12, 2.71] .013
M2 hinting	1.04 [0.95, 1.14] .429	0.90 [0.72, 1.11] .332	0.85 [0.65, 1.10] .209	1.52 [1.20, 1.93] < .001	1.63 [1.09, 2.44] .018	0.76 [0.45, 1.29] .314
M3 instructing	0.89 [0.83, 0.96] .003	0.84 [0.69, 1.04] .103	0.93 [0.77, 1.13] .491	1.36 [1.04, 1.78] .027	1.06 [0.71, 1.58] .779	0.94 [0.45, 1.94] .860
M4 explaining	1.21 [1.09, 1.35] < .001	1.36 [1.08, 1.72] .008	1.77 [1.35, 2.33] < .001	1.81 [1.40, 2.34] < .001	1.12 [0.75, 1.68] .570	1.09 [0.58, 2.06] .789
M5 modelling	0.86 [0.80, 0.93] < .001	0.73 [0.63, 0.86] < .001	0.76 [0.60, 0.96] .021	0.69 [0.47, 1.01] .055	1.04 [0.64, 1.68] .872	1.57 [0.94, 2.61] .084
M6 questioning	0.73 [0.62, 0.86] < .001	0.30 [0.21, 0.44] < .001	0.70 [0.54, 0.90] .006	1.13 [0.74, 1.72] .567	0.66 [0.36, 1.21] .179	1.66 [0.90, 3.08] .106

Table C12 Integrated adjacent-turn regression combining user state, assistant scaffolding and model controls. The outcome is whether the next user turn is constructive. The scaffolded-support row is estimated in a broad scaffolded-support model; M1, M4 and M6 rows are estimated in a decomposed support-form model that includes broad scaffolded-support presence and co-occurring support forms. Estimates are odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors. The primary pooled model uses dataset fixed effects; the model/source sensitivity adds recoverable assistant model or source-family fixed effects.

Predictor	Pooled dataset FE	Pooled model/source FE	WildChat model FE
Scaffolded support	1.50 [1.44, 1.55], < .001	1.47 [1.41, 1.52], < .001	1.56 [1.49, 1.63], < .001
Prior user constructive	8.90 [8.45, 9.38], < .001	8.55 [8.11, 9.01], < .001	9.11 [8.53, 9.74], < .001
Prior user active	2.62 [2.52, 2.72], < .001	2.57 [2.47, 2.67], < .001	2.85 [2.71, 3.00], < .001
Prior user passive	2.08 [1.76, 2.44], < .001	2.07 [1.76, 2.44], < .001	2.32 [1.85, 2.90], < .001
Intentional framing	1.40 [1.35, 1.46], < .001	1.42 [1.37, 1.48], < .001	1.40 [1.33, 1.47], < .001
Coding task	1.37 [1.22, 1.54], < .001	1.37 [1.22, 1.54], < .001	1.21 [1.05, 1.40], .009
M1 feedback	0.79 [0.73, 0.86], < .001	0.80 [0.74, 0.86], < .001	0.74 [0.67, 0.82], < .001
M4 explaining	2.15 [1.99, 2.32], < .001	2.12 [1.96, 2.30], < .001	2.11 [1.92, 2.31], < .001
M6 questioning	0.82 [0.72, 0.93], .002	0.84 [0.74, 0.96], .008	1.08 [0.93, 1.25], .340
Model/source block	Reference	LR $\chi^2_{41} = 714.4$, $p < .001$	LR $\chi^2_{13} = 553.2$, $p < .001$

Table C13 Incremental contribution of scaffolding features in integrated adjacent-turn models. Likelihood-ratio block tests compare nested logistic regressions for whether the next user turn is constructive. All models include prior user state, user framing, task ecology, assistant-turn index and the indicated dataset or model/source fixed effects.

Scope	Added feature block	LR χ^2 (df)	p value
six task settings, dataset FE	Broad scaffolded support after user/context controls	584.6 (1)	< .001
six task settings, dataset FE	M1–M6 descriptors within scaffolded support	896.7 (6)	< .001
six task settings, dataset FE	Full scaffolding block after user/context controls	1481.3 (7)	< .001
six task settings, model/source FE	Broad scaffolded support after user/context controls	509.7 (1)	< .001
six task settings, model/source FE	M1–M6 descriptors within scaffolded support	829.8 (6)	< .001
six task settings, model/source FE	Full scaffolding block after user/context controls	1339.6 (7)	< .001
WildChat only, model FE	Broad scaffolded support after user/context controls	418.6 (1)	< .001
WildChat only, model FE	M1–M6 descriptors within scaffolded support	519.4 (6)	< .001
WildChat only, model FE	Full scaffolding block after user/context controls	937.9 (7)	< .001

Table C14 Prior user state and support form jointly structure adjacent-turn constructive engagement. The first panel reports likelihood-ratio tests adding prior-state $\times$ M1–M6 interactions to the integrated adjacent-turn logistic model. The second panel reports pooled model/source fixed-effect estimates for selected support forms within each prior user state. Estimates are odds ratios with 95% confidence intervals and two-sided p values from conversation-cluster robust standard errors.

Panel	Scope / prior state	Term or test	Estimate
Interaction block	six task settings, dataset FE	Prior state $\times$ M1–M6	LR $\chi^2_{18} = 305.4$, < .001
Interaction block	six task settings, model/source FE	Prior state $\times$ M1–M6	LR $\chi^2_{18} = 297.3$, < .001
Interaction block	WildChat only, model FE	Prior state $\times$ M1–M6	LR $\chi^2_{18} = 186.0$, < .001
State-stratified model/source FE	prior constructive	M1 feedback	0.72 [0.64, 0.81], < .001
State-stratified model/source FE	prior constructive	M4 explaining	1.72 [1.45, 2.04], < .001
State-stratified model/source FE	prior active	M1 feedback	0.83 [0.72, 0.95], .008
State-stratified model/source FE	prior active	M4 explaining	2.13 [1.88, 2.41], < .001
State-stratified model/source FE	prior passive	M1 feedback	2.16 [1.16, 4.04], .015
State-stratified model/source FE	prior passive	M4 explaining	2.88 [1.55, 5.35], < .001

Table C15 Support-intent labels are dominated by cognitive support. Scaffolded assistant turns can receive metacognitive (I1), cognitive (I2) or affective (I3) intent labels. Scaffolded-turn prevalence reports the share of scaffolded assistant turns containing each intent label. Constructive ratios compare scaffolded conversations containing at least one assistant turn with the intent label with scaffolded conversations without that intent label, using turn-weighted constructive user-turn ratios. Intent labels can co-occur, and affective intent is sparse, so these contrasts are interpreted descriptively.

Scope	Intent label	Scaffolded turns (%)	With intent (%)	Without intent (%)	Diff. (pp)	Conv. with / without
Pooled	I1 metacognitive	8.7	5.0	6.9	-2.0	10,924 / 55,892
Pooled	I2 cognitive	91.6	6.8	2.8	+4.0	61,942 / 4,874
Pooled	I3 affective	0.8	12.4	6.4	+6.0	1,052 / 65,764
Coding	I1 metacognitive	8.2	7.0	9.6	-2.6	6,782 / 36,179
Coding	I2 cognitive	92.3	9.4	3.9	+5.6	40,279 / 2,682
Coding	I3 affective	1.0	16.1	8.9	+7.2	778 / 42,183
Writing	I1 metacognitive	9.6	2.5	2.8	-0.3	4,142 / 19,713
Writing	I2 cognitive	90.4	2.8	1.8	+0.9	21,663 / 2,192
Writing	I3 affective	0.6	3.9	2.7	+1.2	274 / 23,581

Table C16 Support-supply profile within scaffolded assistant turns. Values show the percentage of scaffolded assistant turns containing each scaffolding-means label. Support-form labels are not mutually exclusive, so rows do not sum to 100%. Coding-oriented scaffolded turns are dominated by explanation-heavy support, whereas writing-oriented scaffolded turns show a more mixed support profile.

Setting	M1 Feedback	M2 Hinting	M3 Instructing	M4 Explaining	M5 Modelling	M6 Questioning
WildChat coding	8.8	14.7	29.7	84.0	23.9	6.9
LMSYS coding	4.4	7.6	9.7	76.7	16.5	13.9
ShareChat coding	13.3	12.7	25.7	81.2	18.4	18.6
WildChat writing	16.6	9.6	59.2	29.5	15.0	6.9
LMSYS writing	11.3	10.9	52.3	26.9	10.3	13.4
ShareChat writing	24.2	23.2	22.9	45.2	23.5	38.1

Table C17 Boundary-check corpora processed with the common pipeline. SWE-chat and ThoughtTrace were labelled and analysed with the same Level 1–4 workflow, but are not part of the main six-setting analysis because they do not provide stable coding–writing task comparisons or stable non-scaffolded reference variation. Percentages for cognitive, constructive and emotional engagement use user turns as the denominator; scaffolded-support percentages use assistant turns as the denominator.

Panel A: retained sample and descriptive rates.

Corpus / setting	Conv.	User turns	Assistant turns	Scaffolded turns (%)	Cognitive overall (%)	Constructive level (%)
SWE-chat coding	1,536	9,506	6,813	19.45	32.36	17.96
SWE-chat writing	3	18	14	42.86	22.22	11.11
ThoughtTrace coding	56	251	252	85.71	64.54	9.96
ThoughtTrace writing	61	301	300	67.33	51.16	2.66

Panel B: boundary-check interpretation.

Corpus / setting	Use in manuscript	Boundary-check implication
SWE-chat coding	Coding-only appendix check	Scaffolded conversations were more constructive than non-scaffolded conversations (0.173 vs. 0.133) and deeper after the first answer (+6.41 turns), but the corpus lacks a usable writing counterpart.
SWE-chat writing	Excluded from inference	Only three conversations remained after filtering, so estimates are not interpretable.
ThoughtTrace coding	Feasibility check only	All conversations contained scaffolded support, leaving no non-scaffolded conversation-level reference group for support-association estimates.
ThoughtTrace writing	Feasibility check only	Scaffolded support remained very common and only two conversations lacked scaffolded support, making adjusted support contrasts unstable.

Appendix D Supporting figures

The supporting figures in Appendix D separate two roles. Supplementary Figures D1 and D2 extend the main analyses across the six WildChat, LMSYS Chat and ShareChat task settings. Supplementary Figures D3–D7 are WildChat-only robustness displays because WildChat exposes user-facing model-snapshot metadata and a large enough robustness subset for these additional checks. LMSYS Chat and ShareChat do not provide directly comparable model-snapshot structure in the current analysis artifacts. The model-snapshot figures are descriptive robustness checks because model labels in naturalistic corpora are observational rather than experimentally assigned.




Fig. D1 Coarse around-first-scaffold contrasts are exploratory boundary checks across the six task settings. Bars show the difference in constructive engagement after versus before the first scaffolded assistant turn within conversations containing scaffolded support. Error bars show 95% conversation-bootstrap confidence intervals. These before–after contrasts are not used as treatment-effect or accumulation estimates because the first scaffolded turn may be selected by difficulty, uncertainty or an already developing exchange.




Fig. D2 Scaffolded-support associations by user framing across the six task settings. Stratified Poisson count-ratio estimates compare the association between scaffolded-support presence and constructive participation in intentional versus unintentional conversations. Error bars show 95% model-based confidence intervals. The figure is interpreted as descriptive moderation rather than evidence that framing causally changes support effectiveness, because user framing, task complexity and persistence are not randomly assigned.




Fig. D3 WildChat model-snapshot prevalence profiles preserve the main domain contrast. Across user-facing WildChat model snapshots, coding-oriented conversations generally show higher cognitive engagement and higher assistant scaffolding rates than writing-oriented conversations. The model labels are observational and may reflect time, user mix and traffic composition, so the figure is used as a robustness check rather than a model-causal comparison.




Fig. D4 Intentionality gaps are visible across WildChat model snapshots. The figure reports model-specific differences between intentional and unintentional conversations in cognitive and constructive engagement, separately for coding and writing. Most model snapshots preserve the directional pattern that intentionally framed conversations contain richer learning-oriented participation.




Fig. D5 Conversation-level support associations are broadly positive across WildChat model snapshots. The figure repeats major conversation-level analyses by user-facing WildChat model snapshot, including constructive participation, adjusted support associations and exploratory coarse before–after contrasts around the first scaffolded assistant turn. The pattern supports the distinction between stable conversation-level associations and more local adjacent-turn temporal claims.




Fig. D6 Local temporal signals vary across WildChat model snapshots. The figure reports adjacent-turn constructive lifts and emotion-to-support responsiveness by model snapshot. Adjacent-turn constructive lifts are often positive, especially in coding-oriented conversations, but the magnitude varies across snapshots; emotion-to-support responsiveness is more heterogeneous and is treated as exploratory.




Fig. D7 Coarse model-family aggregation preserves the main directional patterns. Wild-Chat model snapshots are grouped into coarse model families to summarize cognitive engagement prevalence, adjusted support associations and adjacent-turn constructive lift. Coding remains higher than writing across families, and the support-association and local-lift patterns remain directionally consistent on average.